

Machine Learning and Compiler Optimization

P. (Saday) Sadayappan

Oct. 28, 2024

Compiler Optimization: Strengths & Limitations

- Compilers very effective in lowering sequential high-level language programs to minimize the number of executed low-level instructions
 - Production compilers like gcc, llvm/clang, Intel icc/icx, IBM xlc can exploit ILP effectively
 - Many mature optimizations: register allocation, common sub-expr elimination...
- However, the dominant cost is data movement and not the arithmetic ops.
 - Data movement can be orders of magnitude more expensive, in terms of energy & time

```
for (i=0; i<n; i++)  
    a[i]= s*a[i]+a[i];
```



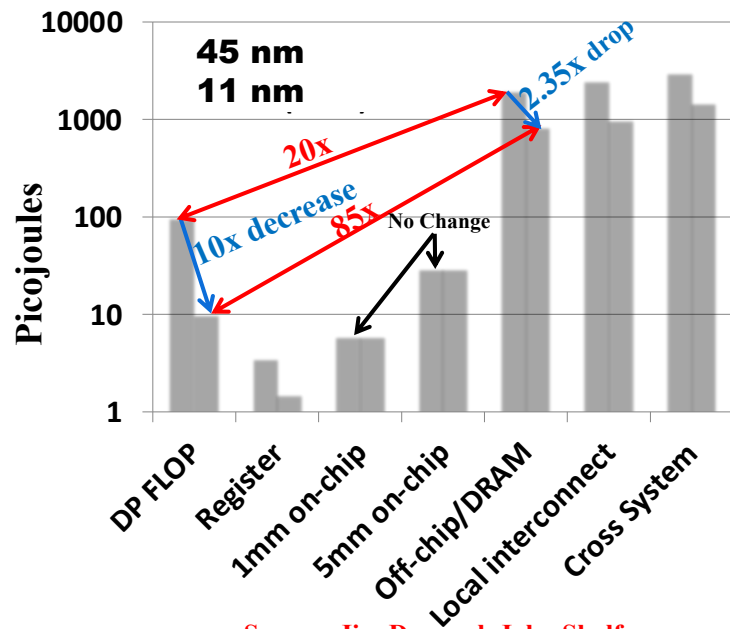
```
.L13:  
    vmovaps (%r11,%rax), %ymm0  
    addl    $1, %ecx  
    vmulps  %ymm2, %ymm0, %ymm1  
    vaddps  %ymm1, %ymm0, %ymm0  
    vmovaps %ymm0, (%r11,%rax)  
    addq    $32, %rax  
    cmpl    %ecx, %r9d  
    ja      .L13
```

AVX code

Data Movement vs. Computation (Energy and Throughput)

- Energy per FLOP is orders of magnitude lower than data movement
 - Technology scaling lowered energy for arithmetic much more than for data movement
- Peak memory bandwidth is rising slower than peak performance

Data Movement Cost: Energy Trends



Source: Jim Demmel, John Shalf

Roofline Chart for 3 Nvidia GPU Generations





Illustration: Data Movement is Expensive

- The main cause of performance loss is data movement overheads
 - Between nodes in a multi-node system
 - Through the memory hierarchy at each node
- Illustrative synthetic example
 - Functionally identical codes with very different performance on my laptop


```
double A[4096][4096];  
// Initialize A[][] to 0.0  
for (t=0; t<64; t++)  
  for (j=0; j<4096; j++)  
    for (i=0; i<4096; i++)  
      A[i][j] += (2*t+j)+(3*t+i);
```

6.4 sec



```
double A[4096][4096];  
// Initialize A[][] to 0.0  
for (i=0; i<4096; i++)  
  for (j=0; j<4096; j++)  
    for (t=0; t<64; t++)  
      A[i][j] += (2*t+j)+(3*t+i);
```


0.26 sec



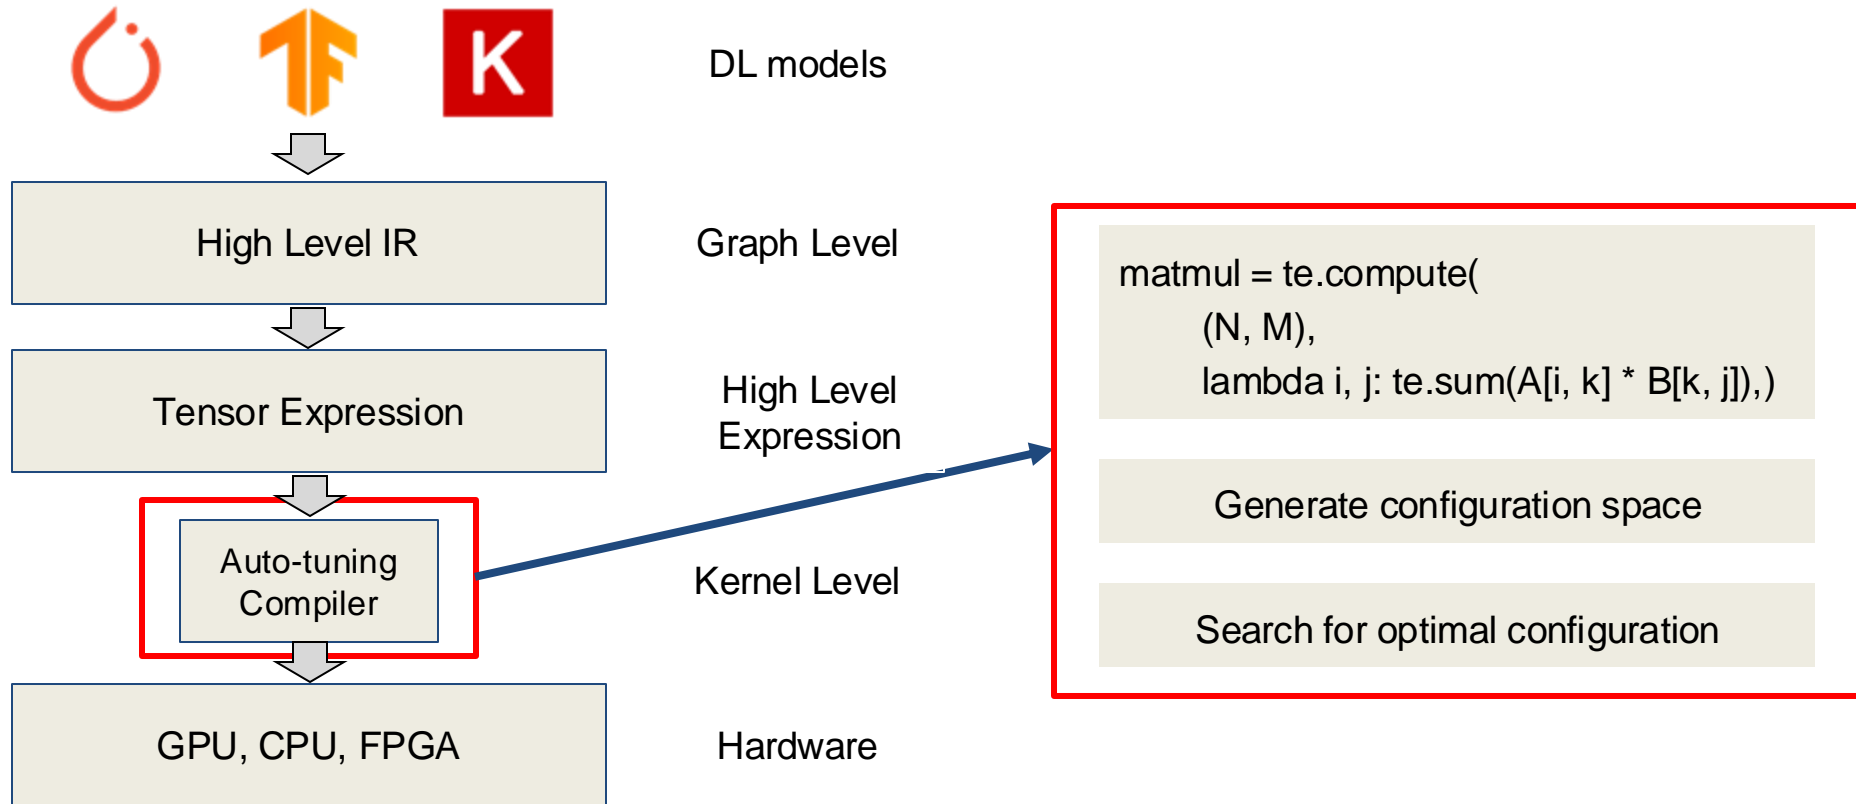
High-Level Summary

- Many efforts on using ML for Compiler Optimization
 - But few (so far) have addressed problems of significance
 - Examples: Best loop-unroll factor; Device selection (CPU or GPU)
 - Or the comparison baselines have been weak
 - Ideally, performance should match/exceed manual optimization
- Significant SW Problem: Developing high-performance applications for parallel/heterogeneous systems
 - Challenges: Application developer productivity & performance portability
- One impressive ML-based compiler: TVM/Ansor
 - Automated synthesis of high-performance code
 - Multi-core CPUs, GPUs, FPGA
 - Input: High-level “Einsum” tensor expression
 - Performance exceeds highly tuned vendor libraries (MKL, cuDNN)
 - But only fixed operator sizes (OK for ML inference pipelines)
- Can we expand the scope of effective ML-based high-perf. code synthesis beyond what TVM/Ansor can now achieve?

TVM/Ansor Performance: ResNet



TVM/Ansor Auto-Tuning Compiler



TVM/Ansor: Code Schema for Nvidia GPU

```
1 // blockIdx.x i.0@j.0@ (None)
2 for i.0 (None)
3   for j.0 (None)
4     //vthread i.1@j.1@ (None)
5     for i.1 (None)
6       for j.1 (None)
7         // threadIdx.x i.2@j.2@ (None)
8         for i.2 (None)
9           for j.2 (None)
10            // thread level code Line 10 - 22
11            for k.0 (None)
12              // shared memory buffer loading Line 13- 14
13              B.shared[...] = B[...]
14              A.shared[...] = A[...]
15              __syncthreads();
16              // register level
17              for k.1 (None) for k.2 (None)
18                for i.3 (None) for i.4 (None)
19                  for j.3 (None) for j.4 (None)
20                    matmul.local = ...
21            // store output to global memory
22            matmul[...] = matmul.local[...]
```

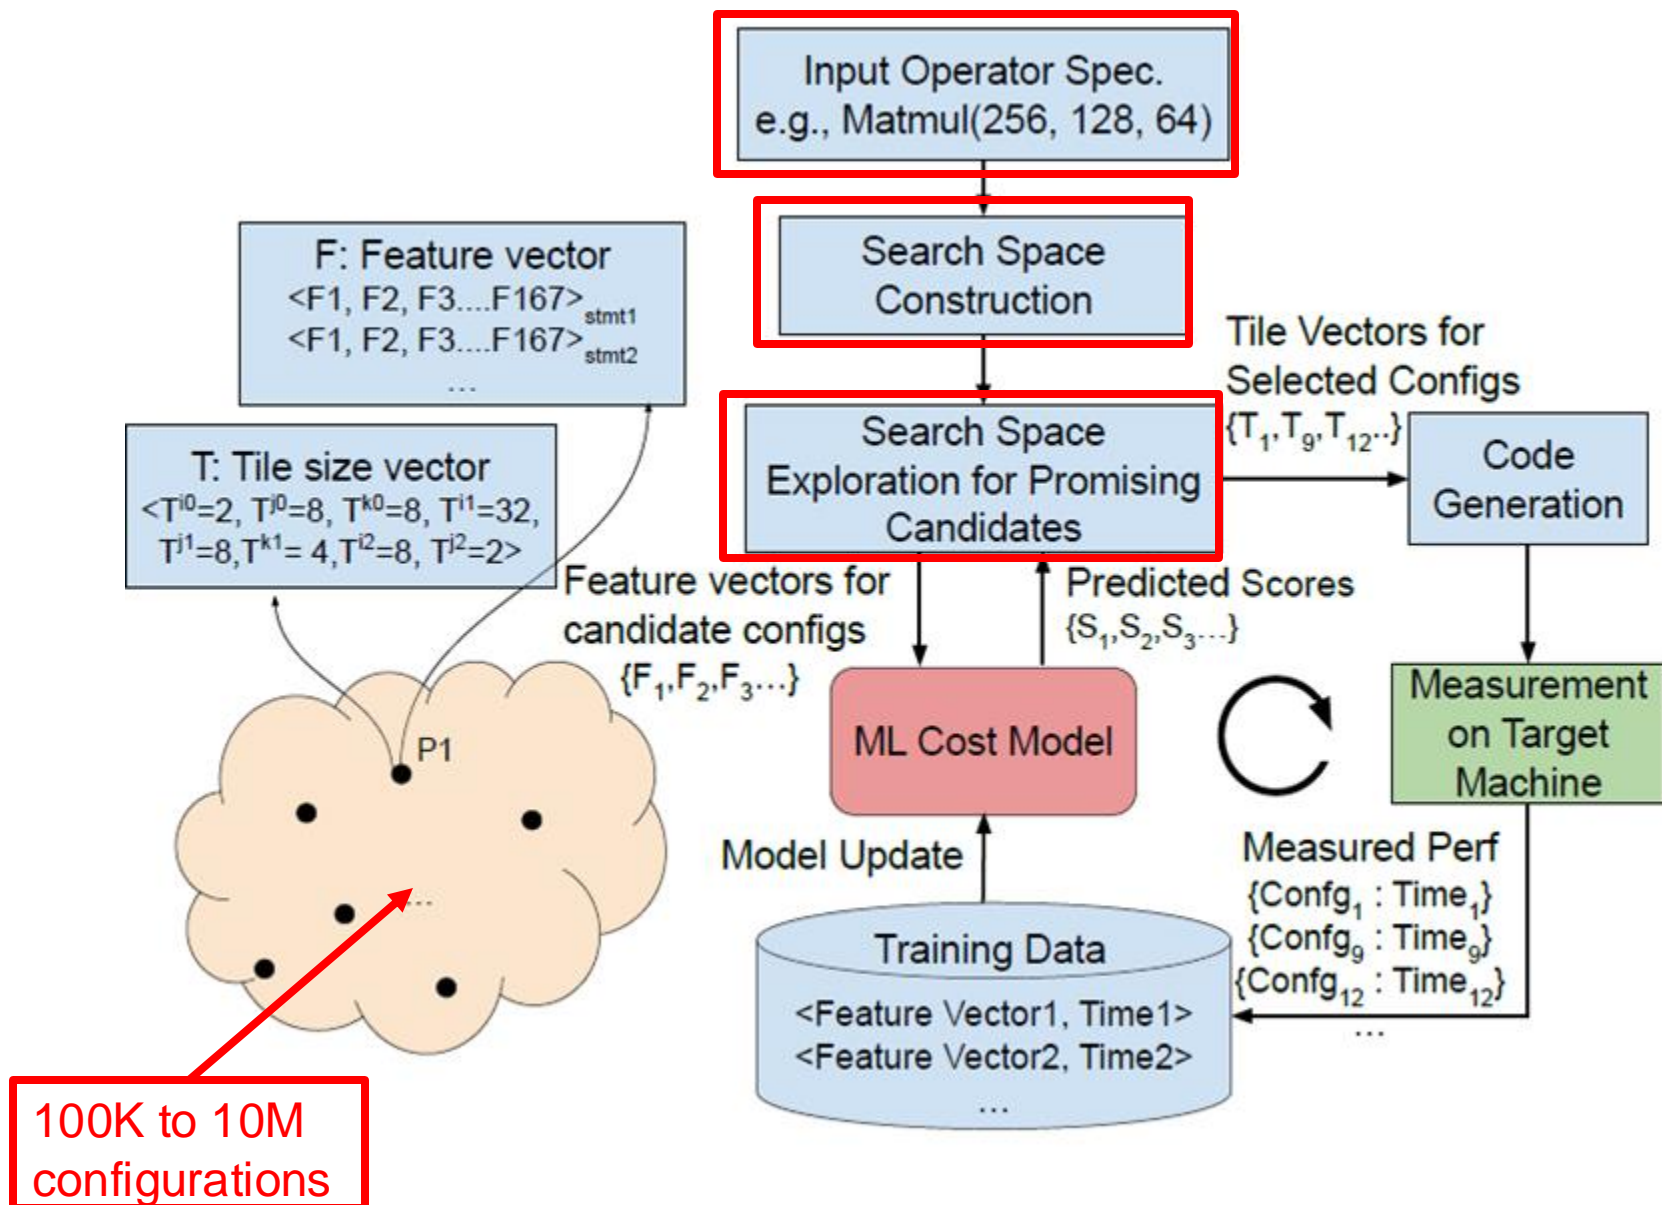
```
matmul = te.compute(
    (N, M),
    lambda i, j: te.sum(A[i, k] * B[k, j]),)
```

Generate configuration space

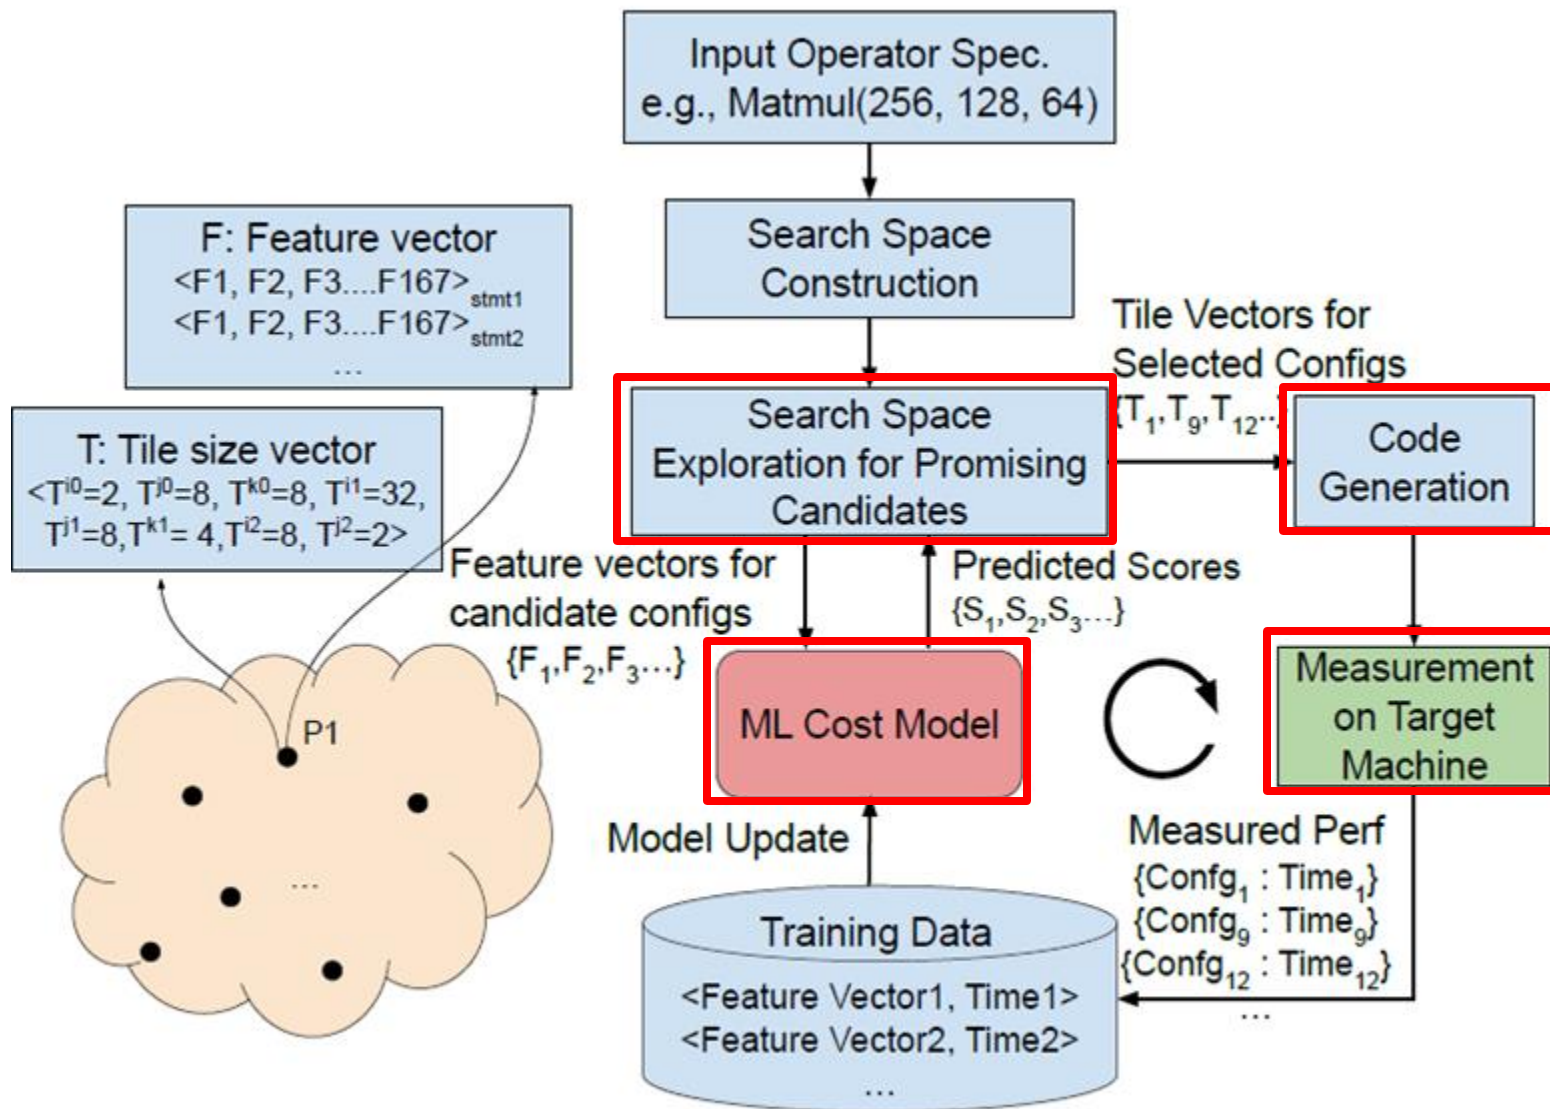
Search for optimal configuration

- Multi-level tiled code
 - 5-level tiling of each "parallel" loop
 - 3-level tiling of "reduction" loops
- Configuration: Set of tile sizes
 - $\langle i_0, i_1, i_2, i_3, i_4, j_0, j_1, j_2, j_3, j_4, k_0, k_1, k_2 \rangle$
- XGBoost proxy perf. model
 - Ansor uses a 167-component feature vector for training model
 - $\langle t_1, t_2, \dots, t_k \rangle \Rightarrow \langle f_1, f_2, \dots, f_{167} \rangle$
 - 167 features generated from AST
 - op count, buffer sizes, ...


TVM/Ansor ML-Driven Compiler



TVM/Ansor ML-Driven Compiler



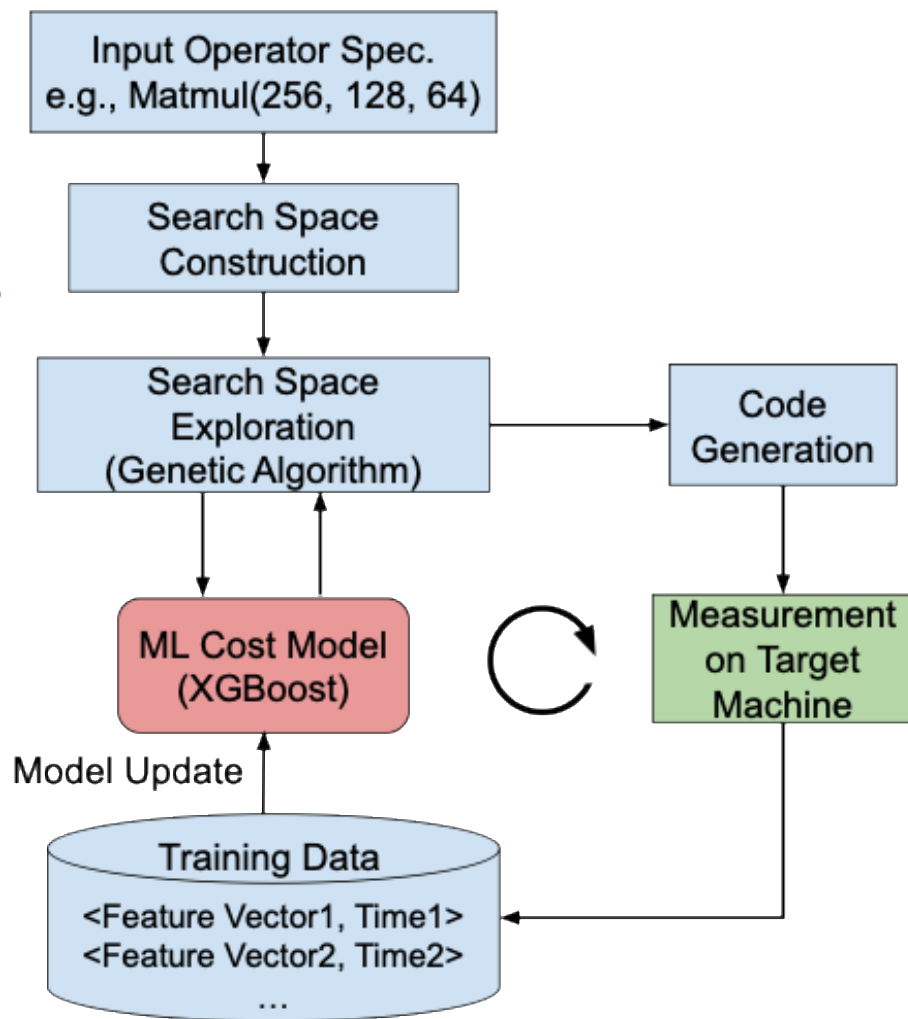
TVM/Ansor Auto-Tuning



- ML perf. model is updated after every batch of 64 configs.
 - Model very effective in filtering out bad configs. After 2-3 batches
- Across-run variability can be as high as 20%

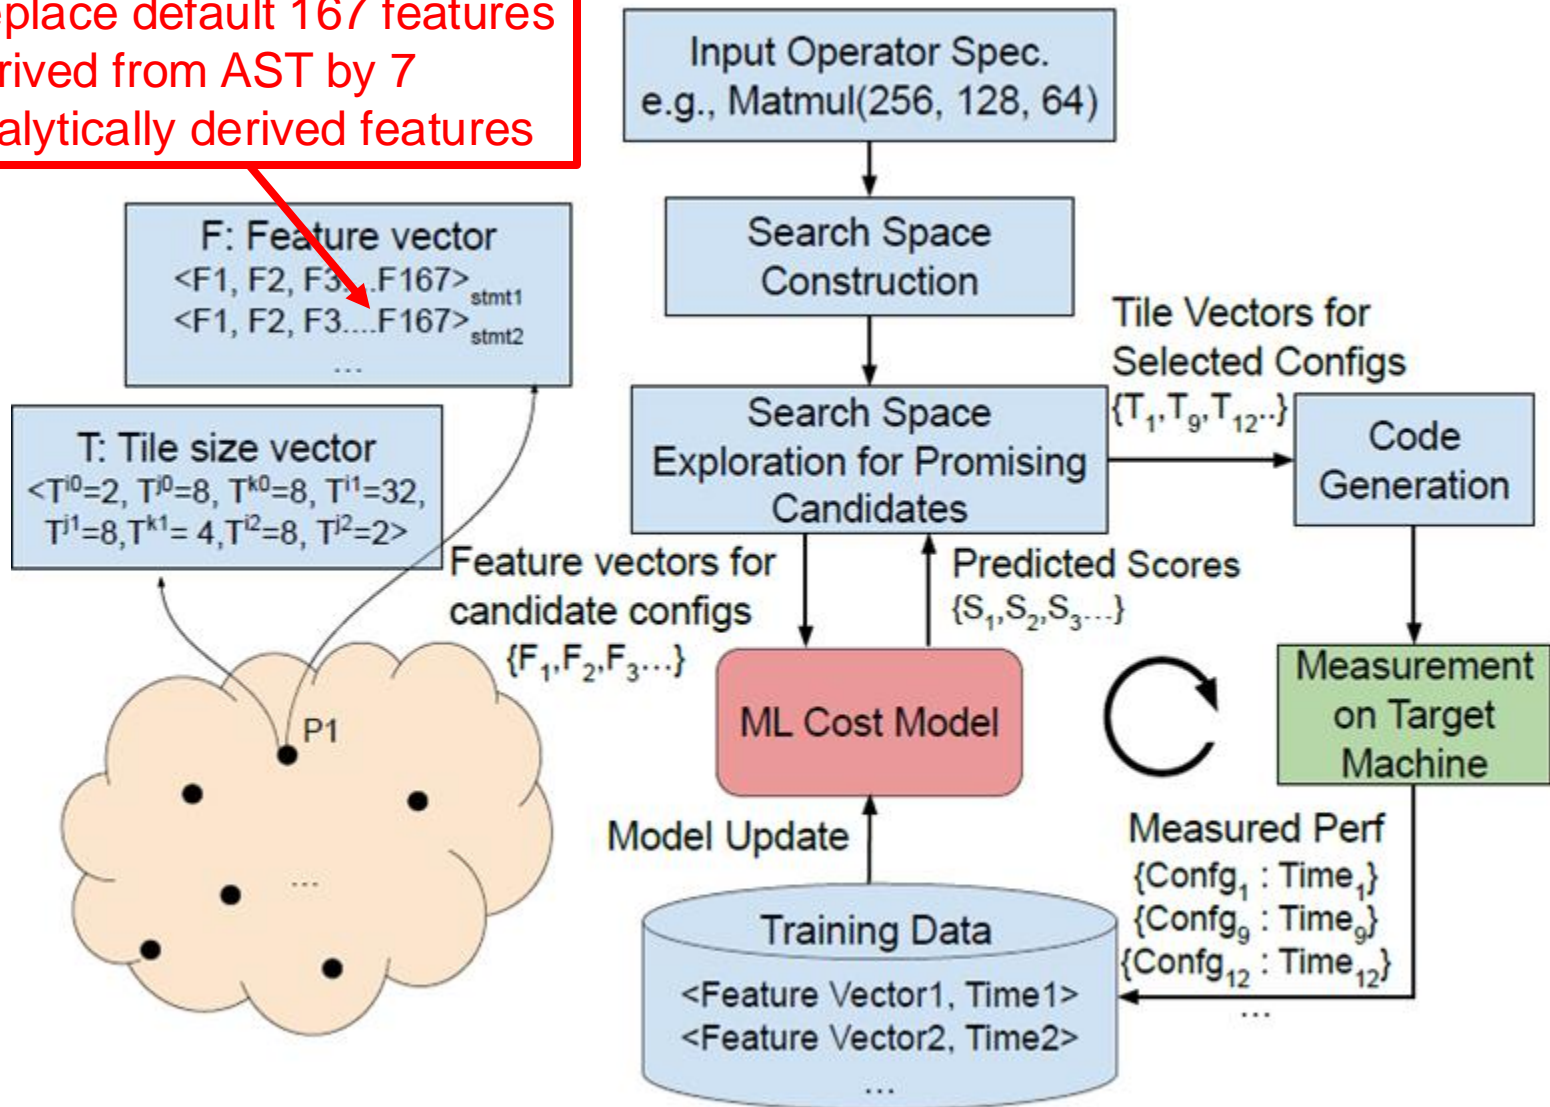
Initial Explorations with TVM/Ansor

- How important is the accuracy of proxy ML model?
 - Can it be improved by using analytically derived features for concurrency, data movement vol...?
- How important is the search technique
 - Can an alternate search strategy improve auto-tuning time and/or performance?
 - Sampled gradient descent
 - Bayesian optimization



Exploration 1: Use Analytical Modeling

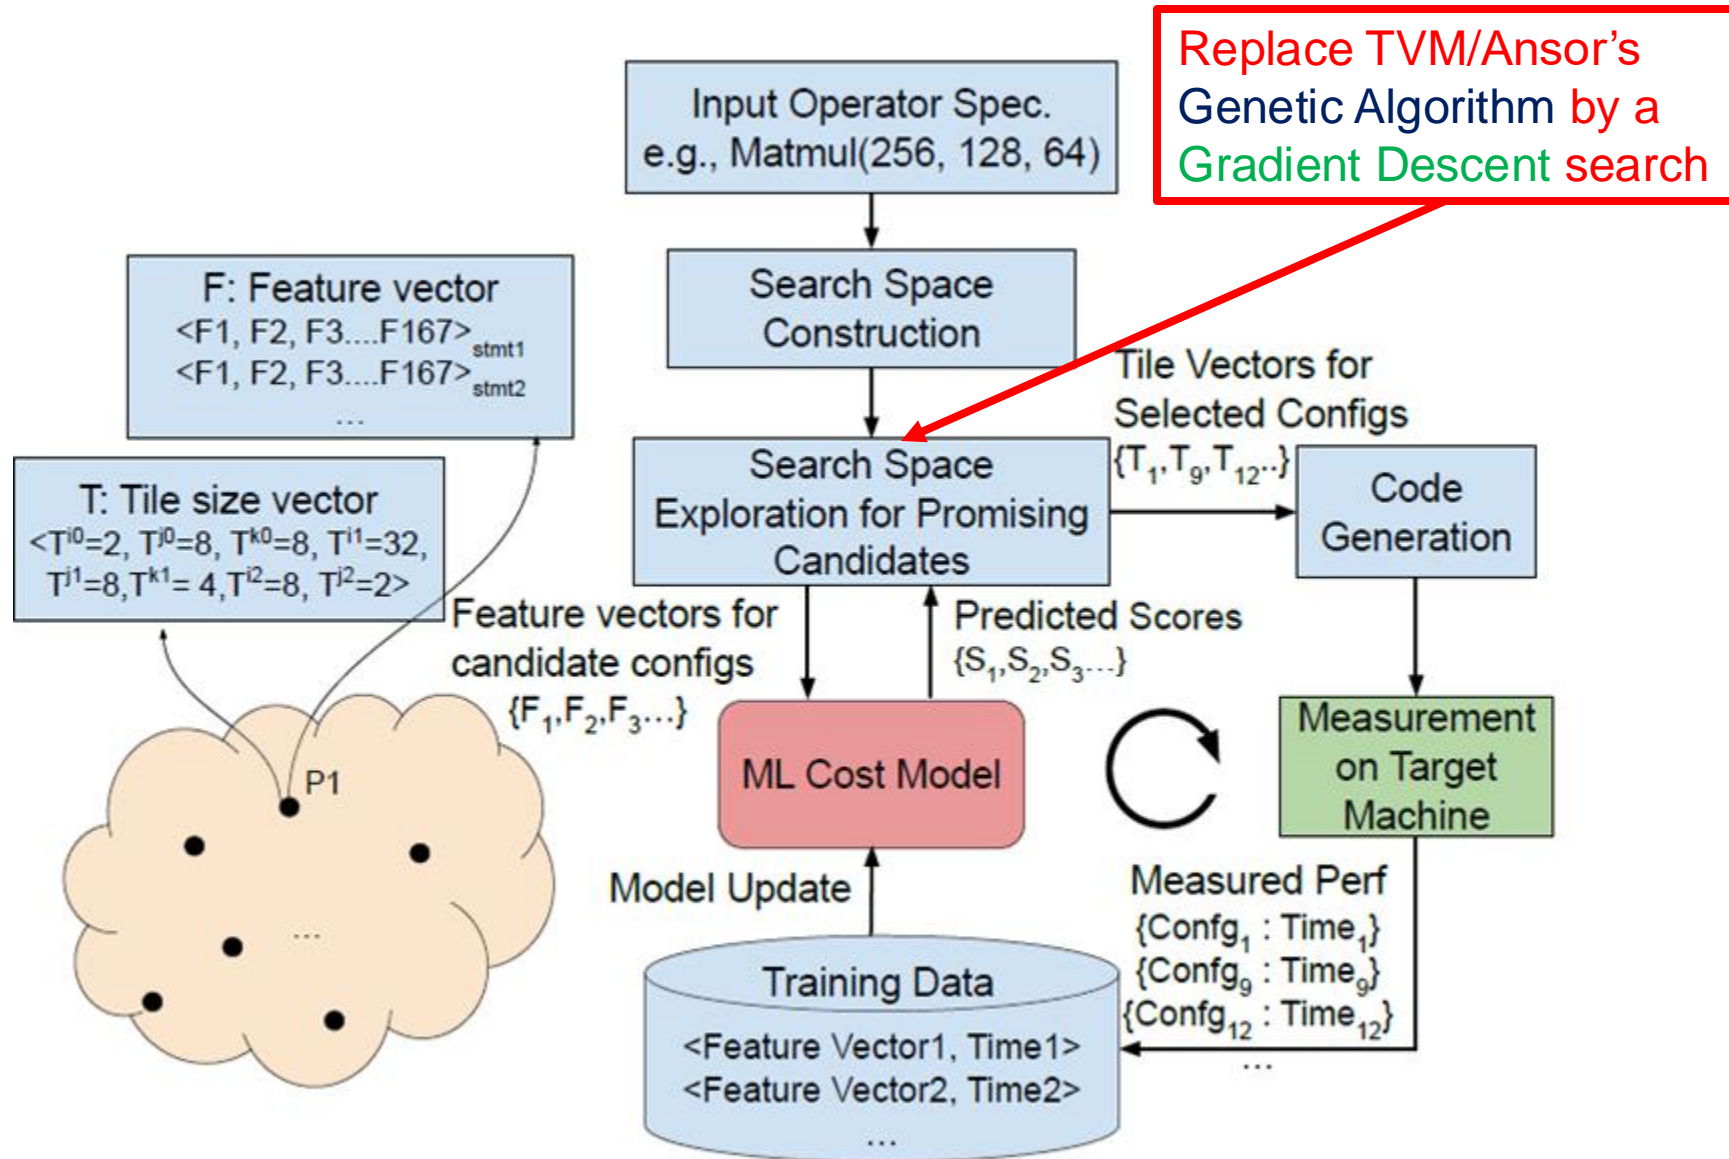
Replace default 167 features derived from AST by 7 analytically derived features



Exploration 1: Use Analytical Modeling


- ⑩ Replace Ansor's 167 metrics from AST with 7 analytically derived metrics from tile sizes
 - Data movement metrics:
 1. Operation Intensity w.r.t. global => shared memory
 2. Operation Intensity w.r.t. shared memory => registers
 3. Operational Intensity w.r.t. registers => global memory
 - Concurrency metrics:
 4. Instruction level parallelism (ILP)
 5. Warp level parallelism (WLP)
 6. Estimated occupancy
 - Load balance metric:
 7. Wave efficiency

Exploration 2: Use Gradient-Descent Search



Exploration 2: Gradient Descent Search

- Replace Genetic Algorithm search by Gradient Descent search
 - Start at a random configuration
 - Use XGB proxy model to predict performance of all 1-hop and 2-hop neighbors in D-dimensional tile-space
 - Evaluate the top-k (k=2 used) by compile/execute/measure
 - If better, move; else “slide window” and evaluate the next-k
 - Abort thresholds to terminate search path and start at new random start



Experimental Evaluation

- Compared TVM/Ansor with Ansor-AF-DS
 - AF: Exploration 1 (Analytical Features for XGBoost model)
 - DS: Exploration 2 (Gradient-Descent Dynamic Search)
- Three runs for each benchmark on two platforms
 - Nvidia RTX 3090 and RTX 4090
 - Mean and variability computed
- Best achieved performance measured after:
 - 1 minute wall-time for auto-tuning
 - 2 minutes wall-time for auto-tuning
 - 1000 total trials for auto-tuning

Experimental Evaluation: Benchmarks

Matrix Multiplication (Bert Base/Large)

Layer	M	N	K
M0	512	64	1024
M1	512	4096	1024
M2	512	64	768
M3	512	3072	768
M4	512	1024	4096
M5	512	768	3072


Layer	F	C	H/W	R/S
Y0	32	3	544	3
Y1	64	32	272	3
Y2	128	64	136	3
Y3	64	128	136	1
Y4	256	128	68	3
Y5	128	256	68	1
Y6	512	256	68	3
Y7	512	256	34	3
Y8	256	512	34	1
Y9	1024	512	17	3
Y10	512	1024	17	1

ResNet-50 conv2d

Layer	F	C	H/W	R/S
R0*	64	3	224	7
R1	64	64	56	1
R2	64	64	56	3
R3	256	64	56	1
R4*	128	256	56	1
R5	128	128	28	3
R6	512	128	28	1
R7*	256	512	28	1
R8	256	256	14	3
R9	1024	256	14	1
R10*	512	1024	14	1
R11	512	512	7	3
R12	2048	512	7	1

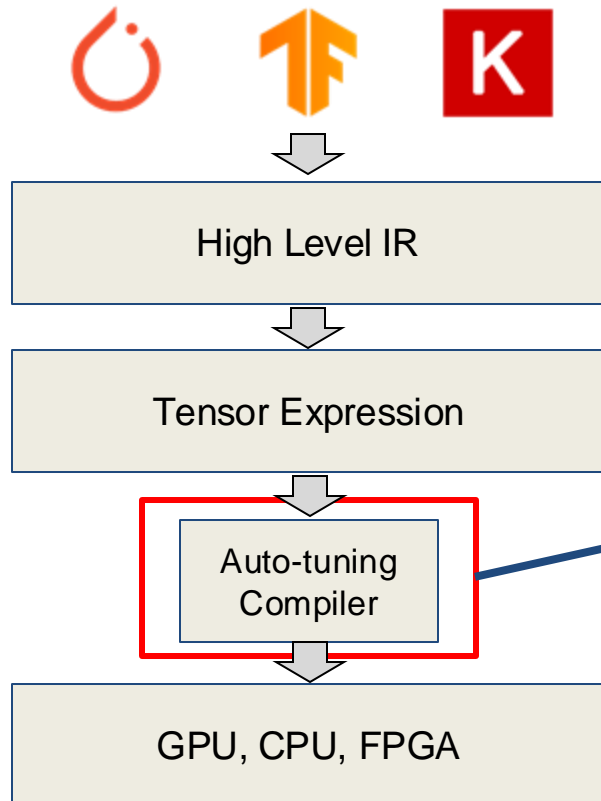
Yolo-9000 conv2d

Experimental Evaluation: Kernel Performance



- Ansor-AF-DS: Comparable (within 5%) or better perf. in 2 minutes (dark maroon) than Ansor 1000 trials (green cross)
- For more than half the benchmarks, Ansor-AF-DS after 1 minute (light maroon bar) is better than Ansor-1000-trials.

Some Open Problems



TVM can only synthesize code for fixed constant values of N,M,K.

How to extend for parametric N,M,K?

```
matmul = te.compute(
    (N, M),
    lambda i, j: te.sum(A[i, k] * B[k, j]),)
```

Generate configuration space

Search for optimal configuration

TVM does not learn and improve from previous tuning runs.

How to progressively refine code schema or search strategy across runs?

Summary

- Significant interest in ML for compiler optimization
 - But most efforts are not (yet) targeting high-impact scenarios
- TVM/Ansor's ML-based framework synthesizes parallel high-performance codes for GPUs and CPUs
 - Higher performance than vendor libraries
 - But currently limited to fixed sized tensor operators
- Can the scope/impact of ML for compiler optimization be enlarged via powerful ML models/methodologies?
 - Reinforcement Learning?
 - Bayesian Optimization?
 - LLMs?