

# KernMLOps

---

Aditya Tewari, Saurabh Agarwal, Christopher J. Rossbach

The University of Texas at Austin

# Outline

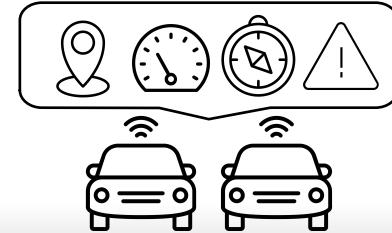
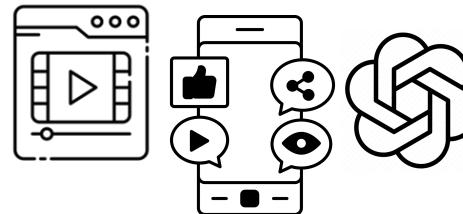
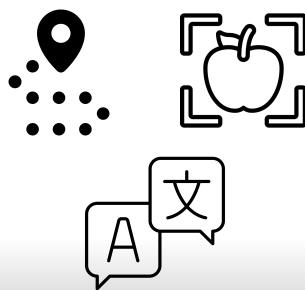
- (Re)vision: Learning-directed OSes
- KernMLOps

# Vision

To build an intelligent, self-adaptive OS  
that optimally meet modern applications' strict needs  
in dynamic and complex scenarios.

Apply Machine Learning to achieve maximal  
efficiency and enable transformative use cases that  
otherwise impose prohibitive cost and effort.

# OSes: Core Computing Infrastructure



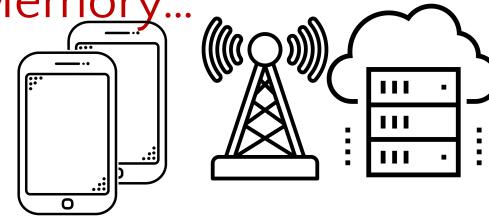
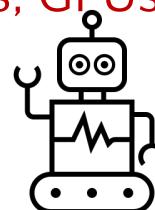
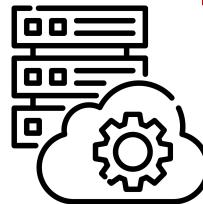
Applications

Operating System (OS)

Resource management: scheduling, allocation

Computer systems

Resources: CPUs, GPUs, Network, Memory...



# OSes: Core Computing Infrastructure



*Real-time*



*High throughput*



*Diverse resource  
needs*

**Increasingly demanding**



**Operating System (OS) / Runtimes**

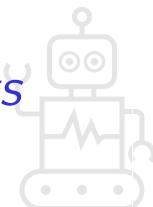
**Effective resource management is crucial!**



**Complex, dynamic conditions**



*Complex environments*



*Constant change*

# OS Resource Management is Broken

Cannot support modern concurrent apps in complex, dynamic settings

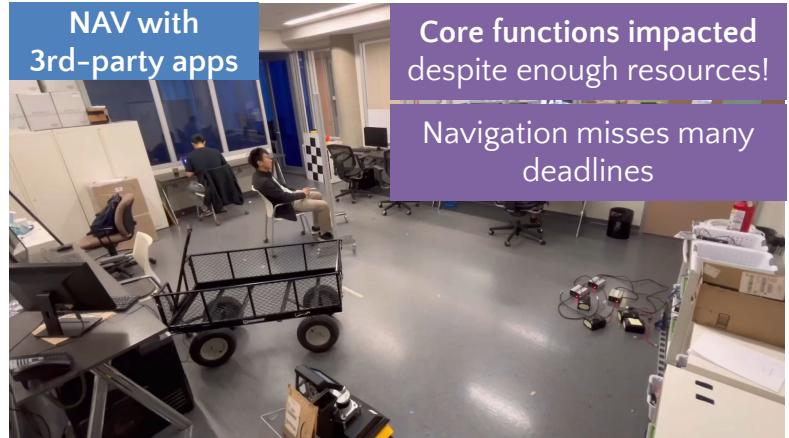
Consumer  
autonomous  
robots



Standalone  
navigation  
(NAV)

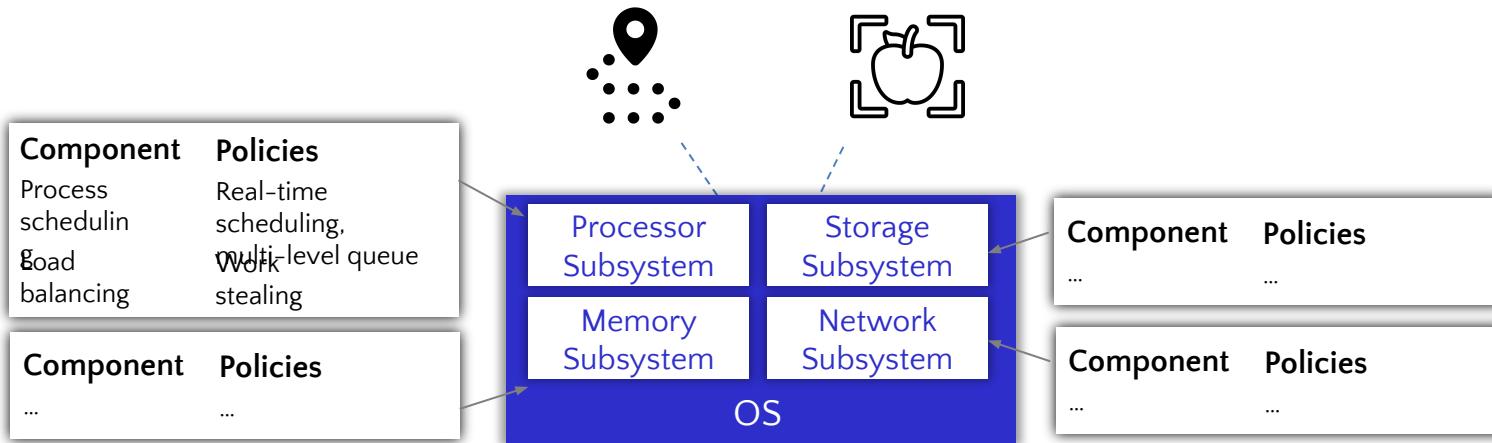


NAV with  
3rd-party apps



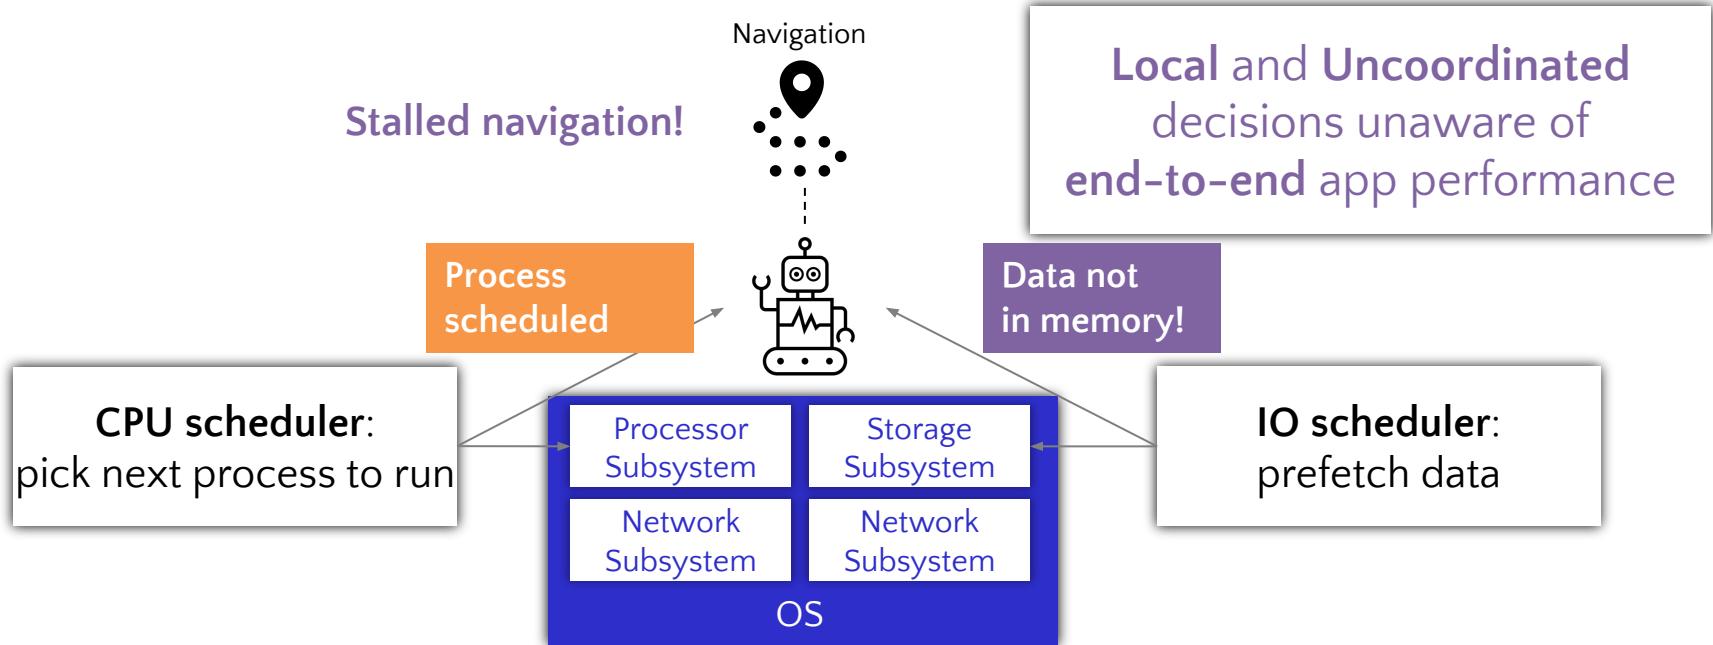
# Why is Resource Management Broken?

## Resource management **policies**



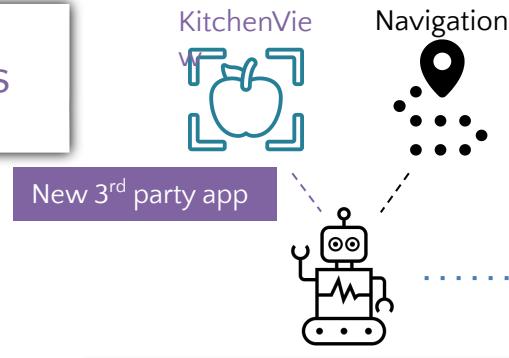
Policy architecture has two fundamental drawbacks

# Poor Composability of Decisions



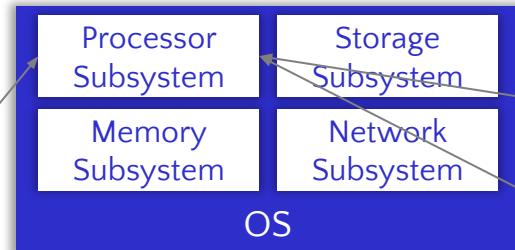
# Limited and Rigid Decision-Making

Cannot adapt to changes



Poor decisions amid complexity

Heuristics with fixed logic



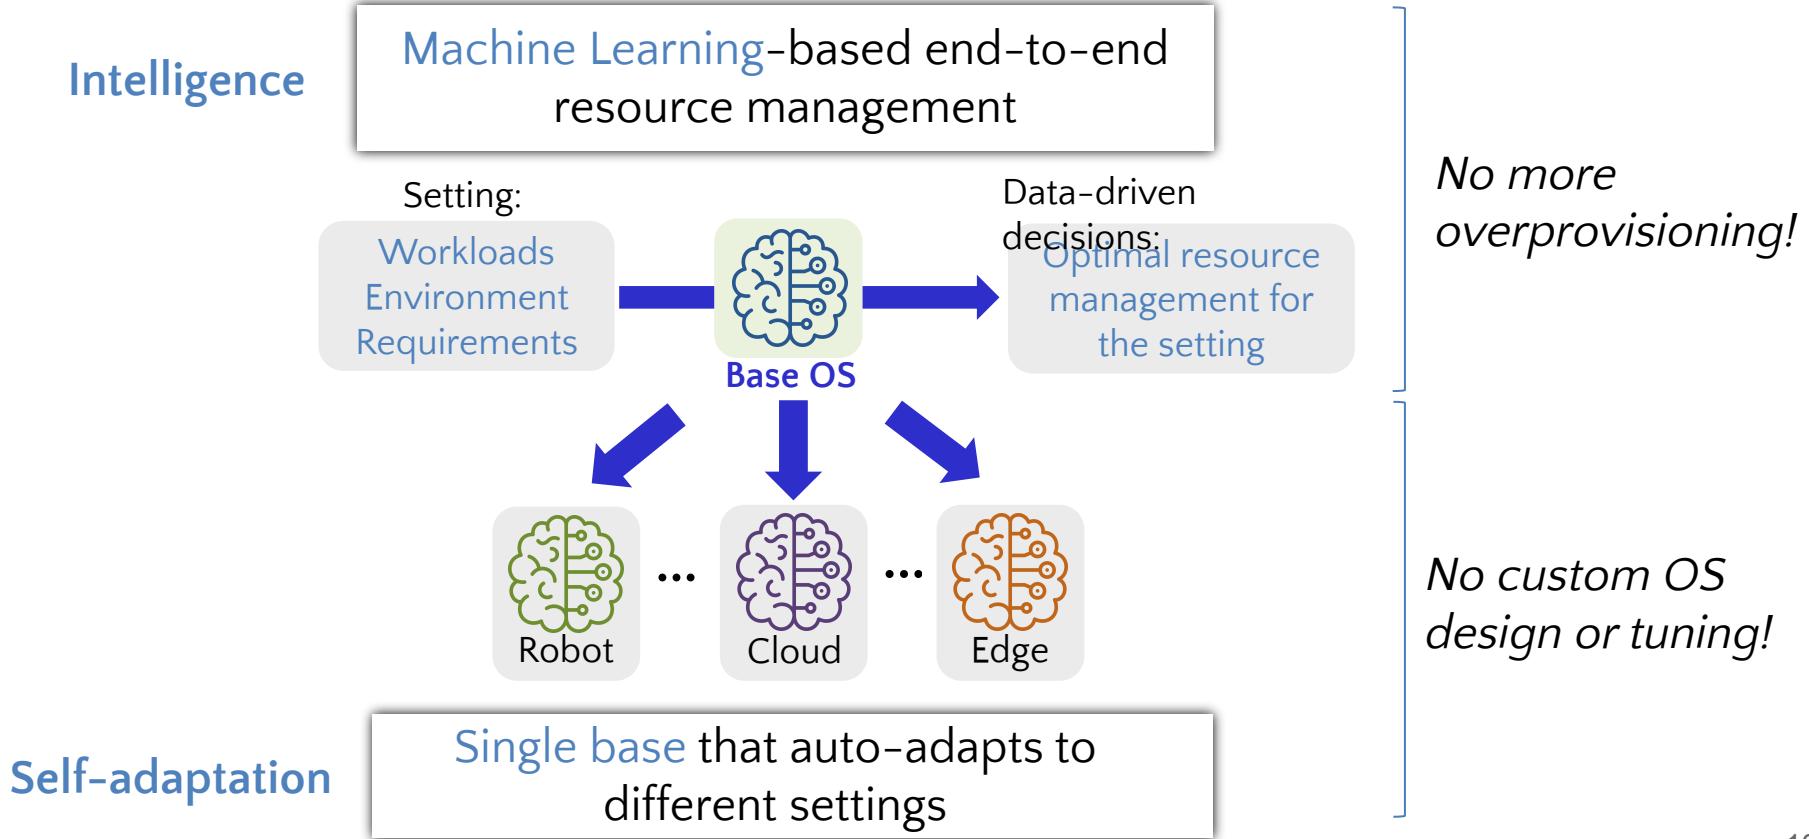
Multi-level feedback queue

Component-specific inputs

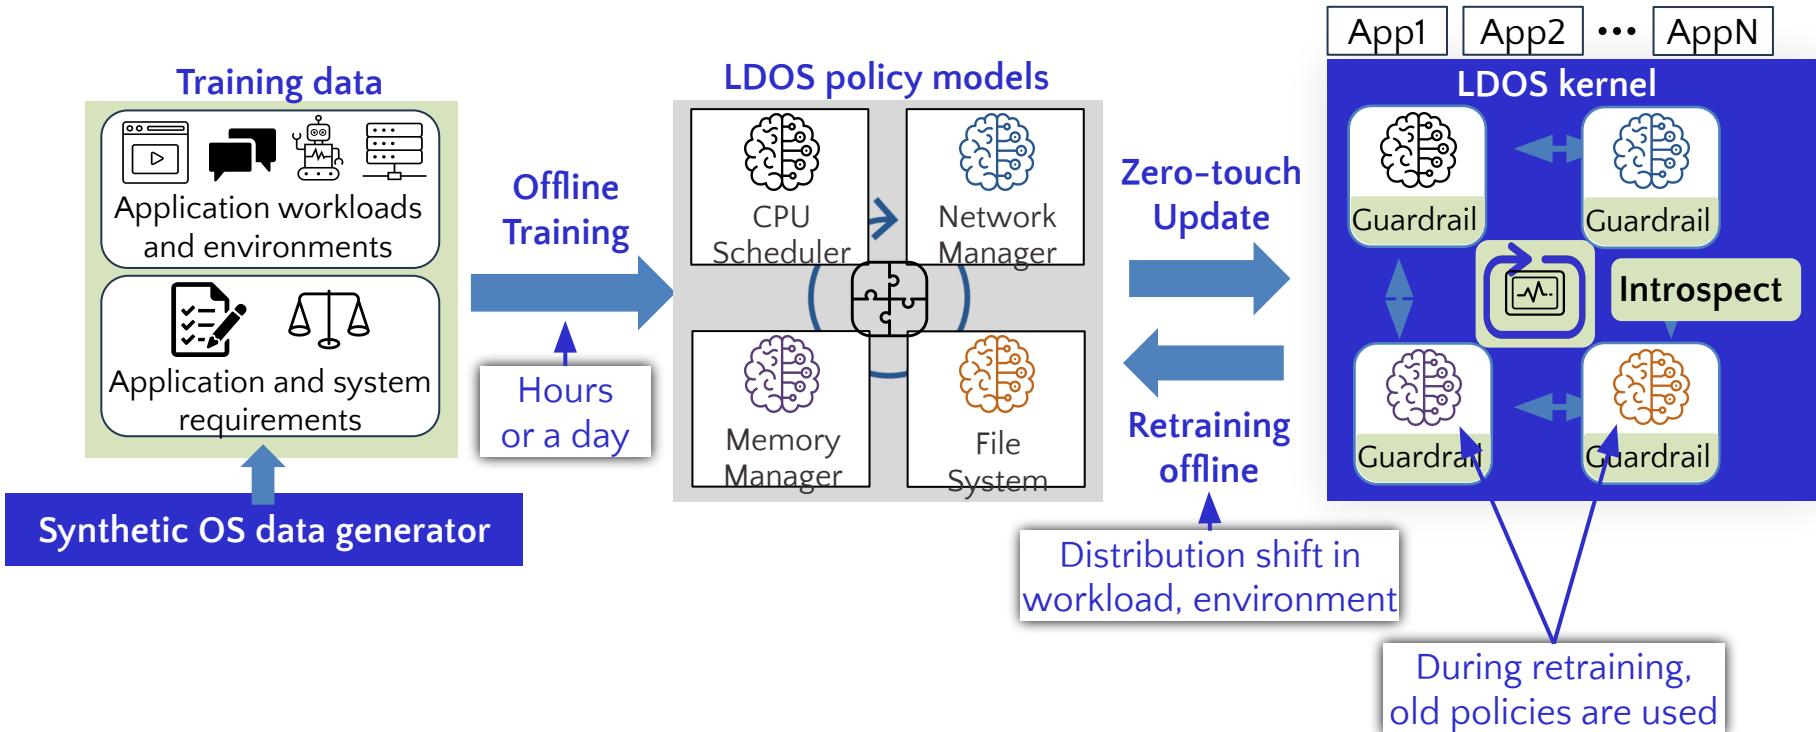
Input: process CPU usage, ...

Decisions for next time step

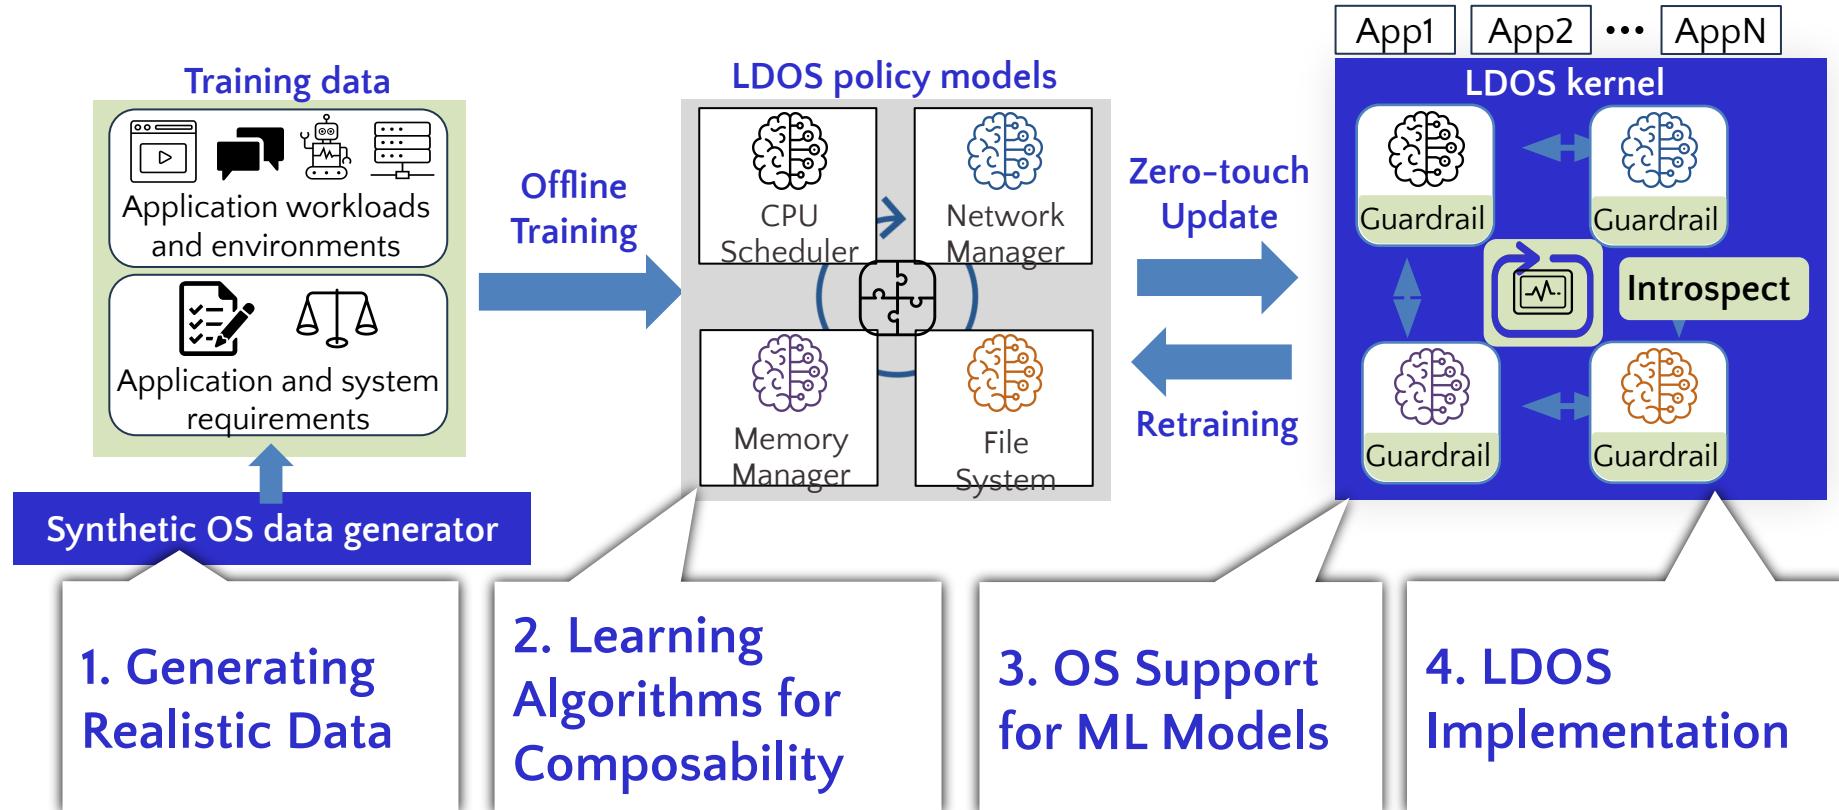
# Intelligent, Self-Adaptive OSes and Runtimes



# LDOS Approach



# Research Thrusts



# Addressing Challenges with LDOS

- Robustness   ○   **Verified Learning**
- Data is hard to come by   ○   **Foundation Model for OS**
- Explainability/Misbehavior   ○   **Guardrails / Adaption**
- Development cycle friction**   ○   **KernMLOps**      Today  

- OS structure   ○   **Clean-slate Design**      
- Choosing learning techniques   ○   <>
- OS+ML: *different* acceleration   ○   <>

# Implementing ML-backed Policies

*Where to start? Many Kernel Subsystems!*

## Core

- Drivers
- Memory Management
- Power Management
- **Scheduler**
- Timers
- Locking

## Networking

- **Protocols**
- NetLabel
- Infiniband/RDMA
- ISDN
- MHI

## Other

- Accounting
- Accelerators
  - FPGAs, TPUs
- Watchdog
- Virtualization
- Crypto
- BPF
- TEEs

## HID

- HID Devices
- Sound
- GPU
- Frame Buffer
- LEDs
- Locking

## Storage

- **Filesystems**
- **Block**
- CD-ROM
- SCSI
- TCM Virtual Device
- VFS Layer

*We've worked with a few subsystems, each of which has taken months-years of effort!*

# Why does ML-ifying the OS take so long?

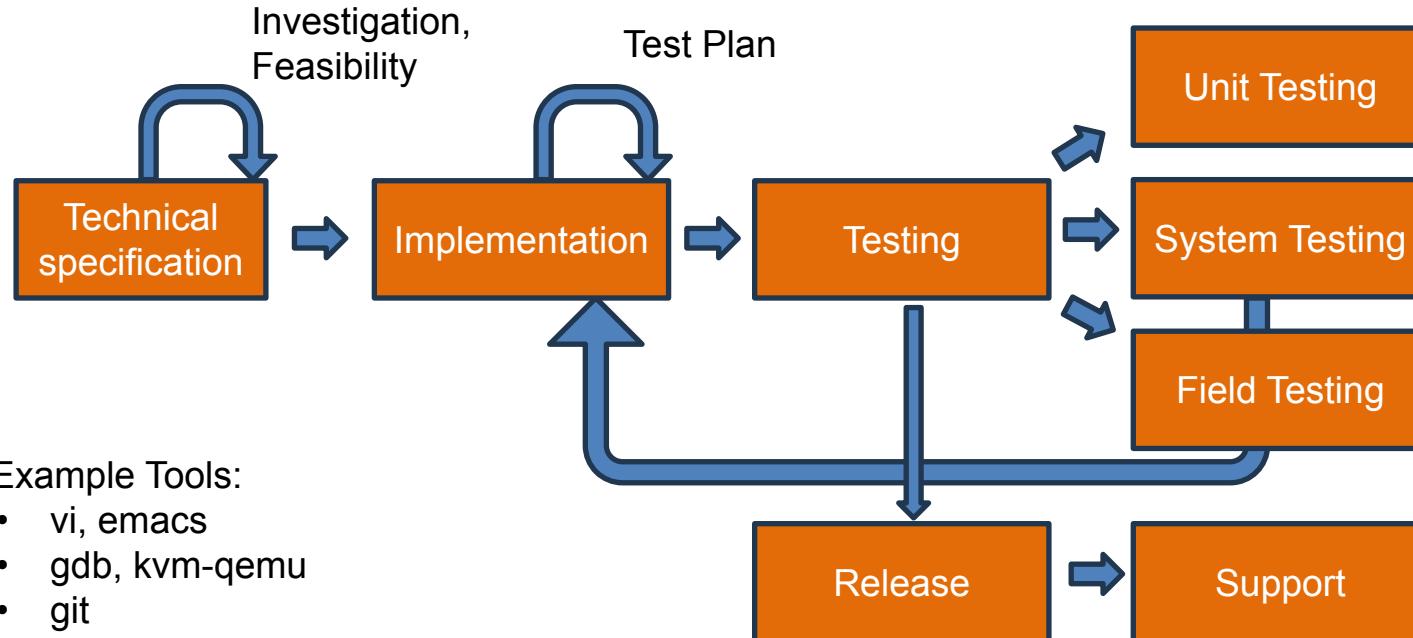
Common Scenario: researcher X investigates using ML in OS subsystem Y

Issues:

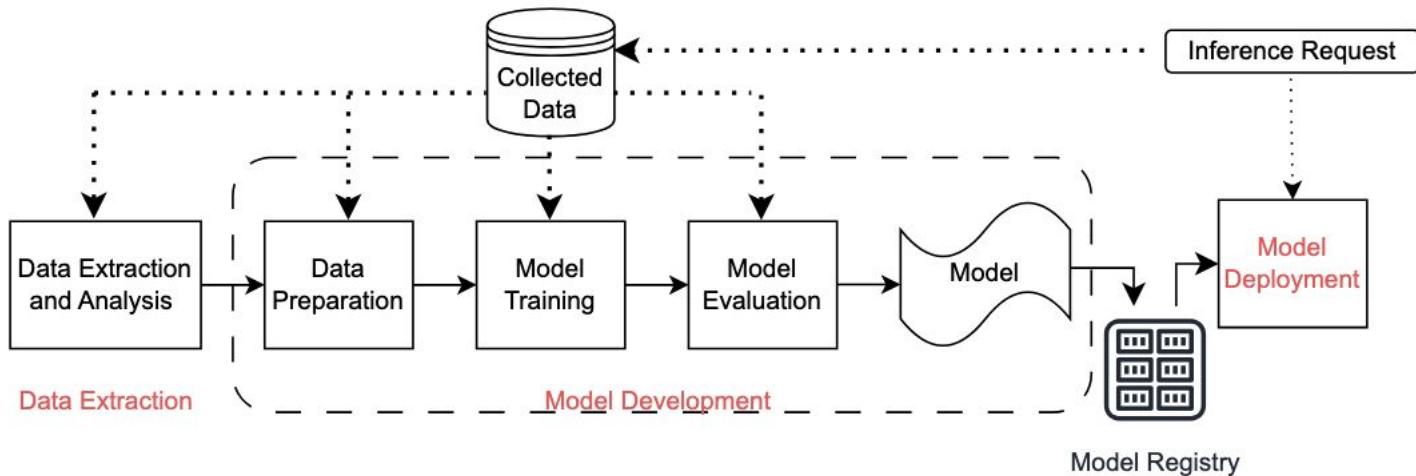
- Researchers seldom have ***both*** ML and OS background
  - Often crash out due to slow progress with one or both
- In general, we do not know if subsystem Y is ML-amenable
  - How to collect data?
  - What ML technique to use?
  - Is my classifier working as expected?
- How does X develop train the model and deploy it in the OS?

***OS developer workflow != MLOps workflow***

# OS Developer Work Flow



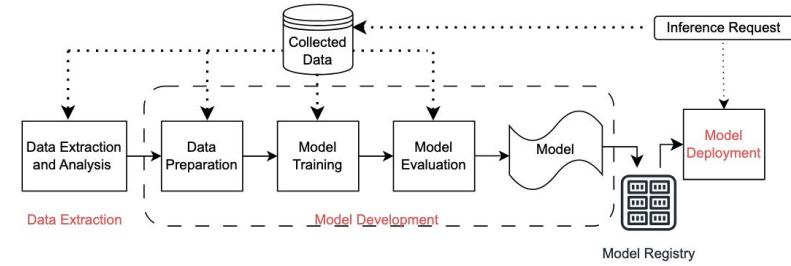
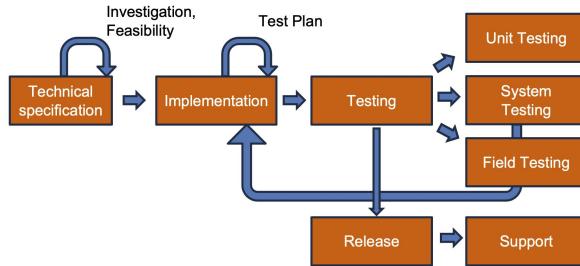
# MLOps Work Flow



Tools:

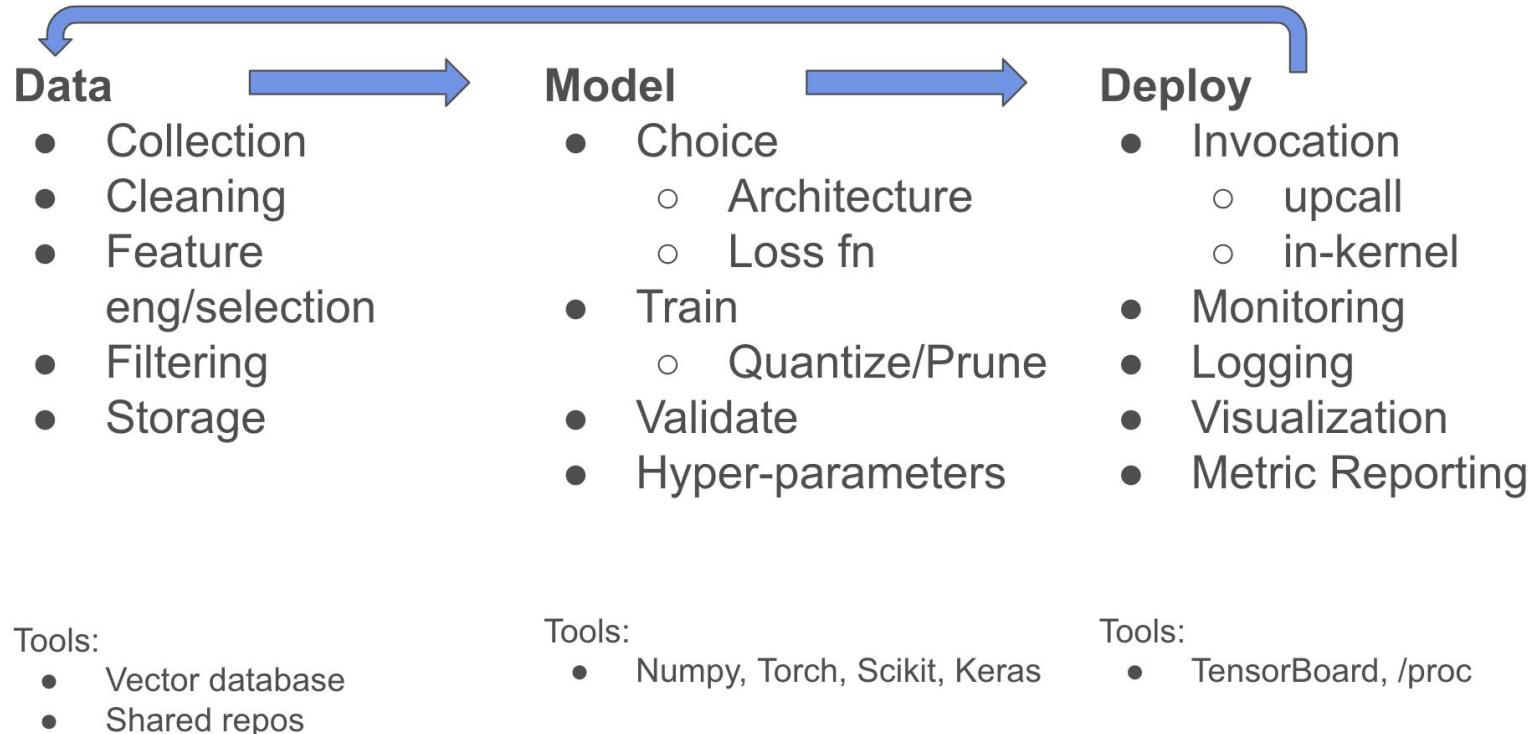
- jupyter
- pyTorch
- pandas/dataframes
- sci-kit learn
- vector database

# OS Flow != MLOps Flow

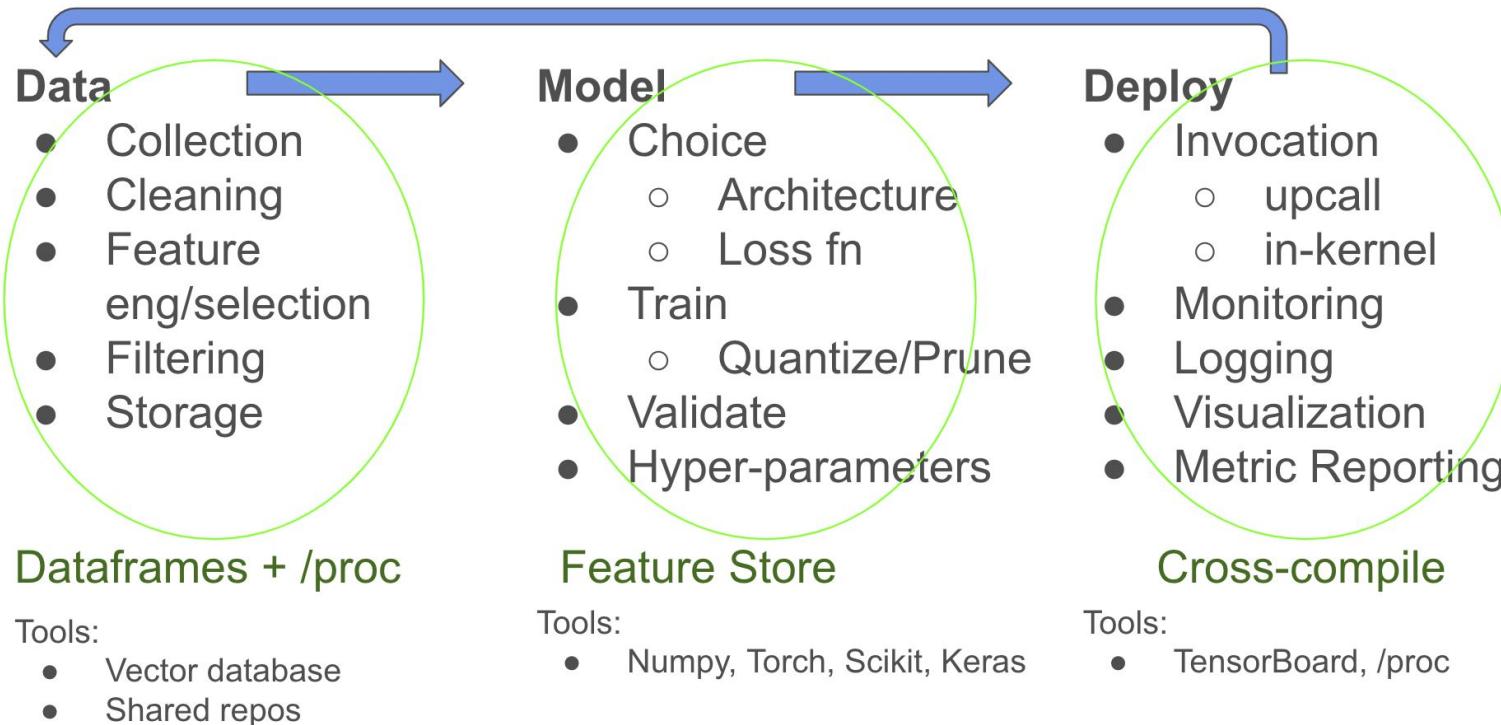


- Different tools
- Different skillsets
- Different backgrounds

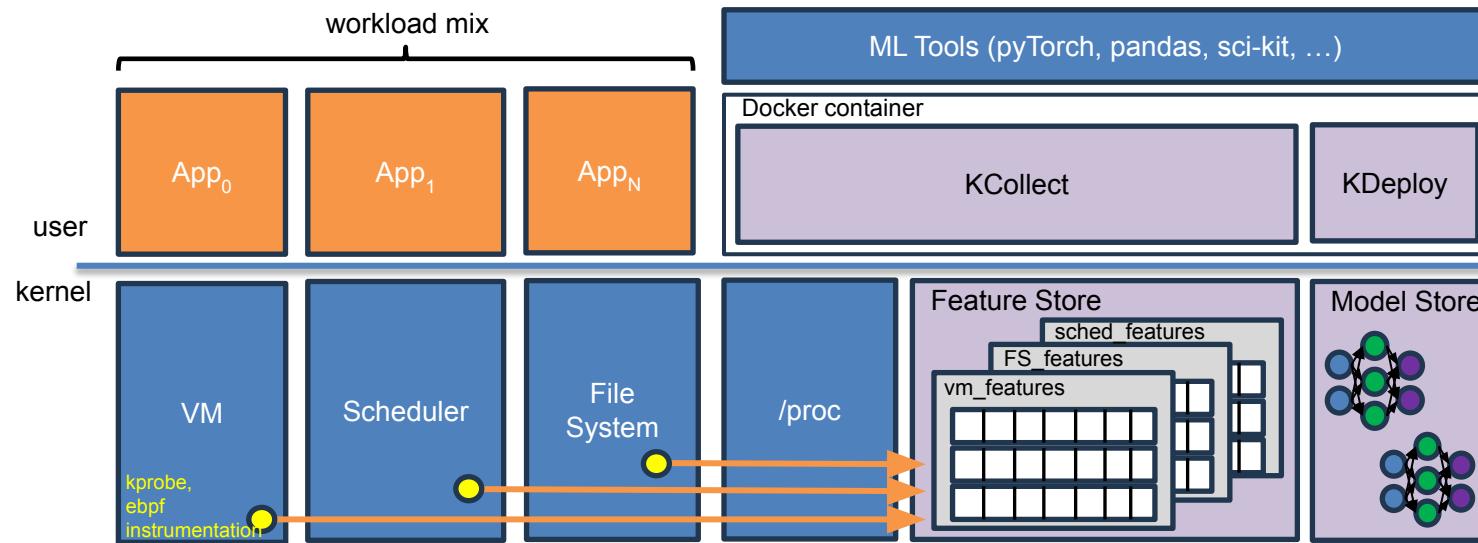
# KernMLOps: MLOps support for LDOS kernels



# KernMLOps: MLOps support for LDOS kernels



# KernMLOps Overview



# KernMLOps Overview

## • KCollect

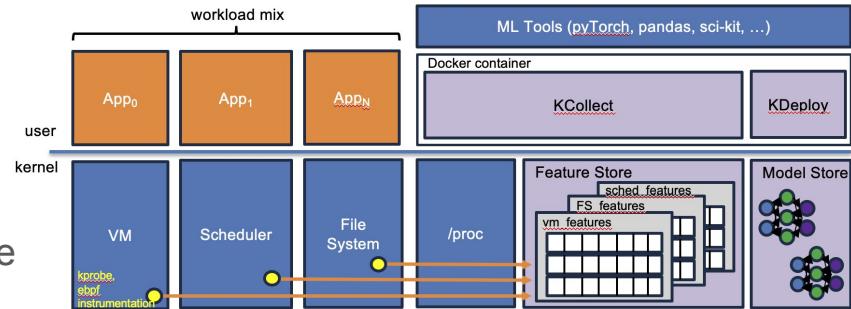
- Exposes Dataframes API
- Compatible with tools of choice
- pandas, pytorch, etc.
- Poll the /proc file system periodically

## • KDeploy

- Deploy ML models  $\square$  kernel
- User trains/validates model in pyTorch
- Deploy tools cross compile to kernel module
- Kernel module deployed automatically

## • Feature Store

- Reimplement LAKE FS using eBPF maps
- Make it easy to instrument kernel and collect data



# Feature Store API

API	Spaces Available	Description
register(name, ebpf_map_fd)	User	Register an eBPF map with the feature store
unregister(name)	User	Unregister an eBPF map with the feature store
capture_feature(ebpf_map_fd, data)	eBPF	Capture a timestamp and a struct whose type is defined with the map in register()
truncate_features(name, timestamp)	eBPF, User, Kernel	Removes all features older than timestamp
get_feature(name, key)	eBPF, User, Kernel	Gets a single feature by key
get_feature_ts(name, ts)	eBPF, User, Kernel	Gets a single feature closest to timestamp
get_features_ts(name, ts)	eBPF, User, Kernel	Gets all features older than a specific timestamp
get_all_features(name)	eBPF, User, Kernel	Gets all features within the map.

# KDeploy API

API	Spaces Available	Description
register(name, ebpf_map_fd)	User	Register an eBPF map with the feature store
unregister(name)	User	Unregister an eBPF map with the feature store
capture_feature(ebpf_map_fd, data)	eBPF	Capture a timestamp and a struct whose type is defined with the map in register()
truncate_features(name, timestamp)	eBPF, User, Kernel	Removes all features older than timestamp
get_feature(name, key)	eBPF, User, Kernel	Gets a single feature by key
get_feature_ts(name, ts)	eBPF, User, Kernel	Gets a single feature closest to timestamp
get_features_ts(name, ts)	eBPF, User, Kernel	Gets all features older than a specific timestamp
get_all_features(name)	eBPF, User, Kernel	Gets all features within the map.

# User- vs Kernel-space Deployment

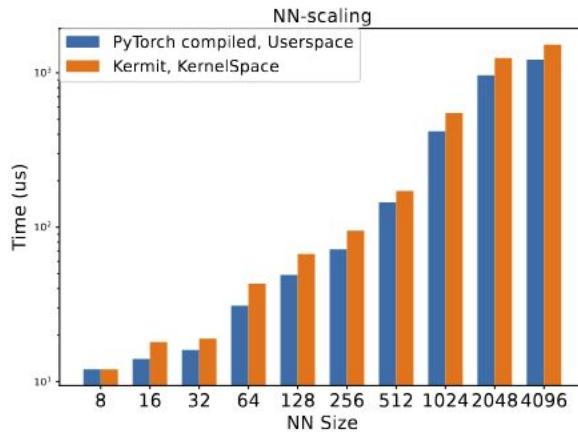


Figure 6: **Evaluating Userspace vs KernelSpace Deployment** We observe that in certain cases it might be more prudent to pay the cost of data movement and run the model in the userspace due to better library support.

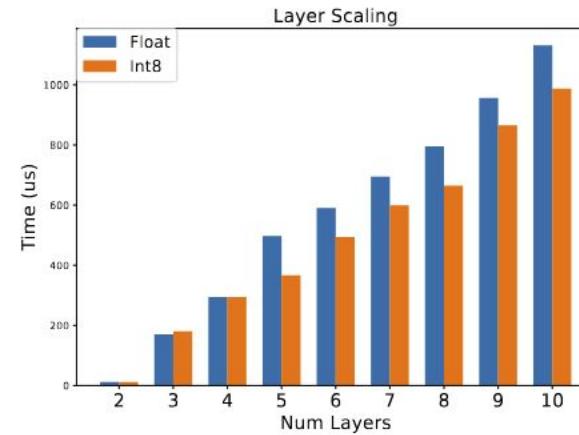


Figure 7: **Evaluating benefit of Quantization:** In the above figure we observe that static quantization for small number of layers and models leads only modest improvements.

# Case Study: Imitation learning CBMM

Work Item	Time	Kernel LoC	eBPF LoC	Python LoC
Expose functions to eBPF	5 minutes	4	0	0
Recompile Kernel	10 minutes	0	0	0
Boot kernel	10 minutes	0	0	0
Examine Eager Paging Code	20 minutes	0	0	0
Deploy Test bpftrace probes	10 minutes	0	0	0
Create eBPF maps for capturing data	45 minutes	0	50	0
Ingest and Ex-filtrate data eBPF	15 minutes	0	126	0
Ingest and Ex-filtrate data Kermit Plumbing	1 hour	0	0	294
Test Ex-filtration in Kermit	5 minutes	0	0	0
Test All CBMM Benchmarks	4 hours	0	0	0
Implement Eager Data Capture	10 minutes	0	19	0
Test Eager Data Capture	1 hour	0	0	0
Implement Prezeroing Data Capture	45 minutes	0	54	0
Test Prezeroing Data Capture	1 hour	0	0	0
Capture All data with CBMM Benchmarks	4 hours	0	0	0
Create Model	1 hour	0	0	109
Train Model	3 hours	0	0	0

# Demonstration

# Conclusion

- Good tooling essential for ML-in-kernel
- KernMLOps replicates pyTorch environment and MLOps flow on Linux
- Wanted: users!

***Questions?***