

Reading Group Fall 2024



# Perspectives on Learning-Directed Operating Systems

---

Christopher J. Rossbach

The University of Texas at Austin

# Outline

- Vision: Learning-directed OSes
- LAKE: early experience supporting ML in the OS
- Some practical challenges
- Addressing those challenges with LDOS

# LDOS: A Clean-Slate Paradigm for OS Design and Implementation

---

Aditya Akella, Alex Dimakis, Swarat Chaudhuri, *Chris Rossbach*,  
Sebastian Angel, Joydeep Biswas, Shuchi Chawla, Isil Dillig,  
Brighten Godfrey, Daehyeok Kim, Sanjay Shakkottai, Michael  
Swift, Shivaram Venkataraman, and Gang Wang



UNIVERSITY OF  
**ILLINOIS**  
URBANA-CHAMPAIGN

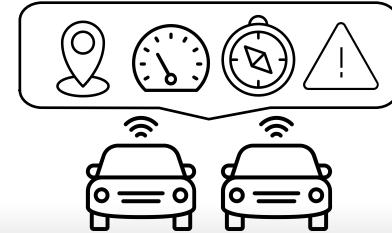
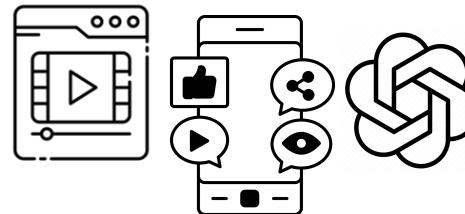
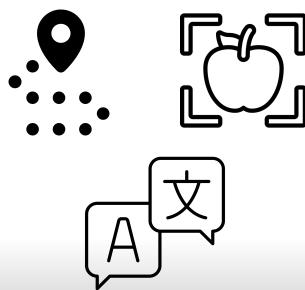


# Vision

To build an intelligent, self-adaptive OS  
that optimally meet modern applications' strict needs  
in dynamic and complex scenarios.

Apply Machine Learning to achieve maximal  
efficiency and enable transformative use cases that  
otherwise impose prohibitive cost and effort.

# OSes: Core Computing Infrastructure



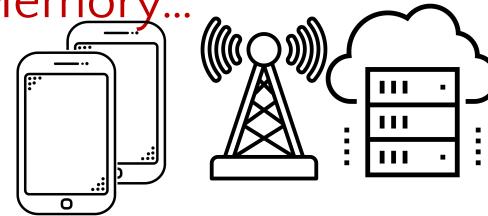
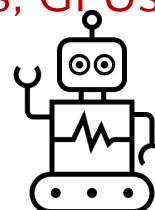
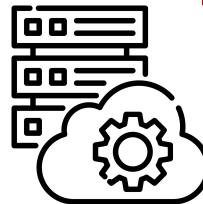
Applications

Operating System (OS)

Resource management: scheduling, allocation

Computer systems

Resources: CPUs, GPUs, Network, Memory...



# OSes: Core Computing Infrastructure



*Real-time*

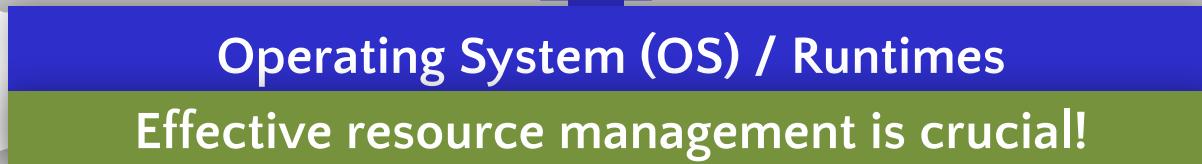


*High throughput*



*Diverse resource  
needs*

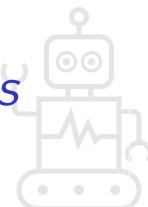
**Increasingly demanding**



**Complex, dynamic conditions**



*Complex environments*



*Constant change*

# OS Resource Management is Broken

Cannot support modern concurrent apps in complex, dynamic settings

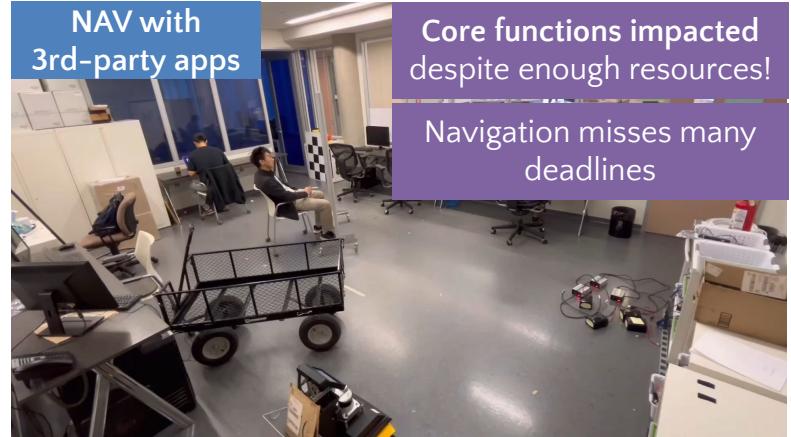
Consumer  
autonomous  
robots



Standalone  
navigation  
(NAV)



NAV with  
3rd-party apps



# OS Resource Management is Broken

Cannot support modern concurrent apps in complex, dynamic settings

Consumer  
autonomous  
robots



3<sup>rd</sup> party added-function apps can't run alongside core apps

Cloud  
computing



Poor dependability

Mobile  
access  
edge



Multiplexed real-time apps suffer

# Overprovisioning and Custom Design



**Extra CPUs, GPUs, memory**

Wasteful, prohibitively expensive, infeasible



**Cloud utilization - 10-20%**

Alibaba SoCC'21; Azure SOSP'17

Untenable ecological impact

Cloud carbon footprint > the aviation industry



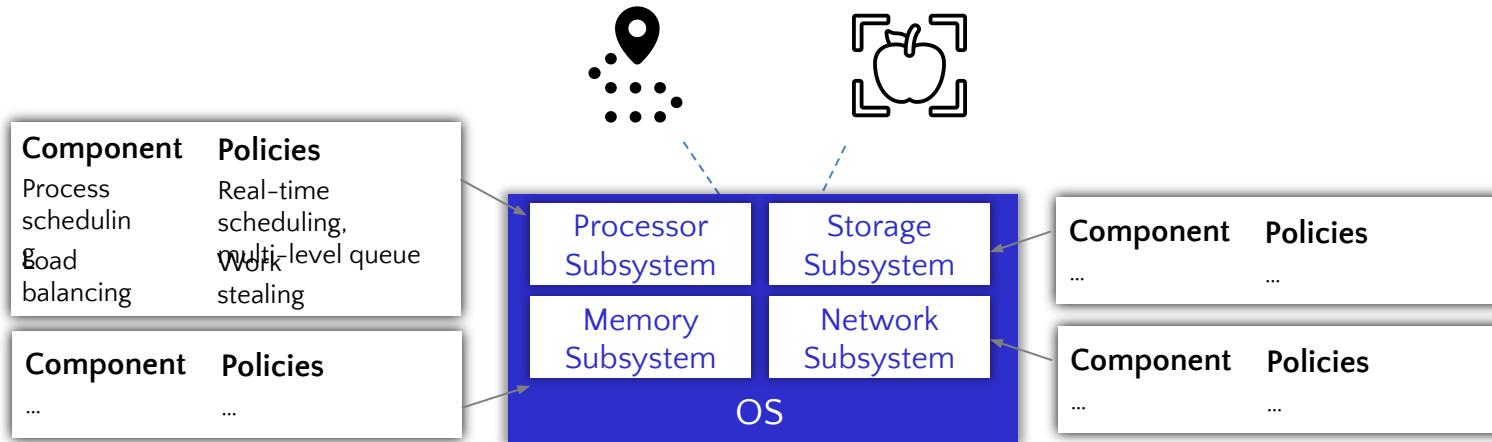
**Near-real-time kernel**

Long time to develop

Easily become obsolete

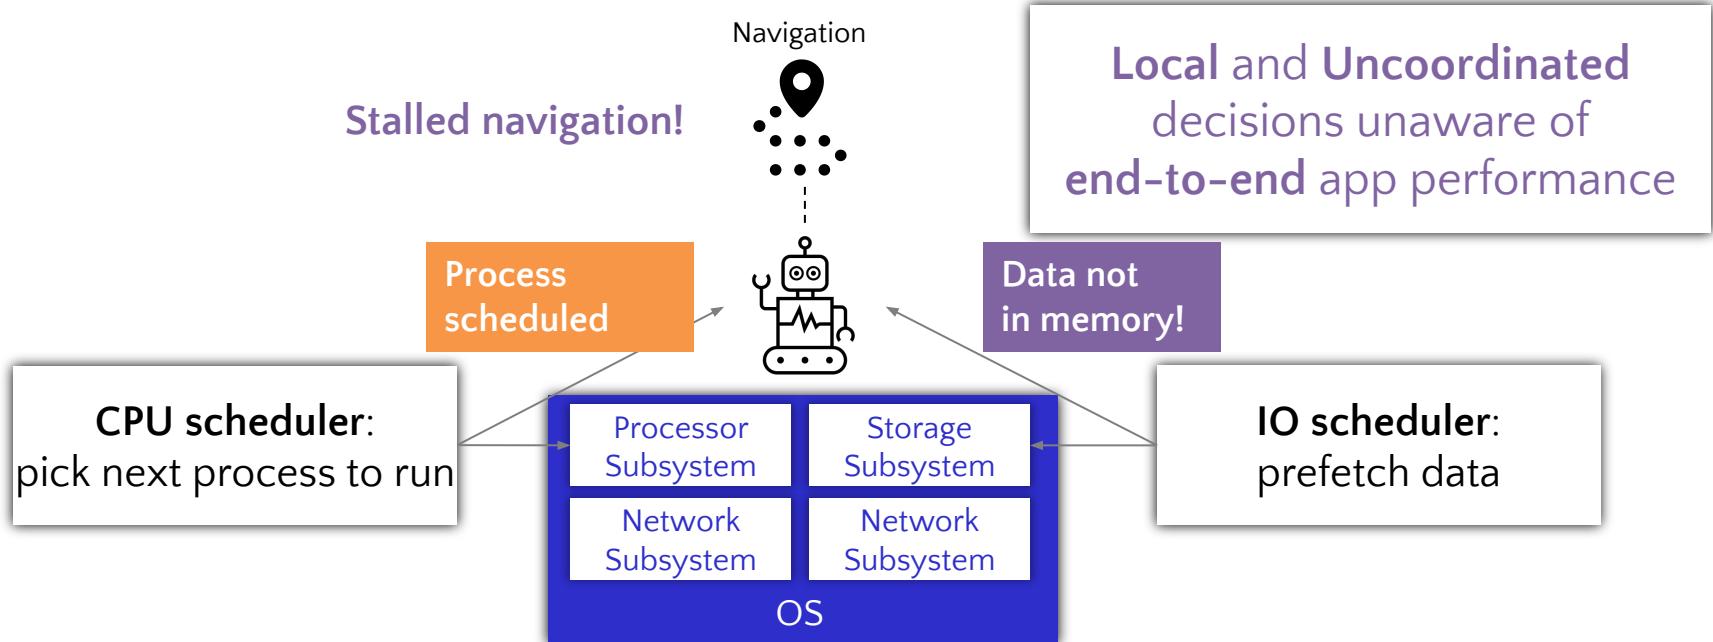
# Why is Resource Management Broken?

## Resource management **policies**



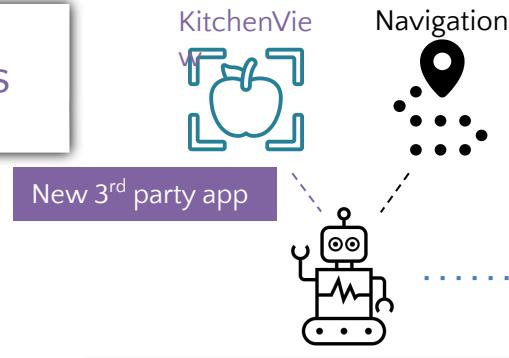
Policy architecture has two fundamental drawbacks

# Poor Composability of Decisions

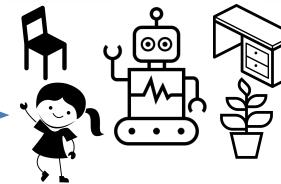


# Limited and Rigid Decision-Making

Cannot adapt to changes

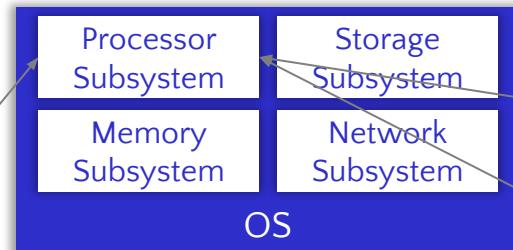


Poor decisions amid complexity



Heuristics with fixed logic

Multi-level feedback queue



Component-specific inputs

Input: process CPU usage, ...

Decisions for next time step

# An Impending Crisis

Optimal decisions exist per setting, but

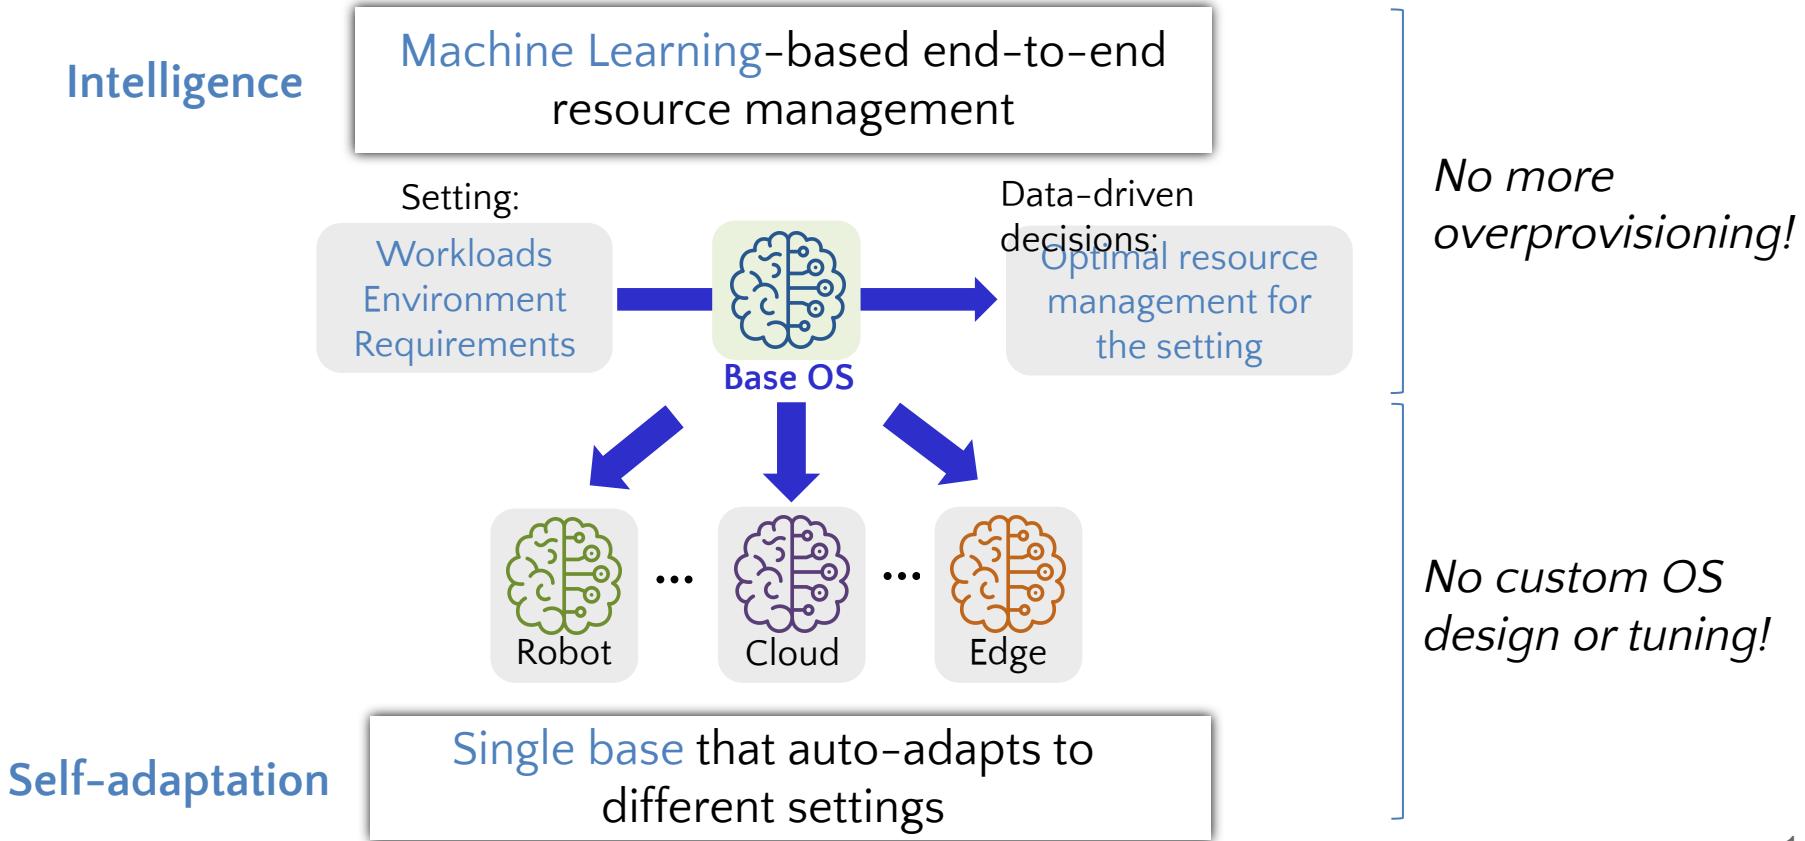
- . **Infeasible** for humans to find
- . OSes **cannot express** them due to policy drawbacks

The best today's OSes can do vs. optimal: **gap can be significant**

An unmanageable crisis is brewing

Performance and efficiency needlessly sacrificed to ever-greater degrees despite substantial human engineering effort

# Intelligent, Self-Adaptive OSes and Runtimes



# Learning-Directed Operating System (LDOS)

**Clean-slate** OSes / runtimes with a new ML-driven resource management policy architecture

**Cross-cutting** approach that fundamentally advances generative AI, ML, formal methods, and systems

## Gen AI, ML, and formal methods innovations

- ML policies driven by rich run-time data, and trained using diverse synthetic data
- Policies that compose to guarantee strict end-to-end app and system needs

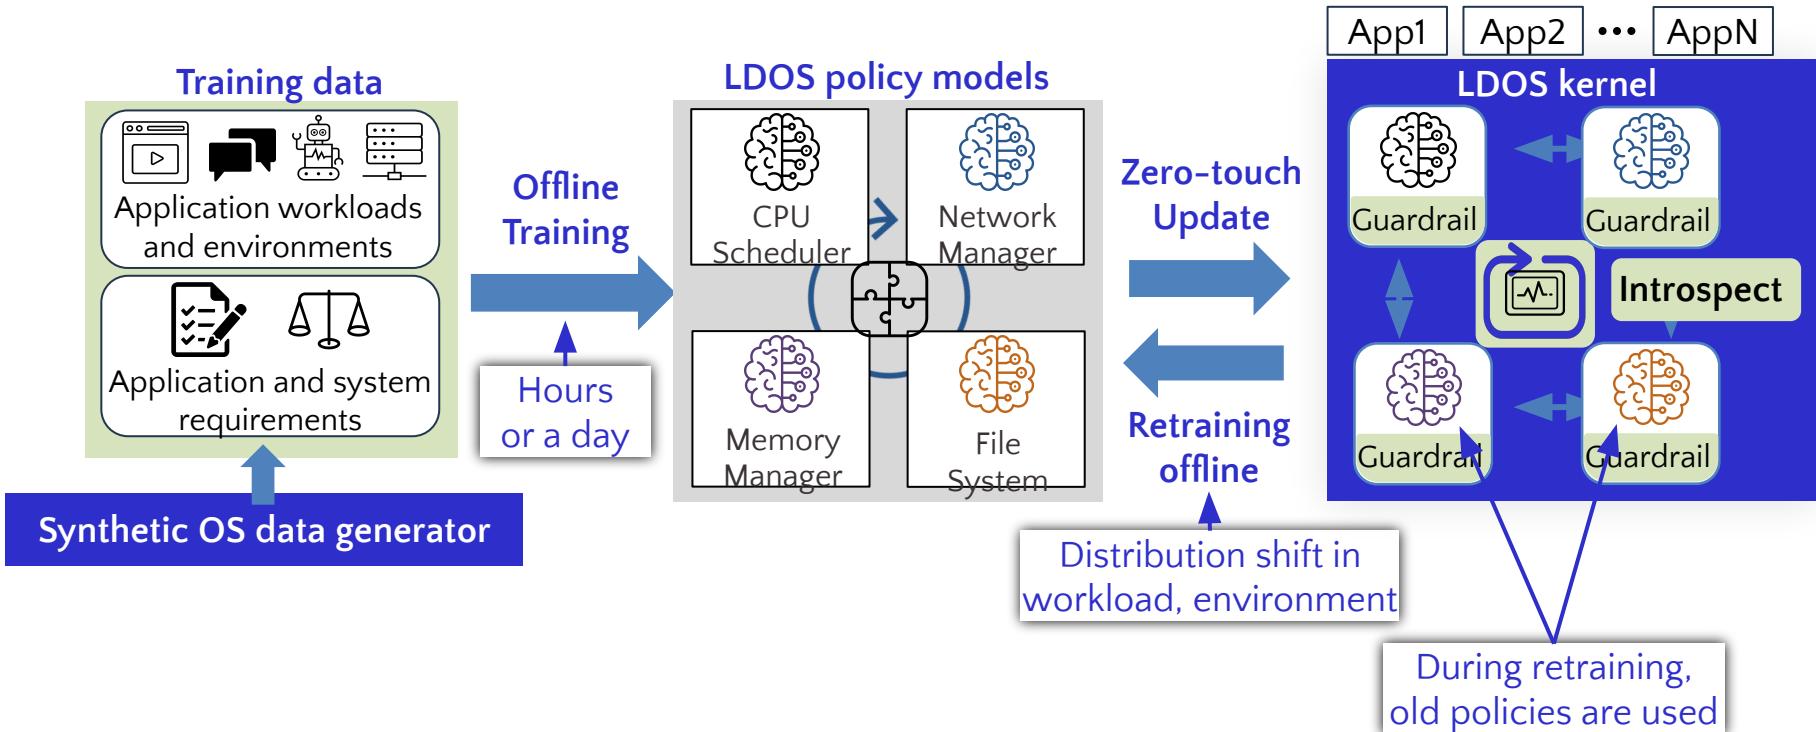
## ML-centric OS design

- Easy policy integration, automatic adaptation
- Low overhead, security, manageability

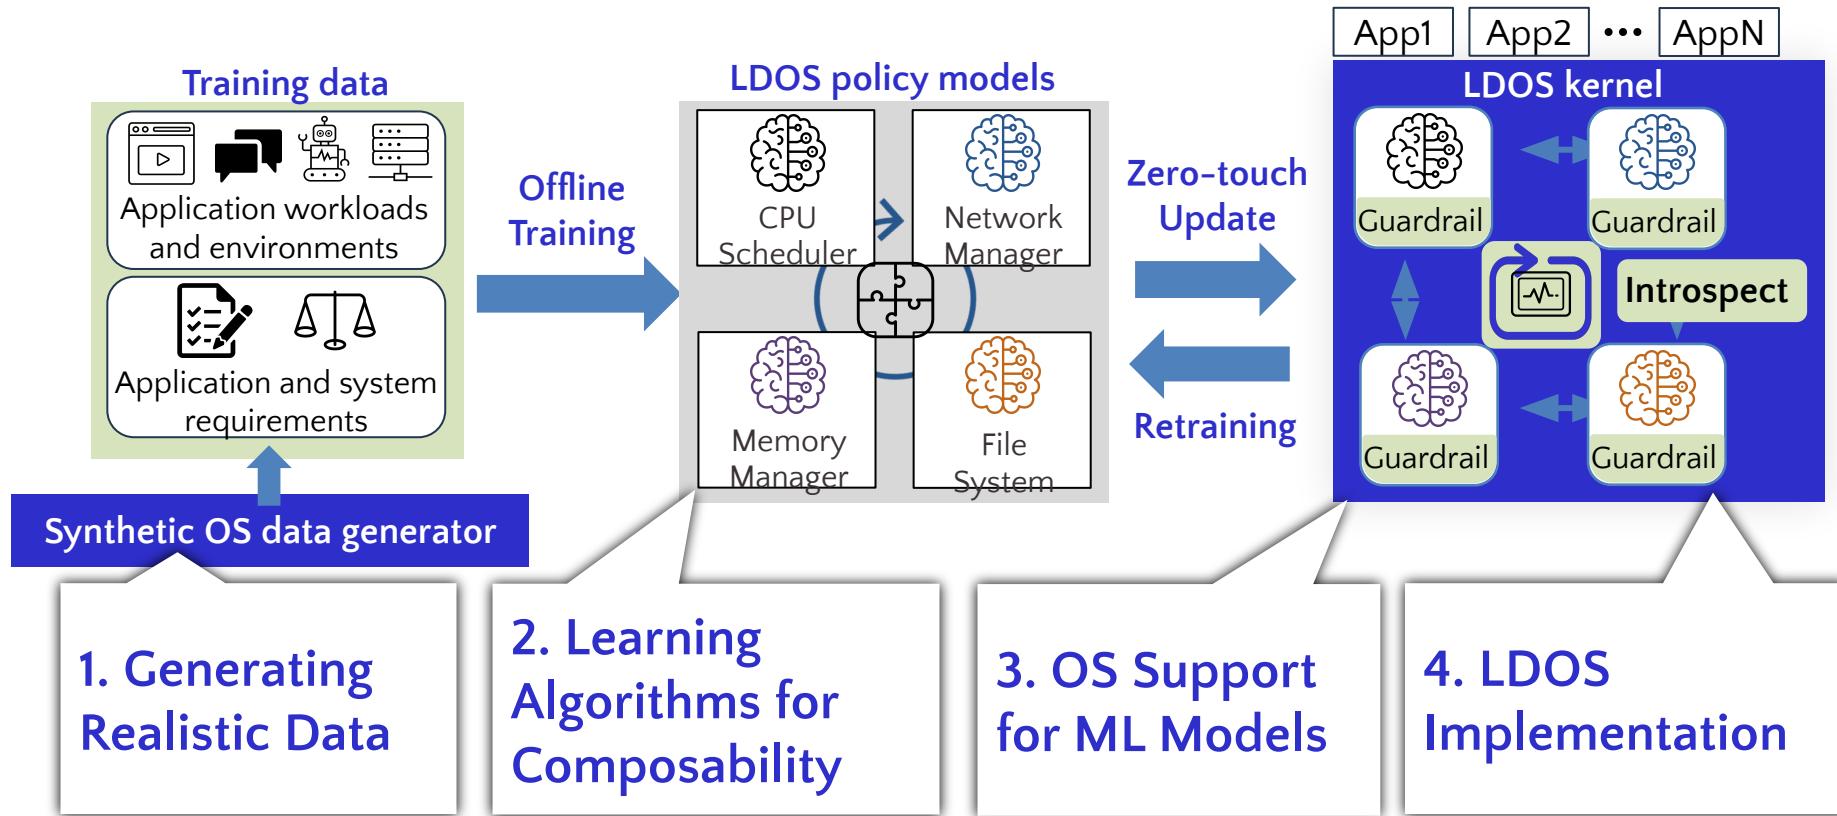
## Pragmatic phased approach

- Substantial benefits per phase

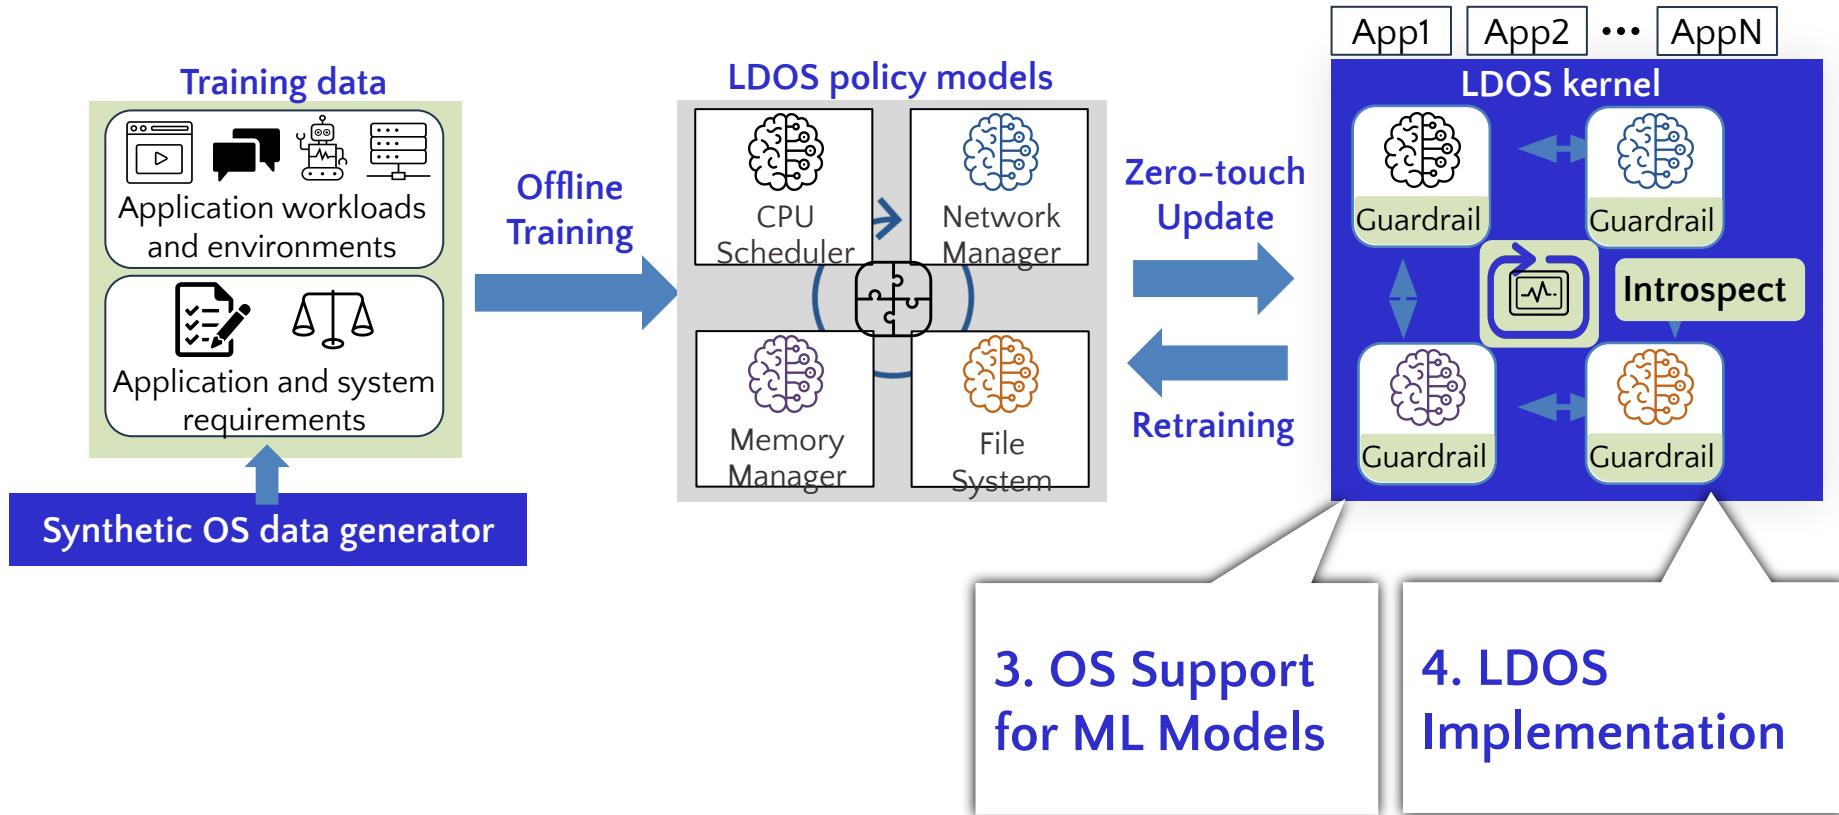
# LDOS Approach



# Research Thrusts



# LAKE: Early Experience



ASPLOS 2023

# Towards a Machine Learning-Assisted Kernel with LAKE



---

Henrique Fingler\*, Isha Tarte \*, Hangchen Yu †, Ariel Szekely ‡,  
Bodun Hu \*, Aditya Akella \*, Christopher J. Rossbach \*§

\*The University of Texas at Austin    †Meta    ‡ Massachusetts Institute of Technology    § Katana Graph

# ML to Replace OS Heuristics

- **KML:** (HotStorage '21) – storage management
- **SmartOS:** (APSys '21) – resource allocation
- **MLLB:** (APSys '20) - load balancing
- **LinnOS:** (OSDI '20) – IO scheduling
- **Kleio:** (HDPC '19) - page scheduling
- **Lynx:** (NVMSA '16) - File prefetching
- On the Feasibility of Online **Malware Detection** with Performance Counters (ISCA '13)
- Others..

System

Properties

OS

Heuristics

## Conjecture:

- ML will be used in OSes
- ML can adapt to dynamic demand

**LAKE:** ask fundamental questions to *explore challenges, implications for OS design*

[1] <https://qumulo.com/blog/non-uniform-memory-access-numa/>

[2] <https://core.vmware.com/blog/understanding-persistent-memory-pmem-vsphere>

# ML in the OS: Challenges

## 1) ML is compute expensive

- Subsystem may be latency-sensitive
- Can not contend for resources with user applications
- We are currently stuck with small and cheap models

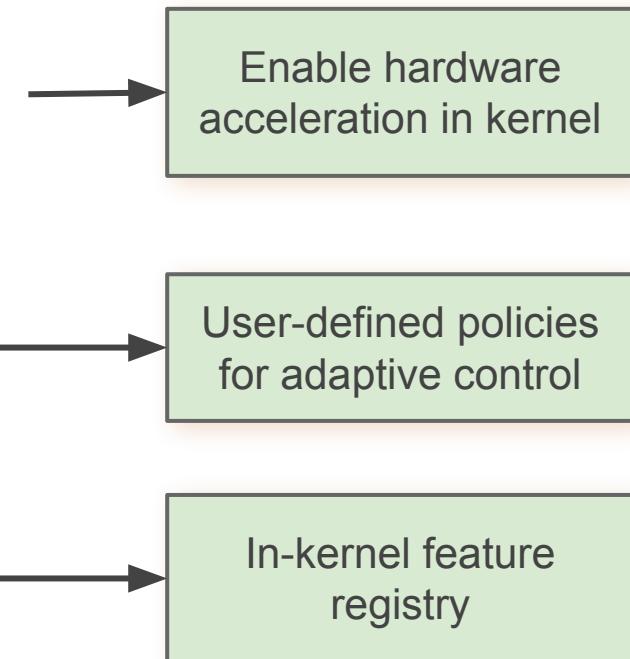
## 2) Capturing features for inference and training is difficult

- Scattered code and data structures
  - Prone to critical system level bugs
  - Aggravated by ML subsystem composition
- Locking discipline
  - Features can be in different kernel subsystems
  - Careless locking may deadlock the system

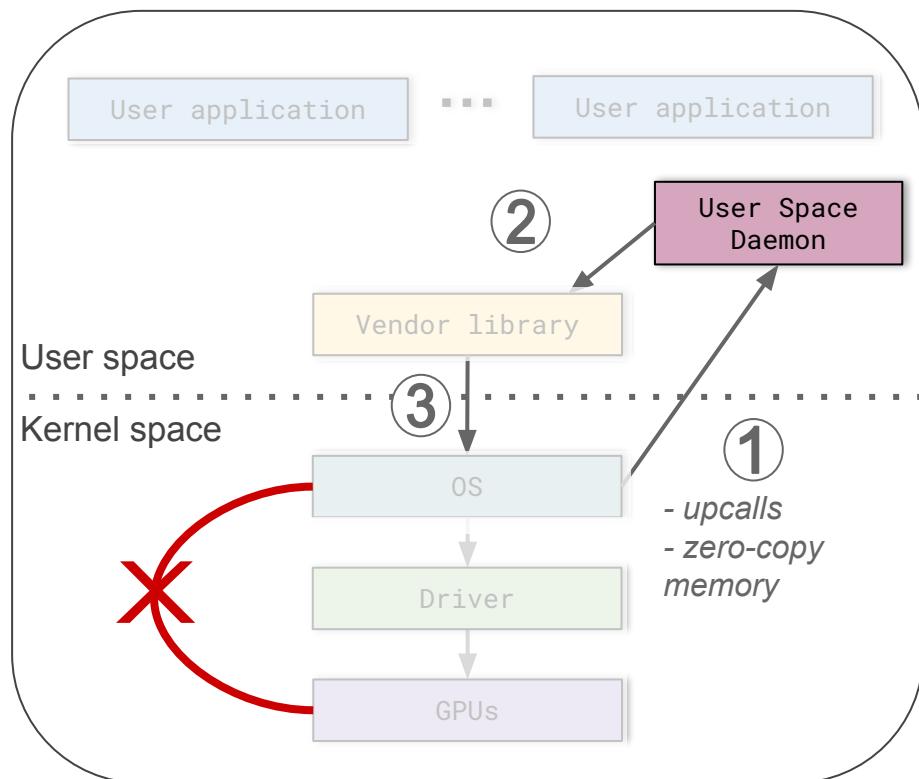
# Key Innovations in LAKE

- ML models are compute expensive
- CPU and accelerator contention
- Difficulty to manage feature vectors

## LAKE's contributions:



# Supporting Acceleration in OSes



GPU's inaccessible because libraries:

- are designed for user space
- may use **kernel bypass** communication
  - ① API calls are remoted to user space
  - ② Bulk transfers are zero-copy

Example: *cudaMemcpy*

- ② Daemon calls original accelerator API
- ④ control calls are opaque
  - e.g. command queues
- ③ Need proxy API then host API

LAKE's API  
provide a Telos

Hardware acceleration: key  
LAKE enables accelerators  
in the OS through  
upcall-based API remoting

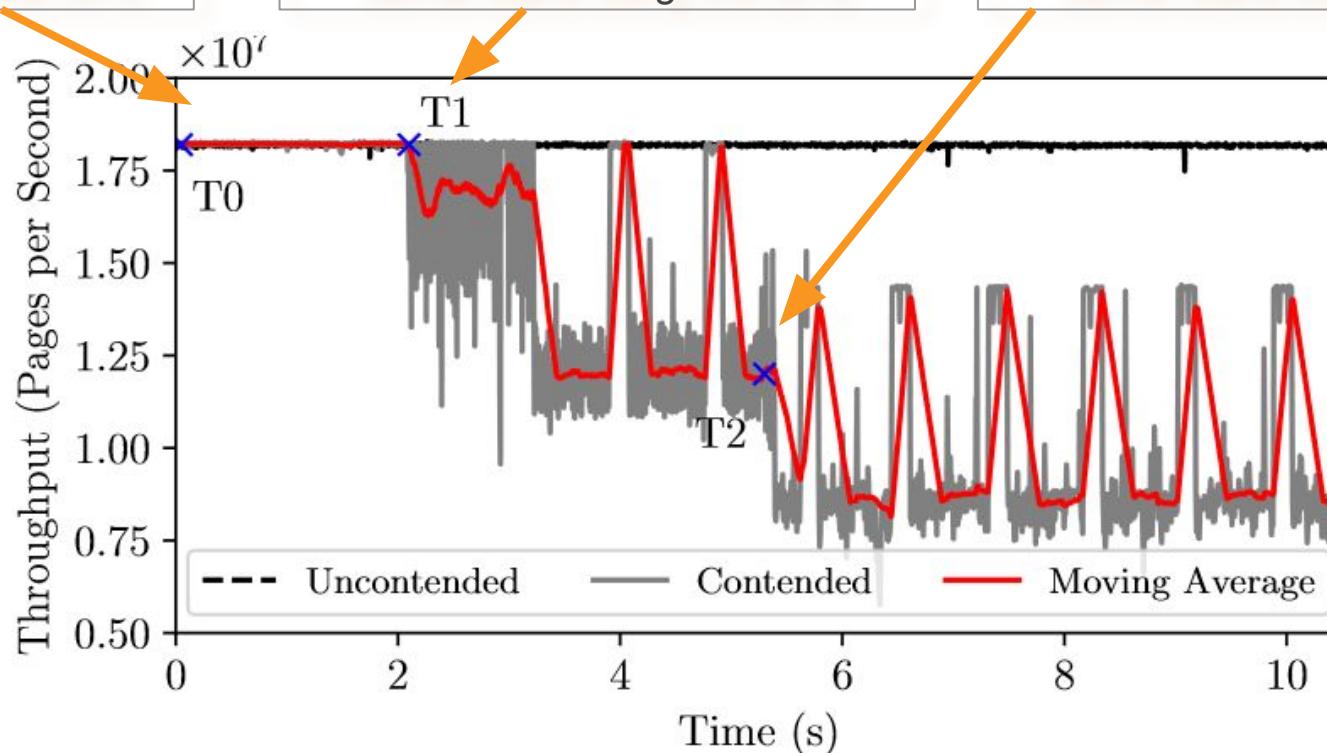
Open research question: will  
such techniques be needed  
by LDoS?

# Managing Contention for the GPU

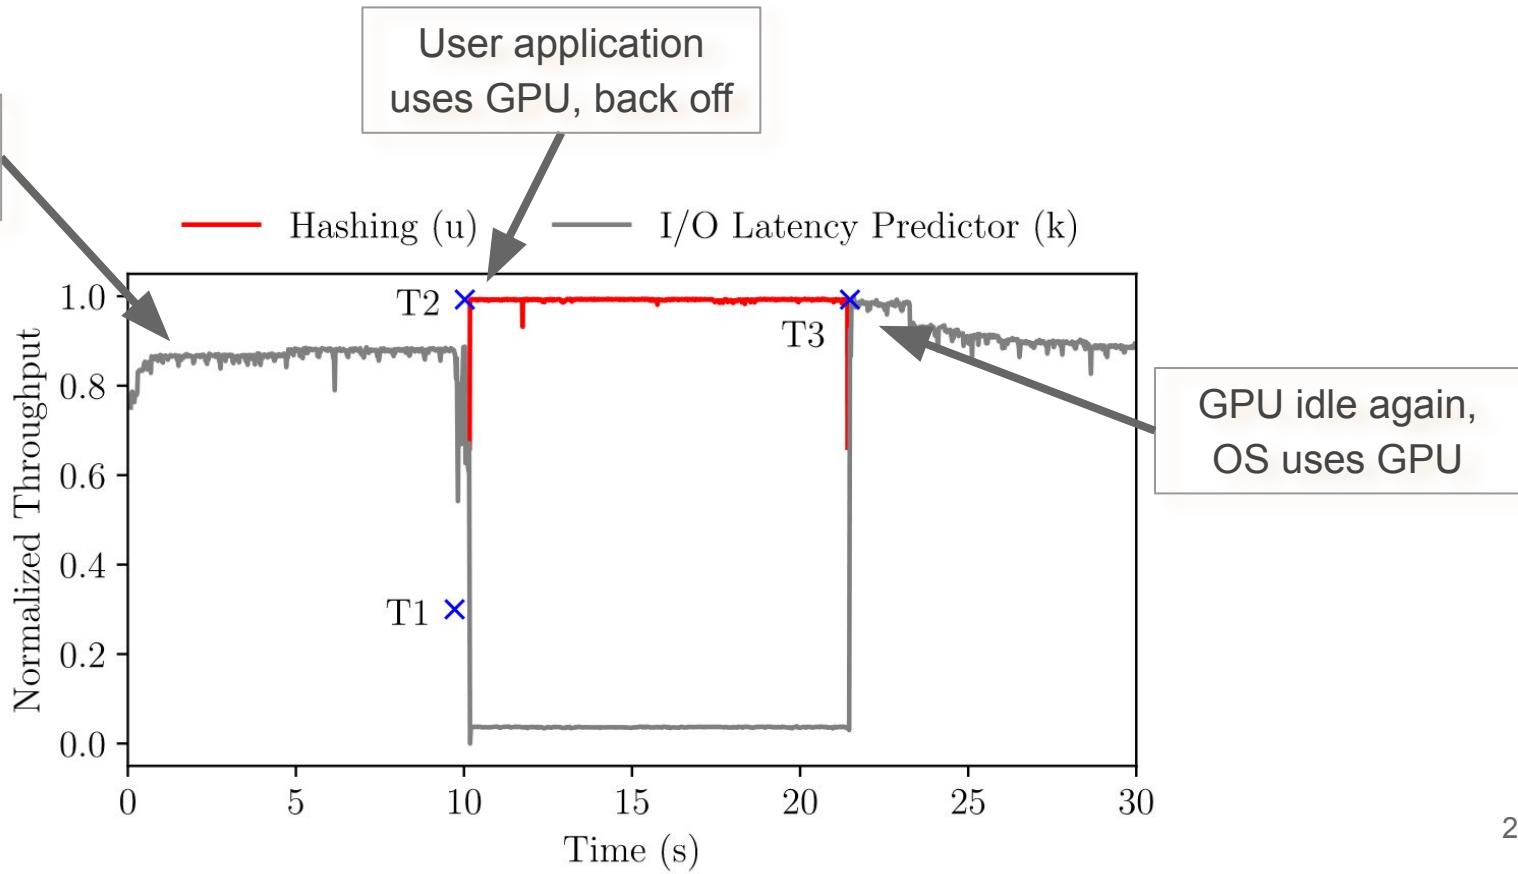
**T0:** user space is using the GPU

**T1:** page warmth classifier in the kernel starts using the GPU

**T2:** I/O latency classifier in the kernel starts using the GPU



# Managing Contention for the GPU



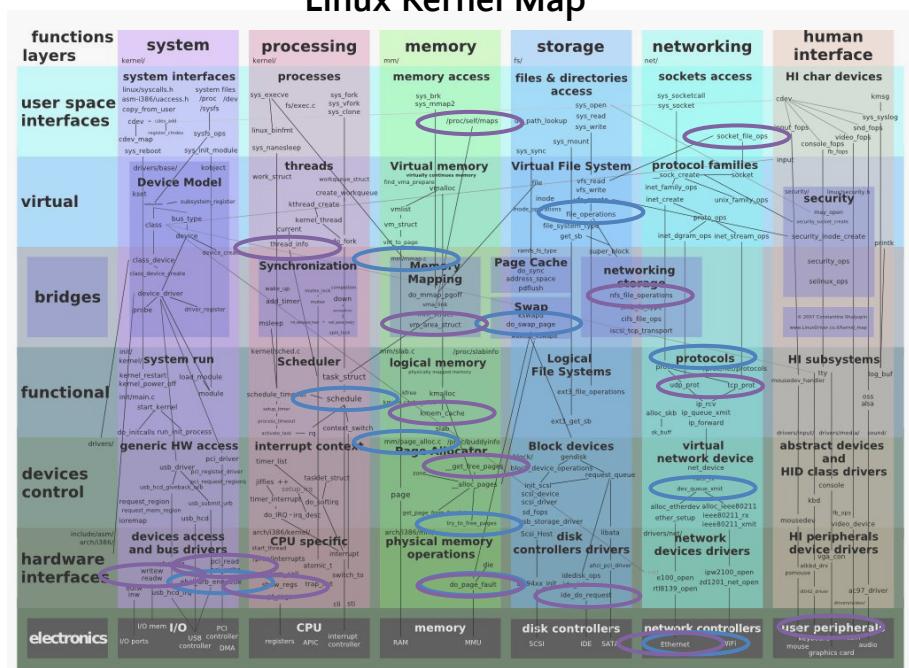
# Expressing Policies using eBPF

```
1 policy cu_policy(offload_func_t dev_func, void
2           *dev_args, offload_func_t cpu_func, void *
3           cpu_args) {
4     static nvmlUtilization_t util;
5     if      5 ms elapsed since last check
6         // LAKE-remoted nvml API
7         nvmlGetUtilization(dev, &util)
8         // compute GPU utilization moving average
9         int exec_rate = mov_avg(util.gpu);
10        // batch size for profitability threshold
11        int batch_sz = get batch size(def_args)
12        if (exec rate < exec threshold
13            && batch_sz >= batch_threshold)
14            return dev_func(dev_args);
15        else
16            return cpu_func(dev_args); }
```

LAKE manages contention and variable profitability through custom policies (eBPF)

What if the ML model produces bad solutions?  
Example: poll utilization to handle contention  
This mechanism can be augmented to support ML guardrails in LDOS  
Example: use GPU if batch is large enough

# Challenge: Poor OS Structure



Source: <https://makelinux.github.io/kernel/map/>

Issues with modern OSes:

Features and policies dispersed

Current structure a major obstacle  
to meeting requirements

# Capturing Feature Vectors

```
// feature vector struct  
  
struct feature_vector_t {  
    int io_size  
    int elapsed  
}  
  
//feature vectors array  
  
struct features_t {  
    int n_fvs  
    struct fv_t* my_fvs  
    ...  
}
```

```
submit_bio():  
    bio->start_ts = now()  
    // create new feature vector  
    new_fv = allocate_fv()  
    new_fv.io_size = bio->io_size  
  
    //store feature vector  
    store_fv(new_fv)
```

```
bio_end():  
    fv = find_fv()  
    ts_end = now()  
    fv.elapsed = ts_end - bio->start_ts  
    //store updated feature vector  
    store_fv(my_struct, new_fv)
```

Locking disciplines:

- Careless locking can cause deadlocks
- Double acquire spinlock
- Sleep/wait while holding a spinlock
- Disabling/enabling interrupt context

# LAKE's Feature Registry

```
// feature vector struct  
  
struct feature_vector_t {  
    int io_size  
    int used  
}  
  
//feature vectors array  
  
struct features_t {  
    int n_items  
    struct fv_t *fv_items  
    ...  
}
```



LAKE's feature registry facilitates feature capture and management

**submit\_bio()**:

```
bio->start_ts = now()  
begin_fv_capture ( now() )  
capture_feature (  
    "io_size", bio->io_size)
```

**bio\_end()**:

```
elaps = now() - bio->start_ts  
capture_feature ("elapsed", elaps)  
commit_fv_capture()
```

API	Description
<code>create_registry(name, sys, schema, window)</code> <code>destroy_registry(name, sys)</code>	Creates feature registry with capacity Destroys a feature registry
<code>create_model(name, sys, path)</code> <code>update_model(name, sys, path)</code> <code>load_model(name, sys, path)</code> <code>delete_model(name, sys, path)</code> <code>register_classifier(name, sys, fn, arch)</code>	Create a new ML model, saved at path Commit a changed model to the file system Load a model from path into memory Delete a model from the file system and memory Provide a function pointer for classifiers/inference Note: arch specifies CPU / GPU / XPU
<code>register_policy(name, sys, fn)</code>	Provide an eBPF policy for contention/batching (\$4.3)
<code>score_features(name, sys, fvs, num)</code> <code>get_features(name, sys, ts)</code>	Run inference on a batch, return batch results Batch retrieves all feature vectors older than ts
<code>begin_fv_capture(name, sys, ts)</code>	Starts the creation of a new feature vector. Subsequent calls to <code>capture_feature</code> for name/subsystem will add/overwrite the current value of that feature
<code>capture_feature(name, sys, key, val, sz)</code> <code>capture_feature_incr(name, sys, key, incrval, sz)</code> <code>commit_fv_capture(name, sys, ts)</code> <code>truncate_features(name, sys, ts)</code>	Sets feature with key, val on the current vector Update a feature with key by incrementing Commits the current feature vector to the registry. Removes all feature vectors older than ts

# Evaluation: I/O Scheduling

**Setup:** server with three NVMeS with replicated data for redundancy

**Goal:** classify I/Os issued to an NVMe as fast or slow

**Actions:**

**Fast** → No-op

**Slow** → Re-issue to different device

**How?** neural network with two layers, 256 and 2 neurons, respectively.

We simulate **more complex models** by adding **+1** and **+2** layers to the NN

Testbed: 16-core Intel Xeon Gold 6226R  
CPUs, 376 GiB DDR4 RAM, two  
NVIDIA A100 GPUs and three Samsung 980  
Pro 1TB (PCIe 4.0) NVMeS

# Evaluation: I/O Scheduling

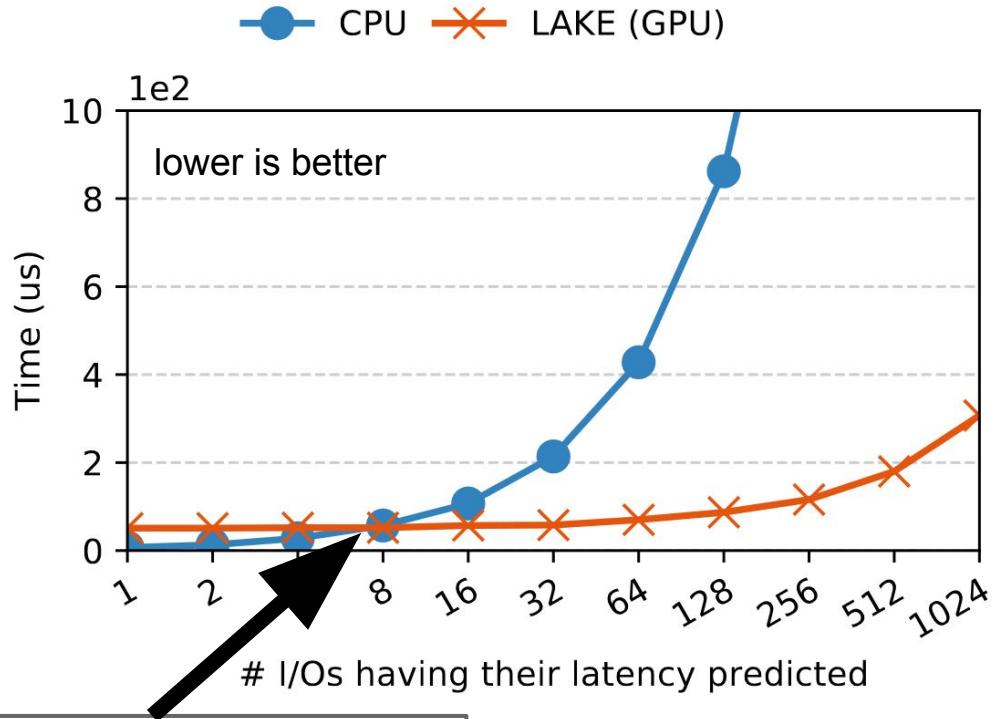
X axis is the number of I/Os we are classifying in **one batch**.

Y is the **total time to perform inference**, the lower the better

**At what point does using a GPU become beneficial?**

(we call this the **crossover point**)

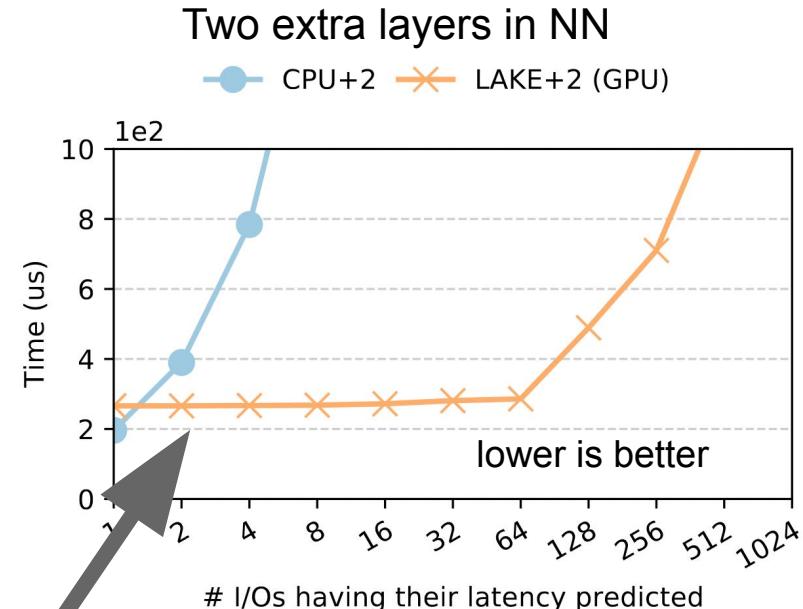
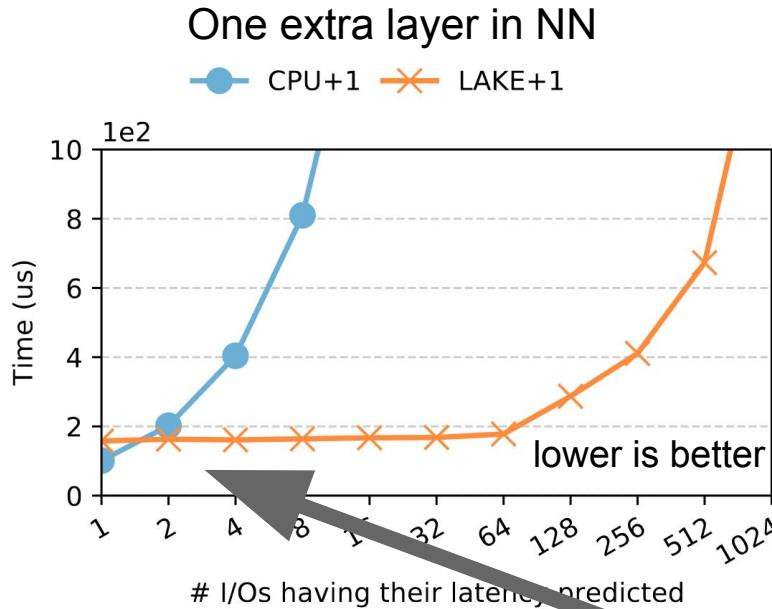
This is a small model, NN with 2 layers



With 8 or more inputs, using a GPU is profitable

Testbed: 16-core Intel Xeon Gold 6226R CPUs, 376 GiB DDR4 RAM, two NVIDIA A100 GPUs and three Samsung 980 Pro 1TB (PCIe 4.0) NVMe

# Evaluation: I/O Scheduling



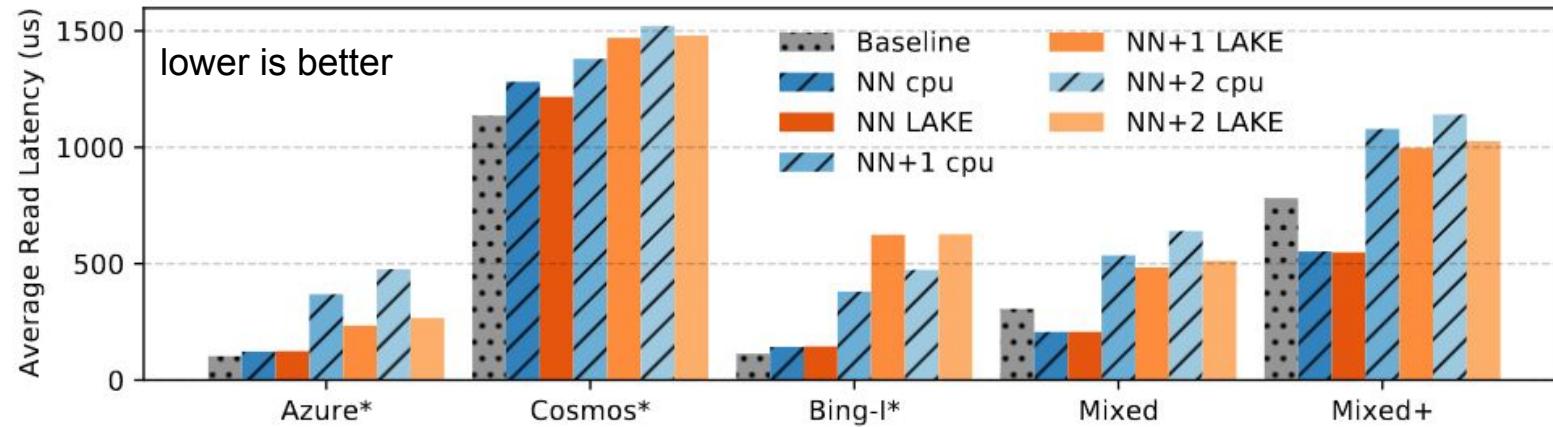
Profitability increases faster for:

- More complex models
- More inputs

For larger models, GPU is faster with just two inputs

Testbed: 16-core Intel Xeon Gold 6226R CPUs, 376 GiB DDR4 RAM, two NVIDIA A100 GPUs and three Samsung 980 Pro 1TB (PCIe 4.0) NVMe

# Evaluation: I/O Scheduling



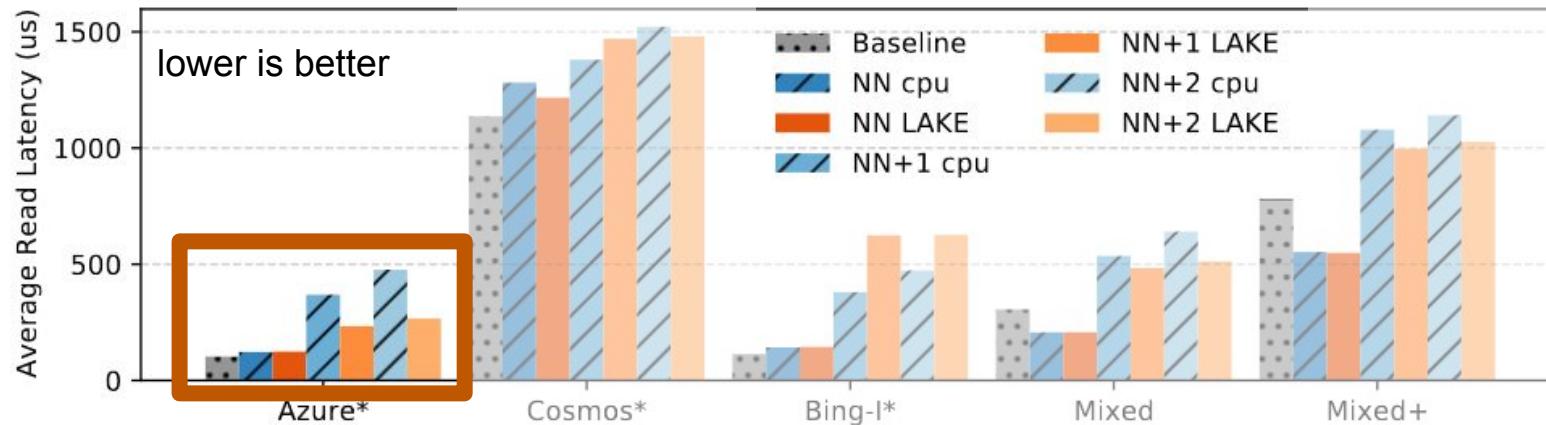
Workload properties:

Trace Name	Avg IOPS	Avg Read/Write I/O size (KB)	Min/Max Arrival Time (us)
Azure	26k	30/19	0/324
Bing-I	4.8k	73/59	0/1.8k
Cosmos	2.5k	657/609	0/1.6k

Testbed: 16-core Intel Xeon Gold 6226R CPUs, 376 GiB DDR4 RAM, two NVIDIA A100 GPUs and three Samsung 980 Pro 1TB (PCIe 4.0) NVMe

# Evaluation: I/O Scheduling

Trace Name	Avg IOPS	Avg Read/Write I/O size (KB)	Min/Max Arrival Time (us)
Azure	26k	30/19	0/324
Bing-I	4.8k	73/59	0/1.8k
Cosmos	2.5k	657/609	0/1.6k



High IOPS (Azure) -> higher batch sizes, great for GPUs

Policy uses the GPU  
when it observes an  
I/O burst

Testbed: 16-core Intel Xeon Gold 6226R CPUs, 376 GiB DDR4 RAM, two NVIDIA A100 GPUs and three Samsung 980 Pro 1TB (PCIe 4.0) NVMe

# Evaluation: I/O Scheduling

Trace Name	Avg IOPS	Avg Read/Write I/O size (KB)	Min/Max Arrival Time (us)
Azure	26k	30/19	0/324
Bing-I	4.8k	73/59	0/1.8k 0/1.6k

Average Read Latency (us)

1500  
1000  
500  
0

## Other LAKE results:

- Five other workloads
- Non-ML: eCryptfs
- CPU utilization

Small I/Os

frequent,  
CPU

on Gold 6226R  
M, two  
980 Pro 1TB (PCIe 4.0) NVMe

# LAKE: what did we learn?

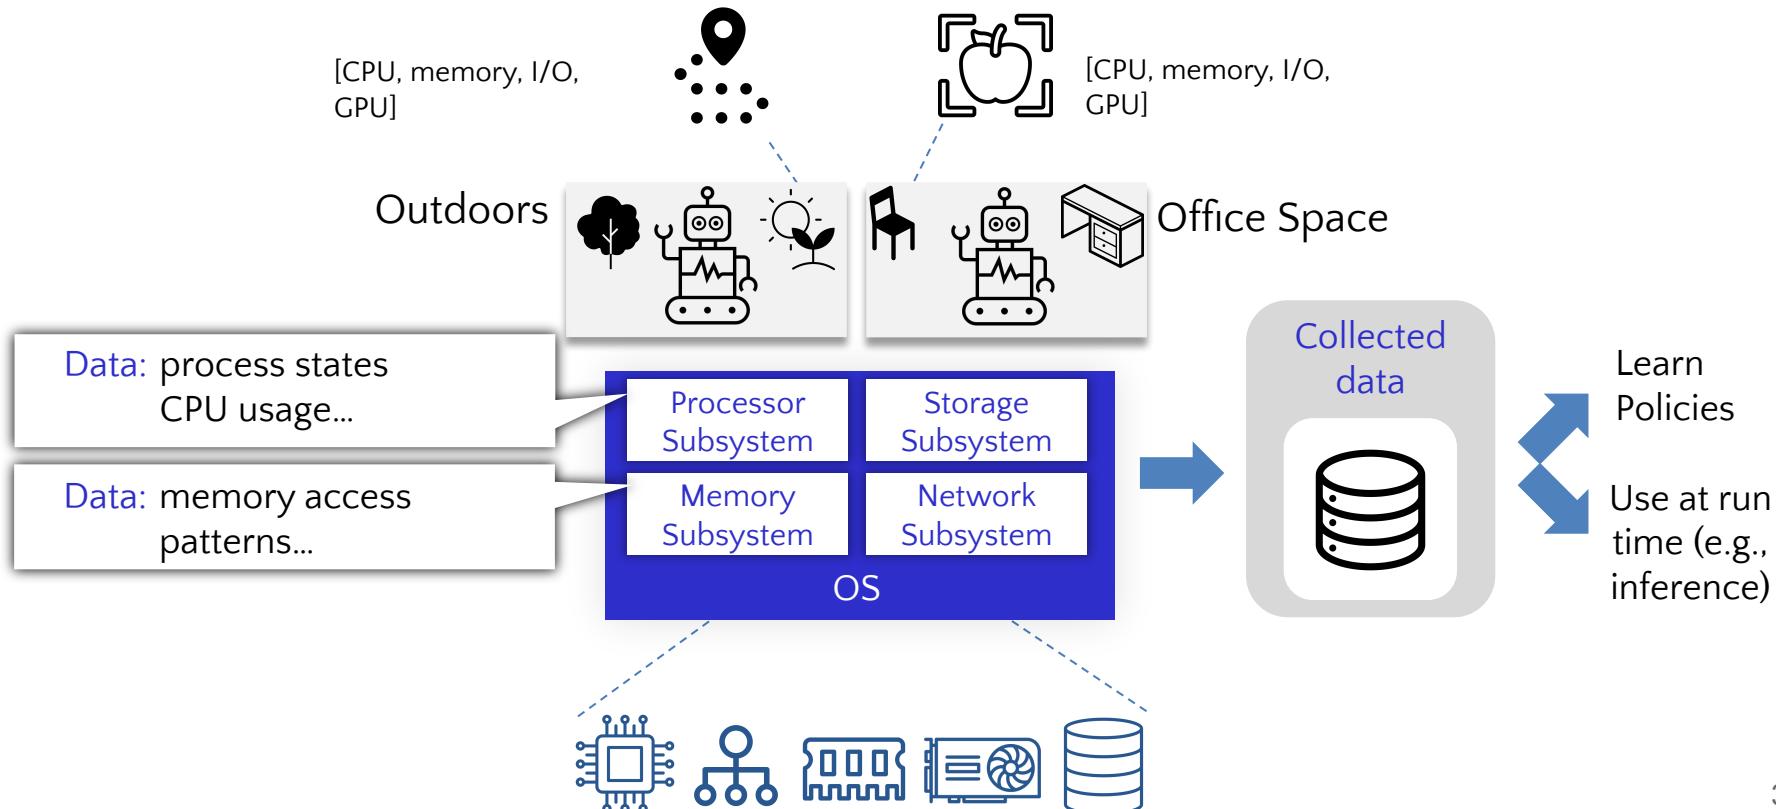
- *Expected to learn:*
  - ML is super-applicable and works great in the kernel
- *Actually learned* about challenges:
  - Robustness
  - Data is hard to come by
  - Choosing learning techniques
  - OS+ML may need *different* acceleration
  - Explainability/Misbehavior/pathologies
  - Development cycle mismatches
  - OS structure

Today

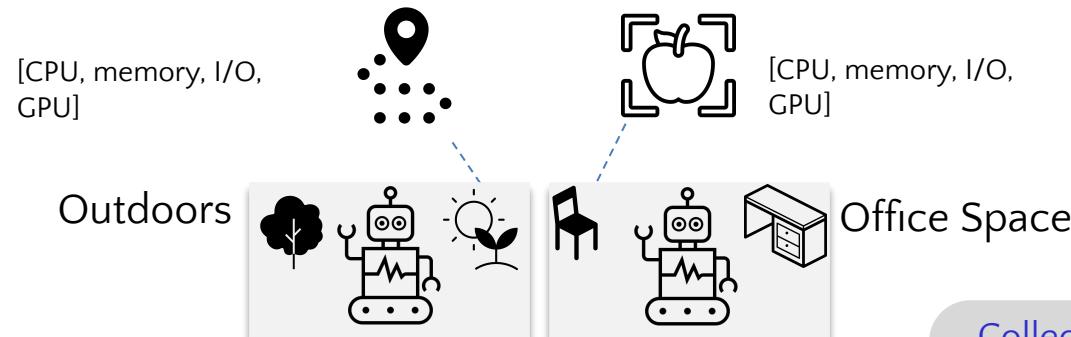
# Robustness

	API?	Contention Control?	Complex models?	End-to-end evaluation?	Reproducible?
KML [1]	✓	*	✗	✓	...
SmartOS [2]	✗	✗	✗	✓	...
LinnOS [3]	<ul style="list-style-type: none"><li>Results difficult to reproduce</li><li>Results very sensitive to env. / exp. design</li><li><b><i>Not a criticism of previous work:</i></b></li><li><b><i>These challenges are fundamental</i></b></li></ul>			✓	...
Lynx [4]				✓	...
Kleio [5]				✗	...
LAKE				✓	...

# Data Challenge



# Data Challenge

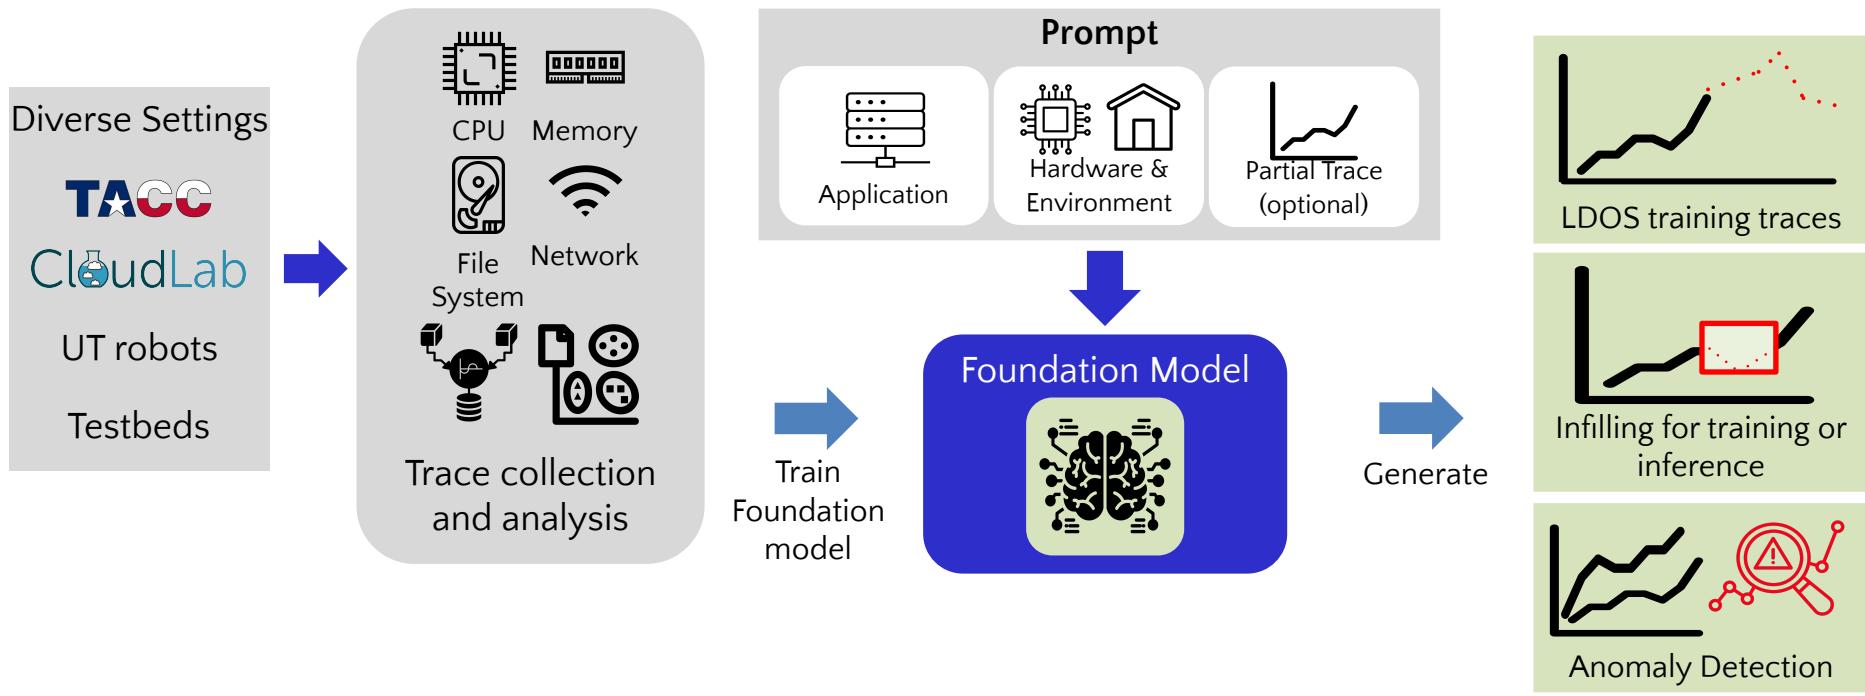


Data collection overhead

Privacy and regulatory constraints

Unseen applications and environments

# LDOS Solution: Foundation Model for OS Data



# Learning Techniques for ML in OS

## Supervised learning techniques

- Require ground truth
- Require lots of data
- Hard to improve over existing heuristics
- Existing sources of ground truth are heuristics
- Ground truth data improving over heuristics: not scalable

## Unsupervised techniques

- Can improve over heuristics (e.g. RL)
- Computationally expensive
- Can be slow to converge on best policy

# ML, OS, Hardware Acceleration

LAKE design adapts to current hardware

- GPUs/TPUs designed for user-level, OS-independence
- GPUs/TPUs designed for high-throughput batch compute
- Kernels care a lot about latency
- In-kernel ML still dominated by parallel compute (e.g. GEMM)

Conjecture: Hardware will evolve if ML+OS benefits

Consequence: near term challenge

- LDOS will likely focus on CPU-amenable ML
- Evaluating benefits will be tough for many ML techniques

# Addressing Challenges with LDOS

- Robustness      ○    **Verified Learning**
- Data is hard to come by      ○    **Foundation Model for OS**
- Explainability/Misbehavior      ○    **Guardrails / Adaption**
- Development cycle mismatches      ○    **KernMLOps**
- OS structure      ○    **Clean-slate Design**
- Choosing learning techniques      ○    <>
- OS+ML: *different* acceleration      ○    <>

# Conclusion

## LDOS

- Clean-slate OS
- ML-driven resource management policy architecture
- Cross-cutting : generative AI, ML, formal methods, and systems

Many exciting research directions!

We welcome collaborators!

<https://ldos.utexas.edu>

Questions?