


Gradient Computation in Neural Networks and Kolmogorov-Arnold Networks

**Keshav Pingali and Lain Mustafaoglu
The University of Texas at Austin**

Organization

- **Problem**
 - Parameter estimation in neural networks
 - Non-linear optimization problem
 - Solved using gradient descent
 - Usual presentations of gradient descent are hard to understand and difficult to extend to irregular network connections
- **Parameterized programs**
 - Abstraction for neural networks
- **Gradient computation in parameterized programs**
 - Backward dataflow analysis
- **Advantages**
 - Compositional algorithm for gradient computation
 - Handles weight-sharing
 - Easy to extend to skip connections and other irregular networks
 - Natural extension to higher dimensional data such as tensors
- **Extension to Kolmogorov-Arnold Networks (KANs) and Recurrent Neural Networks (RNNs)**

Multilayer Perceptron (MLP) Example



Frank Rosenblatt (Cornell)
Inventor of Perceptron

- **Type:** Inputs: $x_1..x_4$, outputs: y_1, y_2 (all \mathbb{R})
- **Scalar view:**
 - Each edge performs a multiplication with a real-valued **parameter**
 - Each neuron adds its input values and applies a non-linear operation f such as tanh, ReLU etc. (known as **activation functions**)
- **Vector view:**
 - Each layer performs a dense matrix vector multiplication
 - Followed by pointwise (map) non-linear operation f
- Gradient computation difficult to understand
- Abstraction of MLP and more complex neural networks: **parameterized programs**

Parameterized program: running example

- Type of desired function: real x real \rightarrow real
- Training data: set of N 3-tuples $\{(p_i, q_i, t_i)\}$
- Model
 - **Composed from**
 - > **base functions** f_i (may be **nonlinear** such as tanh, sigmoid, sin, cos, ...)
 - > **parameters** W_i : *real* (assume no weight sharing so each weight occurs just once)
- Notation: capital letters for variable names, small letters for variable values
 - Function invocation written as $R(w; p_i, q_i)$ where w is (w_0, w_1, w_2)
- Parameter optimization
 - Square error for training sample $(p_i, q_i, t_i) = (t_i - R(w; p_i, q_i))^2$
 - Goal: choose (w_0, w_1, w_2) to minimize mean square error
$$\text{Loss}(w_0, w_1, w_2) = \frac{1}{N} \sum_{i=1}^N (t_i - R(w; p_i, q_i))^2$$

Function $R(P, Q)$ {
 *$A = W_0 * P$*
 $B = f_0(A, Q)$
 *$C = W_1 * B$*
 *$D = W_2 * B$*
 $E = f_1(C)$
 $F = f_2(D)$
 $R = f_3(E, F)$
return R }

Parameter optimization

- Find derivatives of Loss wrt W_0, W_1, W_2

$$\text{Loss}(w_0, w_1, w_2) = \frac{1}{N} \sum_{i=1}^N (t_i - R(w; p_i, q_i))^2$$

$$\frac{\partial \text{Loss}}{\partial W_0}(w_0, w_1, w_2) = -\frac{2}{N} \sum_{i=1}^N (t_i - R(w; p_i, q_i)) \frac{\partial R}{\partial W_0}(w; p_i, q_i)$$

$$\frac{\partial \text{Loss}}{\partial W_1}(w_0, w_1, w_2) = -\frac{2}{N} \sum_{i=1}^N (t_i - R(w; p_i, q_i)) \frac{\partial R}{\partial W_1}(w; p_i, q_i)$$

$$\frac{\partial \text{Loss}}{\partial W_2}(w_0, w_1, w_2) = -\frac{2}{N} \sum_{i=1}^N (t_i - R(w; p_i, q_i)) \frac{\partial R}{\partial W_2}(w; p_i, q_i)$$

$$\nabla_W R(w; p_i, q_i)$$


Function $R(P, Q)$ {

$A = W_0 * P$

$B = f_0(A, Q)$

$C = W_1 * B$

$D = W_2 * B$

$E = f_1(C)$

$F = f_2(D)$

$R = f_3(E, F)$

return R }

Derivatives are complicated, non-linear functions
so use gradient-descent for parameter optimization
-> Focus on gradient computation

Parameterized program as flow graph

Function $R(P,Q)$ {

$A = W0 * P$

$B = f0(A,Q)$

$C = W1 * B$


$D = W2 * B$

$E = f1(C)$

$F = f2(D)$

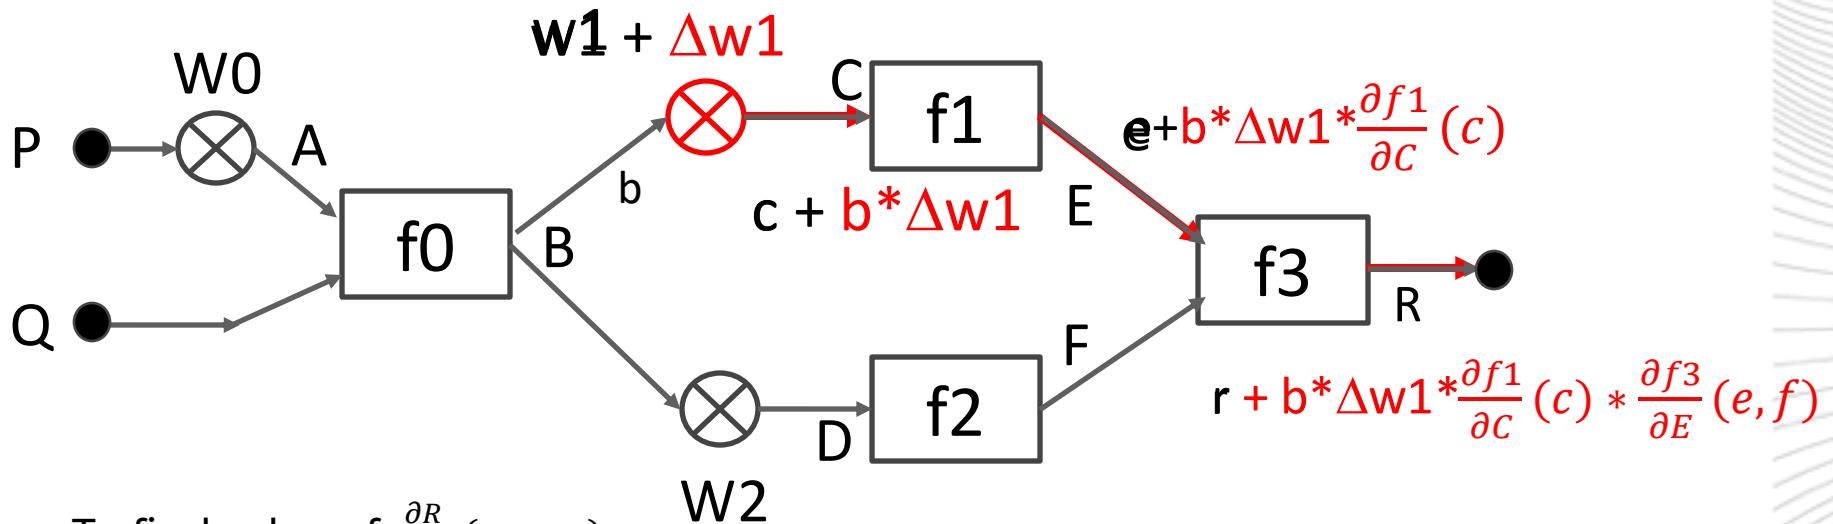
$R = f3(E,F)$

return R }



- Useful to represent multiplication by weights differently from functions f_i
 - Weights change during training
- Forward Propagation: execute nodes in any topological order
- All variable values are stored
 - Needed for gradient computations

Value of $\frac{\partial R}{\partial w_1}(w; p, q)$



- To find value of $\frac{\partial R}{\partial w_1}(w; p, q)$
 - Multiply values of partial derivatives of all vertices on path from w_1 to R
 - Multiply result by value of input to w_1 (i.e., b)
 - Result: $b * \frac{\partial f_1}{\partial c}(c) * \frac{\partial f_3}{\partial E}(e, f)$
 - Compute the product either forwards or backwards along path

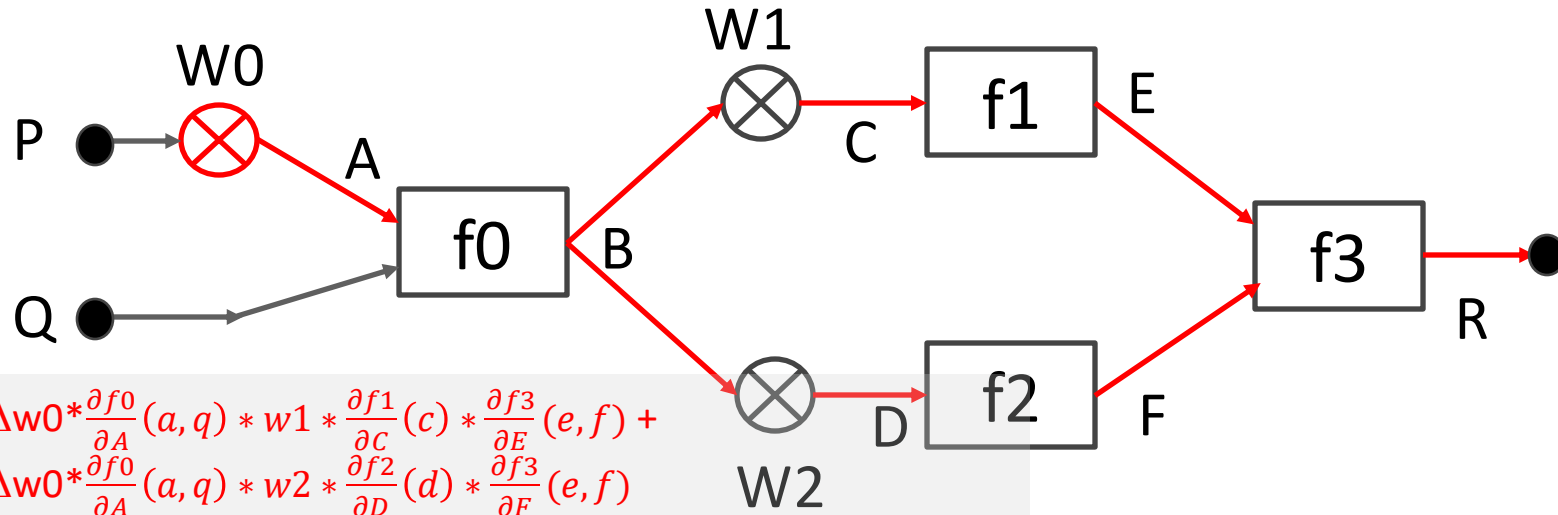
- In general, for path $\rho : X \xrightarrow{*} Y$ *path derivative* $\pi(\rho)$

= product of derivatives of nodes on path excluding X and Y

= 1 for empty path or if there are no intermediate nodes

- $\frac{\partial R}{\partial w_1}(w; p, q) = b * \pi(w_1 \xrightarrow{*} R)$

Value of $\frac{\partial R}{\partial w_0}(w; p, q)$

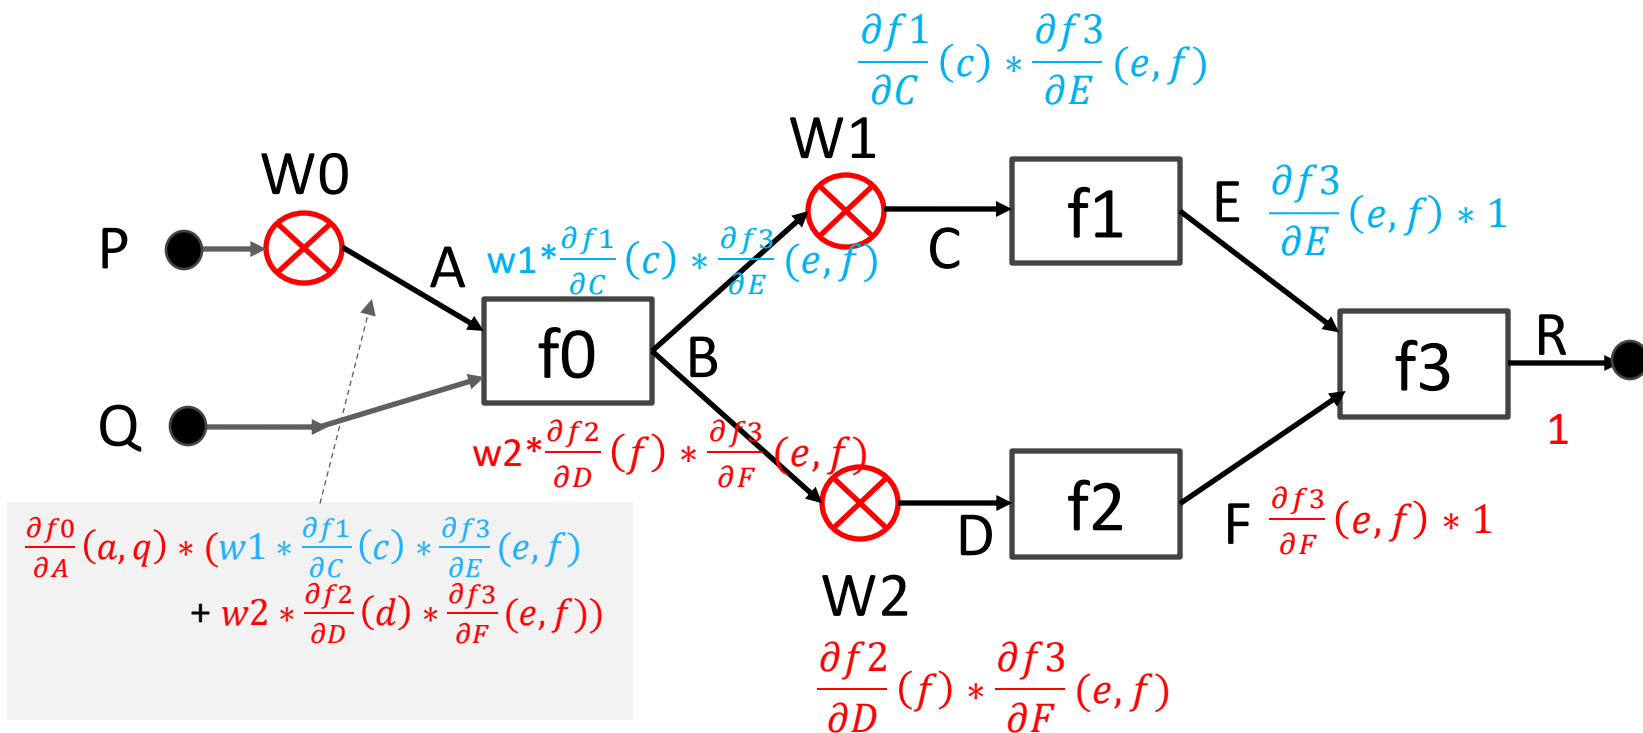


- In general, there is a DAG from weight to the output
- Value of partial derivative:
 - Enumerate all paths from weight to output and add up the contributions of all paths
$$\frac{\partial R}{\partial w_0}(w; p, q) = p * \sum_{\rho \in \mathcal{P}(W_0)} \pi(\rho)$$

(where $\mathcal{P}(W_0)$ is the set of paths from W_0 to exit)

 - Intuition: derivatives make this a linear problem, so *superposition of paths* works
- Problems:
 - Treats DAG like tree so could do exponential computation in size of DAG. More efficient solution?
 - What order should we compute $\frac{\partial R}{\partial w_0}(w; p, q)$, $\frac{\partial R}{\partial w_1}(w; p, q)$ and $\frac{\partial R}{\partial w_2}(w; p, q)$?

Efficient computation of all derivatives



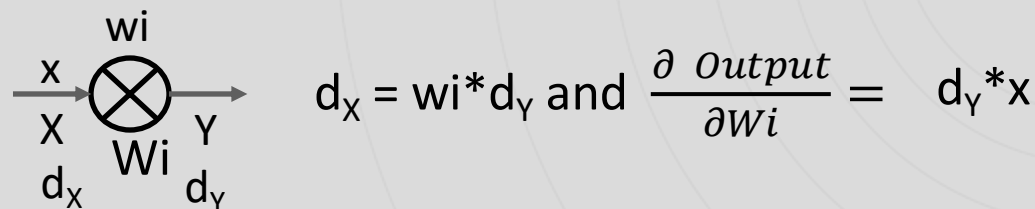
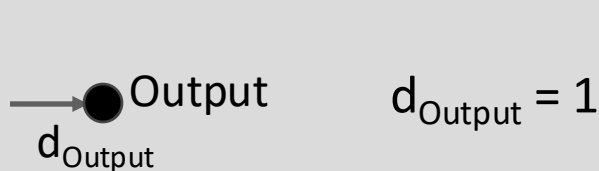



- Compute derivative of output wrt every variable/edge ($\frac{\partial R}{\partial A}, \frac{\partial R}{\partial C}, \frac{\partial R}{\partial F} \dots$)
 - Real number on each edge
 - Derivatives of output wrt weights can be computed from this
- Traverse DAG in reverse topological order of variables for computation
 - $\frac{\partial R}{\partial R} = 1$
 - **Transfer functions** to propagate derivative from function output to its inputs

Summary of derivatives computation

- At each point X , compute $d_x: \Re$ where

$$d_x = \frac{\partial \text{Output}}{\partial X}(w; p, q) = \sum_{\rho \in \mathcal{P}(X)} \pi(\rho)$$
- Small tweak to handle weight-sharing
- Called back-propagation in ML literature



Back to running example

- Optimization problem: choose W_i values to minimize loss

$$\text{Loss}(w_0, w_1, w_2) = \frac{1}{N} \sum_{i=1}^N (t_i - R(w; p_i, q_i))^2$$

Function $R(P, Q)$ {

$A = W_0 * P$

$B = f_0(A, Q)$

$C = W_1 * B$

$D = W_2 * B$

$E = f_1(C)$

$F = f_2(D)$

$R = f_3(E, F)$

return R }

$$\frac{\partial \text{Loss}}{\partial W_0}(w_0, w_1, w_2) = -\frac{2}{N} \sum_{i=1}^N (t_i - R(w; p_i, q_i)) \left(\frac{\partial R}{\partial W_0}(w; p_i, q_i) \right)$$

Initialize weights to random values

for #epochs do {

GradientVector = 0

 for each training sample (p_i, q_i, t_i) do {

 perform forward propagation and compute

 perform backpropagation and compute weight derivatives

 update GradientVector with products}


 scale GradientVector by $-2/N$

 use GradientVector to update weights using gradient descent

step

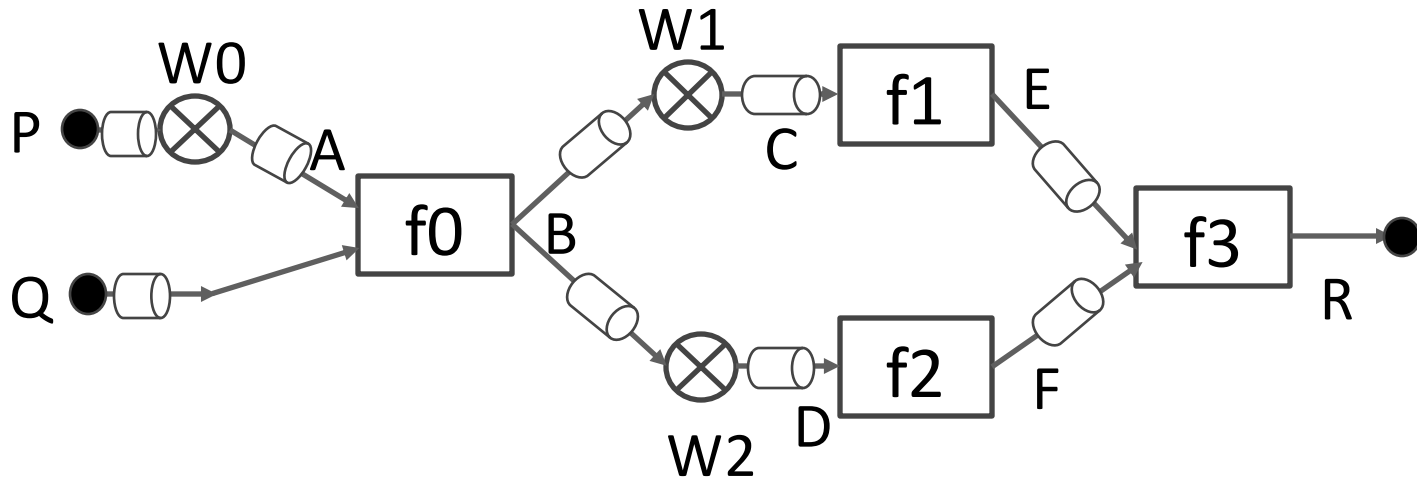
}

Generalization to vectors, matrices, tensors




How to understand data with dogs
Karl Stratos ([Reddit post](#))

Generalization to vectors (matrices, tensors are similar)



- **Programs**
 - Inputs can be scalars or **vectors**, but **output is still a scalar** (loss)
 - Weights → Weight matrices
 - > Weight matrix need not be square
 - Function type: vector → vector, vector x vector → vector, etc.
- **Compute gradients instead of derivatives**
 - Variable is vector of size n → gradient of output wrt this vector is also vector of size n
 - > Intuition: each dimension of gradient vector is derivative of output wrt value in that dimension
 - Transfer functions: Jacobians instead of derivatives
- Useful to know [matrix derivatives](#) notation

Transfer functions



- General function f

- $\underline{d}_X = J_f^* \underline{d}_Y$

- Linear function W

- $J_f = W^T$

- $\underline{d}_X = W^T * \underline{d}_Y$

- Derivatives of output wrt weights in W

$$\frac{\partial \text{Output}}{\partial W} = \underline{d}_Y \otimes \underline{x} \quad (\otimes \text{ is outer-product})$$



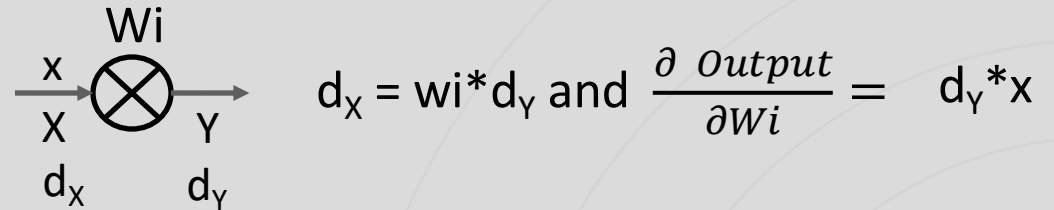

$$\frac{\partial \text{Output}}{\partial W}(i, j) = \underline{d}_Y(i) * \underline{x}(j)$$

(If $w(i,j)$ changes a small amount, how much does the output change?)





$$\begin{aligned} y_1 &= w_{11}x_1 + w_{12}x_2 + \dots + w_{1m}x_m \\ y_2 &= w_{21}x_1 + w_{22}x_2 + \dots + w_{2m}x_m \\ &\dots \end{aligned}$$

Transfer Functions (scalar and vector)


Scalar case




Vector case



Important special cases



map



reduce

- **f: map** ($g: \mathbb{R} \rightarrow \mathbb{R}$)

- J_f = diagonal matrix with $J_f(i,i) = g'(x(i))$

$$\underline{d}_X = \begin{pmatrix} g'(\underline{x}(1)) & 0 \\ 0 & g'(\underline{x}(2)) \end{pmatrix} \underline{d}_Y$$

- **f: reduce** ($r: \mathbb{R} \times \mathbb{R} \rightarrow \mathbb{R}$)

- Called **pooling** in ML literature

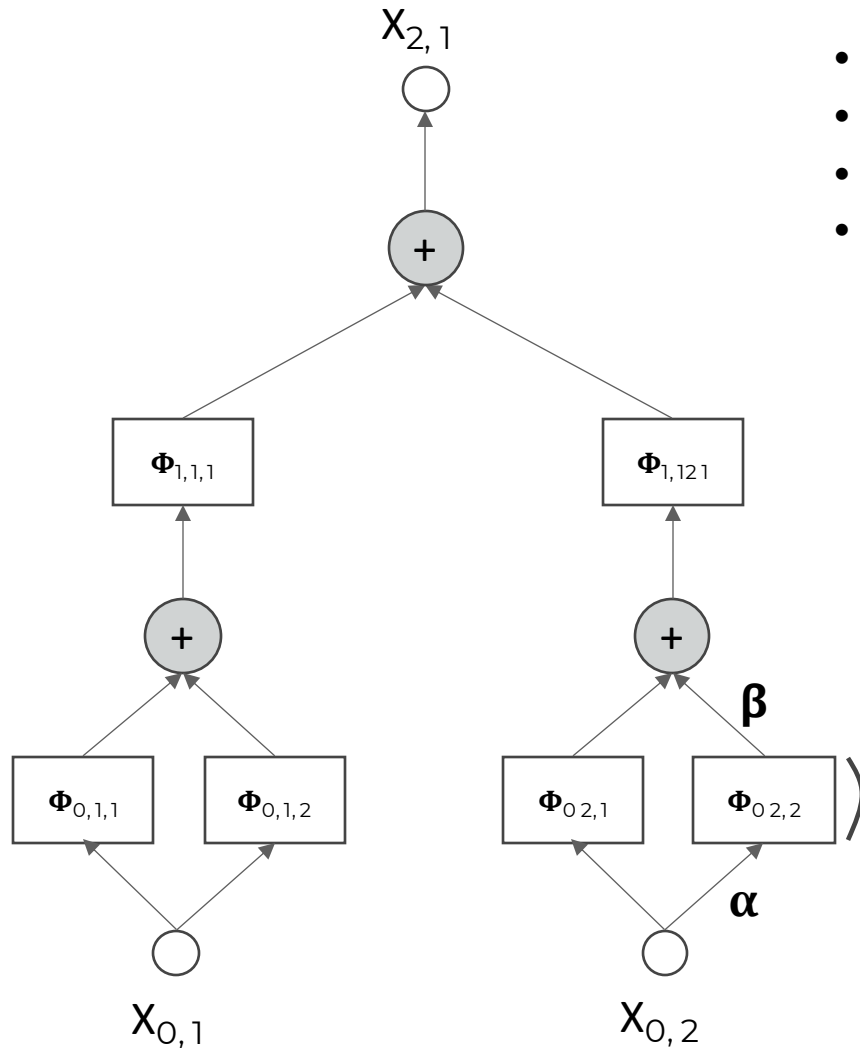
- $r = +, *, \dots$

> $J_f = (1, 1, \dots, 1)^T$ for +

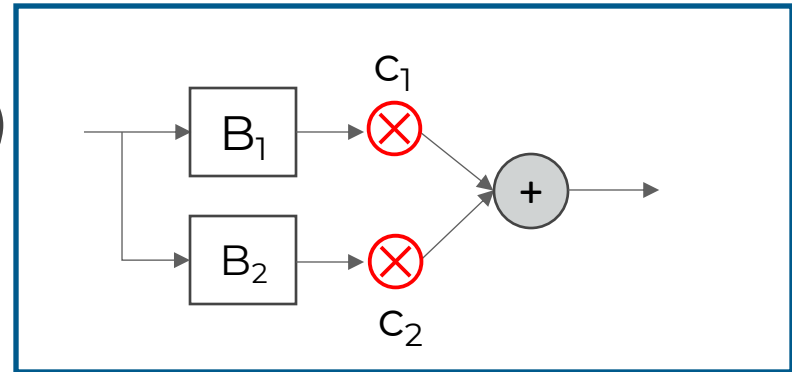
- $r = \max, \min, \dots$

> If $\max(x) = x(j)$, $J_f = l(j)$ (where $l(j)$ is the indicator vector with 1 in j^{th} position)

Kolmogorov-Arnold Networks (KANs)



- Alternative to MLPs
- B_i : **B-spline**
- $\Phi = \sum_i c_i B_i$: **spline**
- Training: backpropagation to adjust B-spline control points c_i



Recurrent Neural Networks (RNNs)

Running example


- RNNs: input is sequence
- Machine translation
 - Input: sentence (sequence of words) in English
 - Output: sentence (sequence of words) in French
- Training data
 - Set of sentence pairs: (sentence in English, sentence in French)
- Abstractly we want a function of this type
 - $F: [x_1, x_2, \dots, x_m] \rightarrow [y_1, y_2, \dots, y_n]$ (m,n can be different for different sentences)
 - Input sequence can be of arbitrary length
 - Assume $m=n$ for simplicity
- Questions
 - How do we encode (represent) words?
 - How do we encode sequences of words?
 - How do we handle arbitrarily long sequences?
 - How is output produced?

Encoding words and sequences of words

Dictionary



cat:	$[1,0,0,0]^T$
horse:	$[0,0,1,0]^T$
is:	$[0,1,0,0]^T$
small:	$[0,0,0,1]^T$

lookup: word \rightarrow one-hot encoding



- **Vector space model:** words and sequences of words **embedded** as points in $H = \mathbb{R}^m$
- Embedding of word x : $h = W_{xh} * \text{lookup}(x)$
 - $\text{lookup}()$ uses dictionary to map word x to its one-hot encoding
 - W_{xh} is learned: column m is embedding of word with 1 in the m^{th} position of one-hot encoding
- Embedding of sequence (e.g.) $[x_1 \ x_2]$
 - One possibility: add embeddings of x_1 and x_2
 - Drawback: $[\text{small cat}]$ will have same embedding as $[\text{cat small}]$
 - Better idea: $W_{hh} * h_1 + h_2$ (where W_{hh} is learned)
- In general, H : sequence of words $\rightarrow \mathbb{R}^m$
 - $H([\]) = \mathbf{0}$
 - $H([x_1 \ x_2 \dots x_{i-1} \ x_i]) = W_{hh} * H([x_1 \ x_2 \dots x_{i-1}]) + W_{xh} * \text{lookup}(x_i)$


In practice



$$\frac{dy}{dx} = 1 - \tanh^2(x)$$

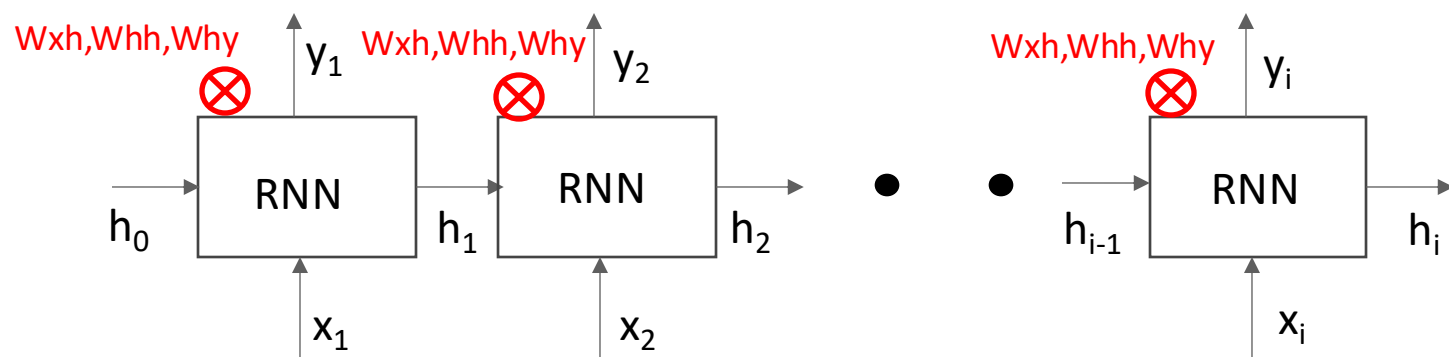
- Use tanh to squash H encodings into unit hypercube to prevent blow-up
 - $H([x_1 \ x_2 \dots \ x_{i-1} \ x_i]) = \tanh(W_{hh} * H([x_1 \ x_2 \dots \ x_{i-1}]) + W_{xh} * \text{lookup}(x_i))$
- Output for simple RNN produced “online”
 - y_i depends only on $[x_1, \dots, x_i]$
 - $y_i \sim \text{softmax}(W_{hy} * H([x_1 \ x_2 \dots \ x_{i-1} \ x_i]))$
 - > Output of softmax = probability vector for next output word
 - > y_i is sampled from output distribution of softmax

RNN implementation



- RNN is single-layer neural network with feedback loop
- $H([x_1 \ x_2 \dots x_{i-1} \ x_i])$ represented as vector h_i
 - Fed back to next iteration
- Details
 - W_{xh} can be initialized to embeddings from Word2Vec
 - RNNs can be chained to form “multi-layer” RNNs


Training RNNs



- Training data: $\{[x_1, x_2, \dots, x_n] \rightarrow [Y_1, Y_2, \dots, Y_n]\}$
 - Notation: Y_i is training data, y_i is output produced by RNN during “inference”
- Training: at each step
 - Compute cross-entropy between ground truth Y_i and computed value y_i
 - > Strictly speaking, between one-hot encoding of Y_i and output of softmax at step i
 - Back-propagate using weight-sharing to update weights
 - In practice, limit the size of the “look-back” window to 3-4
- Analogy: path-sensitive dataflow analysis

Improving RNNs: Encoder-decoder architectures

- Requiring output to be produced online means
 - No “look-ahead” in input stream is possible when determining how to produce next output word
 - Input and output sequences must have same length
- Solution
 - First encode entire input sequence (encoder)
 - Then produce output one word at a time (decoder)
- Two architectures
 - Baseline encoder-decoder architecture based on RNNs
 - Transformers



Baseline encode-decoder architecture

- We want to learn a function $F: [x_1, x_2, \dots, x_n, y_1, y_2, \dots, y_{i-1}] \rightarrow y_i$
- Training input: $[x_1, \dots, x_n], [Y_1, \dots, Y_n]$
- Encoder
 - $h_i = f_1(h_{i-1}, x_i)$ $f_1 = \tanh(W_{hh} * h_{i-1} + W_{xh} * \text{lookup}(x_i))$
> $h_0 = 0$
 - h_n = thought vector (embedding of $[x_1, x_2, \dots, x_n]$)
- Decoder (training)
 - $h_{n+i} = f_1(h_{n+i-1}, Y_{i-1})$ (embedding of $[x_1, x_2, \dots, x_n, Y_0, Y_1, \dots, Y_{i-1}]$)
> $Y_0 = \text{_START}$
 - $y_i = f_2(h_{n+i})$ (used to compute loss between y_i and Y_i) $f_2 \sim \text{softmax}(W_{hy} * h_{n+i})$
- Decoder (inference)
 - $h_{n+i} = f_1(h_{n+i-1}, y_{i-1})$ (embedding of $[x_1, x_2, \dots, x_n, y_0, y_1, \dots, y_{i-1}]$)
> $y_0 = \text{_START}$
 - $y_i = f_2(h_{n+i})$

Remarks on RNNs

- Drawbacks of RNN-based translation

- (1) Encoding and decoding are sequential

- (2) Information loss for long sequences

- > In principle, encoder-decoder RNN architectures allow the decoder to see the entire input sequence before producing any output
 - > However, signal from first few words is lost by end of long sequence
 - > Experience: RNN-based translation works only for sentences of 4-5 words and if languages are well aligned

- Solution: transformer

- (1) Create encoding of sequence in **parallel**

- (2) **Attention**: pick up important signals for a given word from *anywhere* in input sequence

- > Example: The **boy** stood on the burning deck whence all but **he** had fled.



Summary

- Standard presentations of gradients and back propagation
 - Biological metaphors like neurons and synapses come in the way
 - Properties of activation functions are distraction: enough to know we can compute value and gradient at any point
- Presentation in this lecture
 - **Abstraction** for neural networks: parameterized programs
 - Gradient computation
 - > Abstractly (what?): **sum over paths** (sum of products)
 - > Efficient computation (how?): compositional algorithm on dataflow graph representation
 - Handles complex neural networks with weight-sharing and irregular interconnections (such as “skip connections”) smoothly
- Extension to **Kolmogorov-Arnold Networks (KANs)** and **Recurrent Neural Networks (RNNs)**