

MONTH 20XX



PREFIXRL

Using RL to Optimize Parallel Prefix Circuits
Work done By: R. Roy et al. @ NVIDIA

PRESENTED BY: KAYVAN MANSOORSHAHI

The University of Texas at Austin

Outline

- Prefix sum
- PrefixRL
 - Goals
 - Problems + solutions
- Results

Scan

- For input \mathbf{X} , and operator $*$, Scan results in \mathbf{Y} defined as:

$$y_i = x_i * x_{i-1} * x_{i-2} * \dots * x_0 \text{ for } 0 \leq i < N$$

- If we take $*$ as addition we get prefix sum

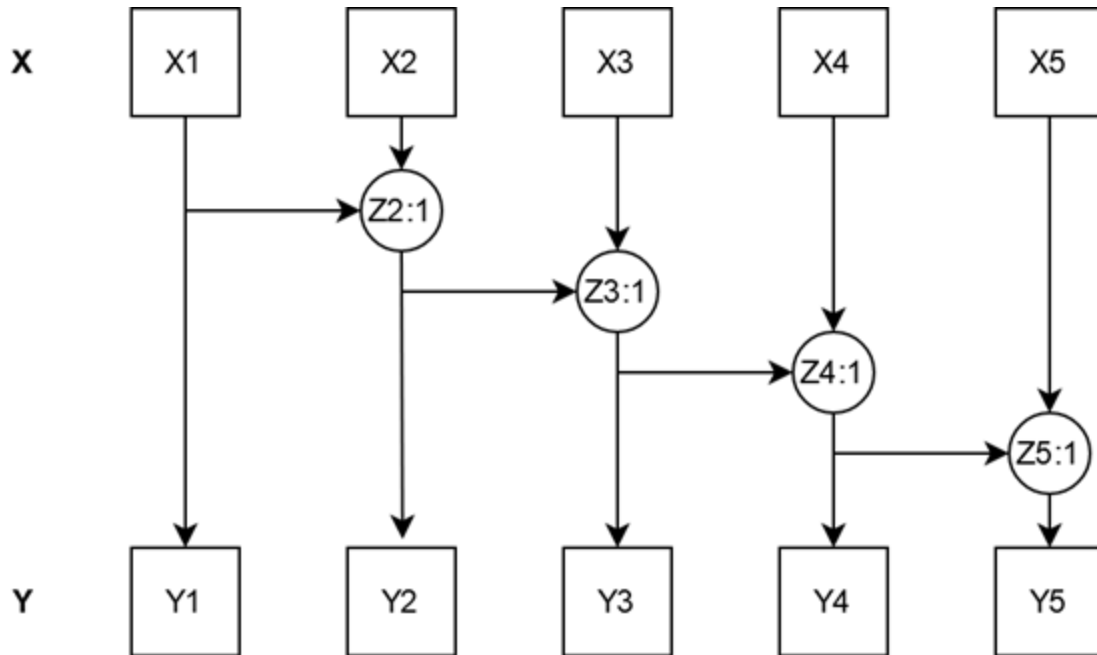
Prefix Sum Example

- $X = [1, 2, 3, 4]$
- $Y = [1, 1+2, 1+2+3, 1+2+3+4]$
- Really easy sequentially
 - $Y[i] = Y[i-1] + X[i]$ for $2 \leq i \leq |X|$
 - $Y[1] = X[1]$


Prefix Sum

- Since $*$ is associative we can divide and conquer, and reuse results
 - We will call these intermediate results $z_{j:i} = x_j * \dots * x_i$
- Inputs (**X**) can be thought of as intermediate product $z_{i:l}$
- Outputs (**Y**) can be thought of as intermediate product $z_{i:0}$
- All $z_{i:j}$ can be split into $z_{i:k}$, and $z_{k-1:j}$
 - $z_{i:k}$ is referred to as the “upper parent” and $z_{k-1:j}$ as the “lower parent”

Prefix Sum Graphs



Prefix Sum Graphs




Adders as Prefix sum

- Adders are ubiquitous circuits. Optimizing PPA is crucial
 - Can be expressed as a prefix sum of bits
 - Can also be parallelized to improve performance
 - SOTA has been Kogge-Stone adder since 1973. Can we do better?
-
- State space is huge ($O(2^{N^2})$), need to explore intelligently
 - PPA of a circuit is expensive to collect (~36 sec per circuit)

PrefixRL

- Large unknown space? RL!
- State space is all LEGAL circuits
 - A matrix: (i,j) means generating $z_{i,j}$
- Action is add or remove (i,j)
- Value is $\langle \text{area}, \text{delay} \rangle$
 - Getting reward is expensive
 - More than one reward




Legal Circuits Only

- Adding arbitrary intermediate can result in invalid circuit
- Add a legalization step before evaluating the new state
- For each node (MSB_{node} , LSB_{node}):
 - Upper parent = get node (MSB_{node} , max)
 - Lower parent = find node ($LSB_{Upper\ parent}-1$, LSB_{node})
- Upper parent will always exist since starting adder is legal
- Lower parent will be added by legalizer if it doesn't exist


Legal Circuits Only

- Starting state
- Sequential RCA




Legal Circuits Only

- Agent says add 3,2
- Parents are picked
 - (3:3) is upper parent
 - (2:2) is lower parent
- Adder is now invalid



Legal Circuits Only

- Visit 3,0 node
- Apply rules to validate
- Legal again



Legalizer

- Minlist: non-lower parent nodes
- Nodelist: all nodes
- Adding existing nodes is banned
- Delete can only touch min list
 - Legalizer will reverse it otherwise

Algorithm 1: PrefixRL N-input prefix graph actions

```

Function Initialize:
    nodelist  $\leftarrow \emptyset$ , minlist  $\leftarrow \emptyset$ ;
    for m  $\leftarrow 0$  to (N - 1) do           // add in/out nodes
        | add (m, m), (m, 0) to nodelist
    end
Function Add(msb,lsb):
    add (msb,lsb) to minlist;
    // remove new node's and child's lps from minlist
    for l  $\leftarrow$  (msb - 1) to 0 do
        | if (msb, l) is in minlist then
            | | delete lp(msb, l) from minlist
        | end
    end
    Legalize()
Function Delete(msb,lsb):
    delete (msb,lsb) from minlist;
    Legalize()
Function Legalize:
    nodelist  $\leftarrow$  minlist;
    for m  $\leftarrow 0$  to (N - 1) do           // add in/out nodes
        | add (m, m), (m, 0) to nodelist
    end
    for m  $\leftarrow$  (N - 1) to 0 do           // add missing lps
        | for l  $\leftarrow$  (m - 1) to 0 do
            | | if (m, l) is in nodelist then
                | | | add lp(m, l) to nodelist
            | | end
        | end
    end
    end
    
```

Multiple Objectives

- Both area and delay are important
- Often at odds with each other
- Simple Norm won't work
- Scalarized deep Q-learning

$$\mathbf{r}_t = [\text{area}(s_t) - \text{area}(s_{t+1}), \text{delay}(s_t) - \text{delay}(s_{t+1})]$$

$$\mathbf{y}_t = \mathbf{r}_t + \gamma \mathbf{Q}(s_{t+1}, \underset{a}{\operatorname{argmax}}[\mathbf{w}^\top \mathbf{Q}(s_{t+1}, a; \theta_t)]; \theta'_t)$$


$$\text{Loss}(\mathbf{Q}(s_t, a_t; \theta_t), \mathbf{y}_t)$$

$$a_t = \underset{a}{\operatorname{argmax}}[\mathbf{w}^\top \mathbf{Q}(s_t, a; \theta_t)]$$

- Network learns each separately
- \mathbf{w} used for action selection
- \mathbf{w} picked for objective preference

Circuit Synthesis in the Loop

- What is the area and delay of a graph?
- Could count nodes
 - Doesn't account for optimizations
 - Doesn't account for route
 - Doesn't account for fanout



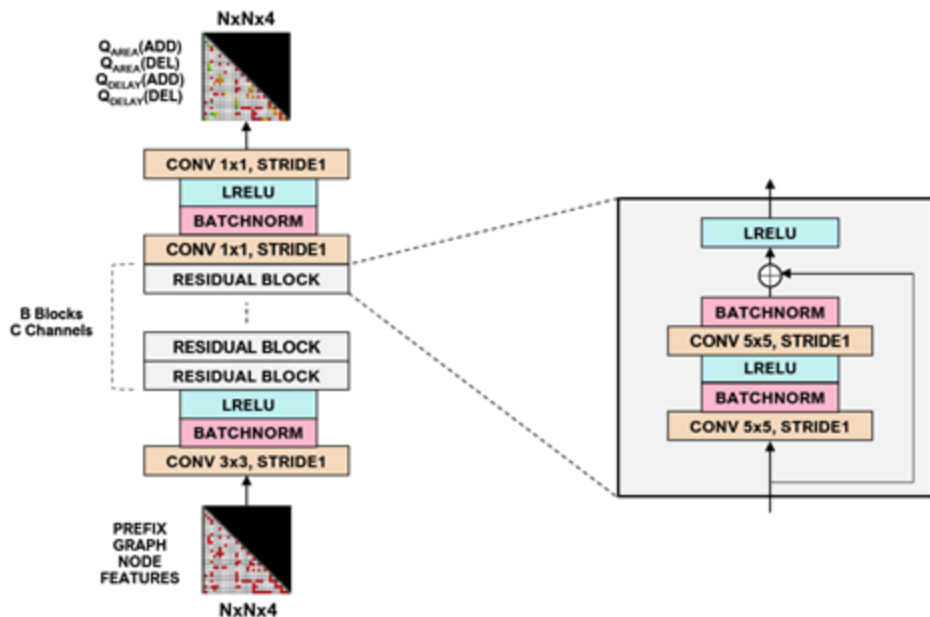
- Synthesis in the loop
- Circuit synthesis is slow

Circuit Synthesis in the Loop

- Synth is long running, independent work
- Distribute, collect results, and learn
- They use Deep Q network (DQN)
- DQN is off policy, can use samples after Θ changes
- DQN can separate sample collection and training
 - Need to read more. Don't fully understand why this is the case
- picked 192 workers to cover latency

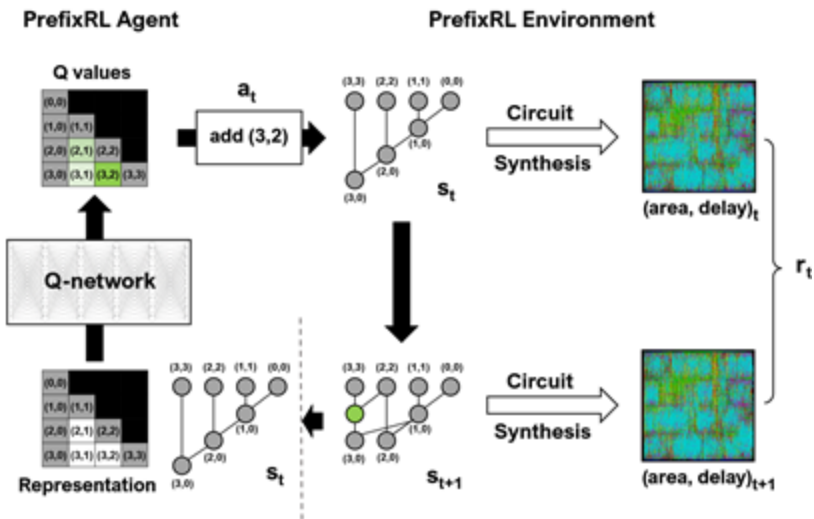
Q-Network Architecture

- Q network shown left
- Inputs are $N \times N \times 4$
 - node list
 - min list
 - level/depth of node
 - fanout of node
- Outputs are $N \times N \times 4$
 - Q_{Area} for add/delete
 - Q_{Delay} for add/delete

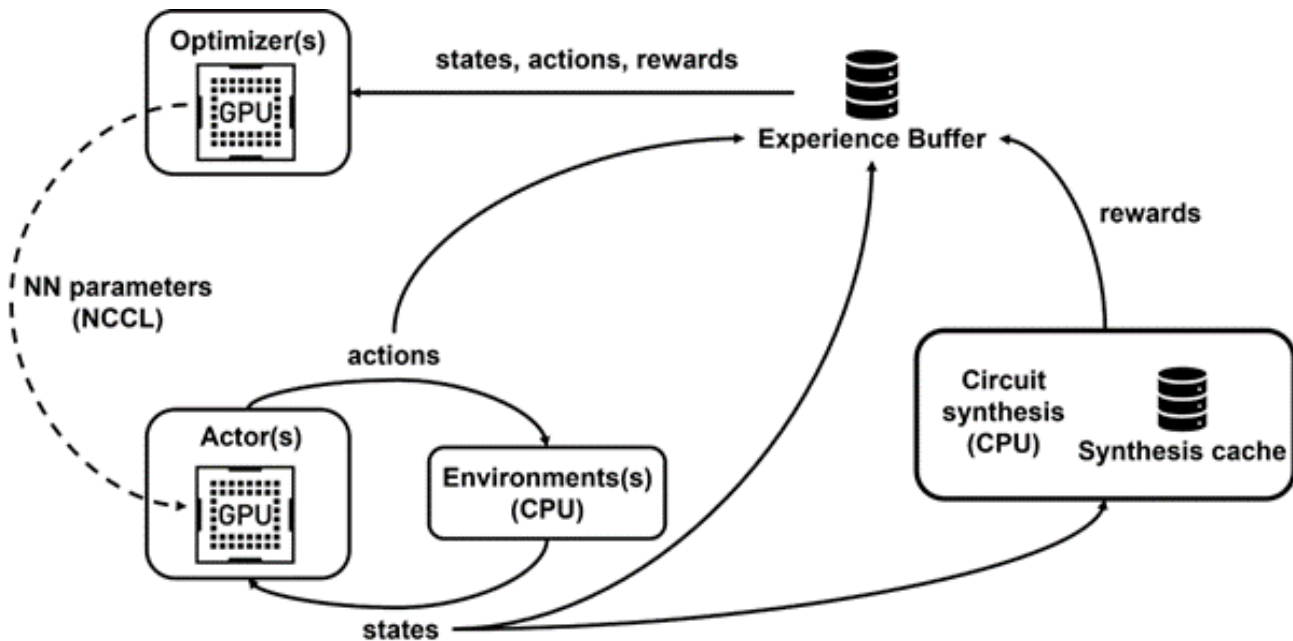


Training

- Aren't specific, but my best guess
- Main loop will
 - Query Q values of Actions using Θ
 - Act
 - Legalize
 - Ship to synth farm
 - Update Θ' using sample from buffer
 - After 60 steps, copy Θ' to Θ
- The synth farm will
 - Return synthesis result to replay buffer
- Again, need to understand DQN




Training

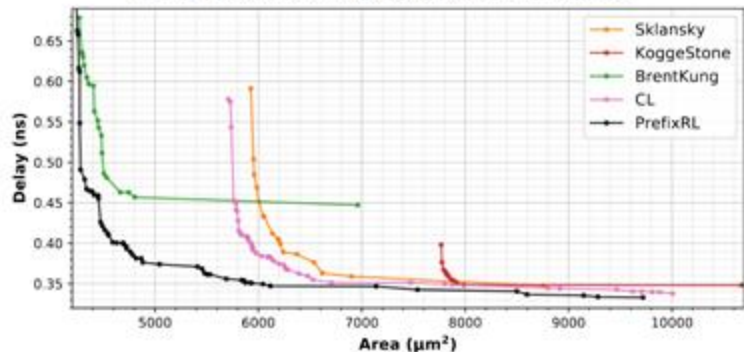


Results

- Sklansky, BrentKung, KoggeStone are all SOTA human results
 - “Sit down and think”
- SA is simulated annealing
- PS is pruned search
- CL is machine learning
- 32b saves max 16% area
- 64b saves max 30% area



(a) 32b Adder Synthesis. OpenPhySyn, Nangate45.



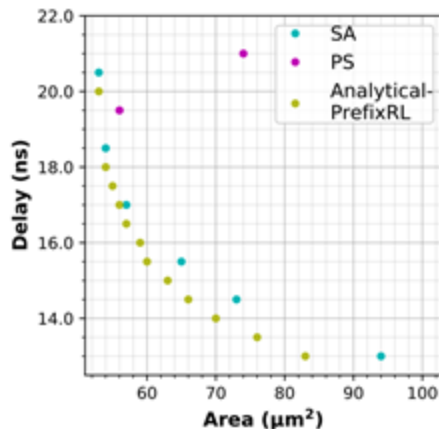
(b) 64b Adder Synthesis. OpenPhySyn, Nangate45.

Conclusions


- RL is effective for design space exploration
- Analytical steps can reduce the space (Legalization)
- Can learn complex post processing steps (Synthesis)
- Can amortize cost of expensive rewards (Distributed training)
- When learning multiple objectives
 - scalarize for action selection
 - otherwise keep them separate

Results mod Synthesis

- RL dominates with heuristics too
 - RL beats SA with its own heuristics
- These heuristics don't translate
 - Still RL search the space "better"



(a) Analytical Metrics.



(b) Synthesized Metrics.

Extras

- Reward of area and delay are different units, so can't be compared
 - They simply scale each by some factor. I guess selected empirically
 - They scale area by .001 and delay by 10
- Synth results change based on delay targets
 - Synth at 4 delay targets, interpolate, and return optimal based on \mathbf{w} vector
- Synth results are cached
 - If exploration factor is low the hit rate is 10%;50% for 64b;32b respectively

Extras

- How much did synthesis in the loop really buy?
 - Purely analytical method beats Kogge-Stone adder
 - Worse than synthesis in the loop (11% vs 30% reduction in area)
- Reinforcement learning agent can be used to search design spaces
- Expensive reward functions can be parallelized if off-policy method

Credits

[1] Roy, Rajarshi & Raiman, Jonathan & Kant, Neel & Elkin, Ilyas & Kirby, Robert & Siu, Michael & Oberman, Stuart & Godil, Saad & Catanzaro, Bryan. (2021). PrefixRL: Optimization of Parallel Prefix Circuits using Deep Reinforcement Learning. 853-858. 10.1109/DAC18074.2021.9586094. // (<https://arxiv.org/pdf/2205.07000>)